



COVID-19 Forecast Hub

Jarad Niemi

15 Oct 2020

COVID-19 Forecast Hub Team

Team: Nutchawattanachit, Serena Wang,
Nicholas Reich, Evan Ray, Jarad Niemi, Khoa Le,
Abdul Hannan Kanji, Yuxin “David” Huang,
Katie House, Estee Cramer,
Matt Cornell, Andrea Brennen, Johannes Bracher

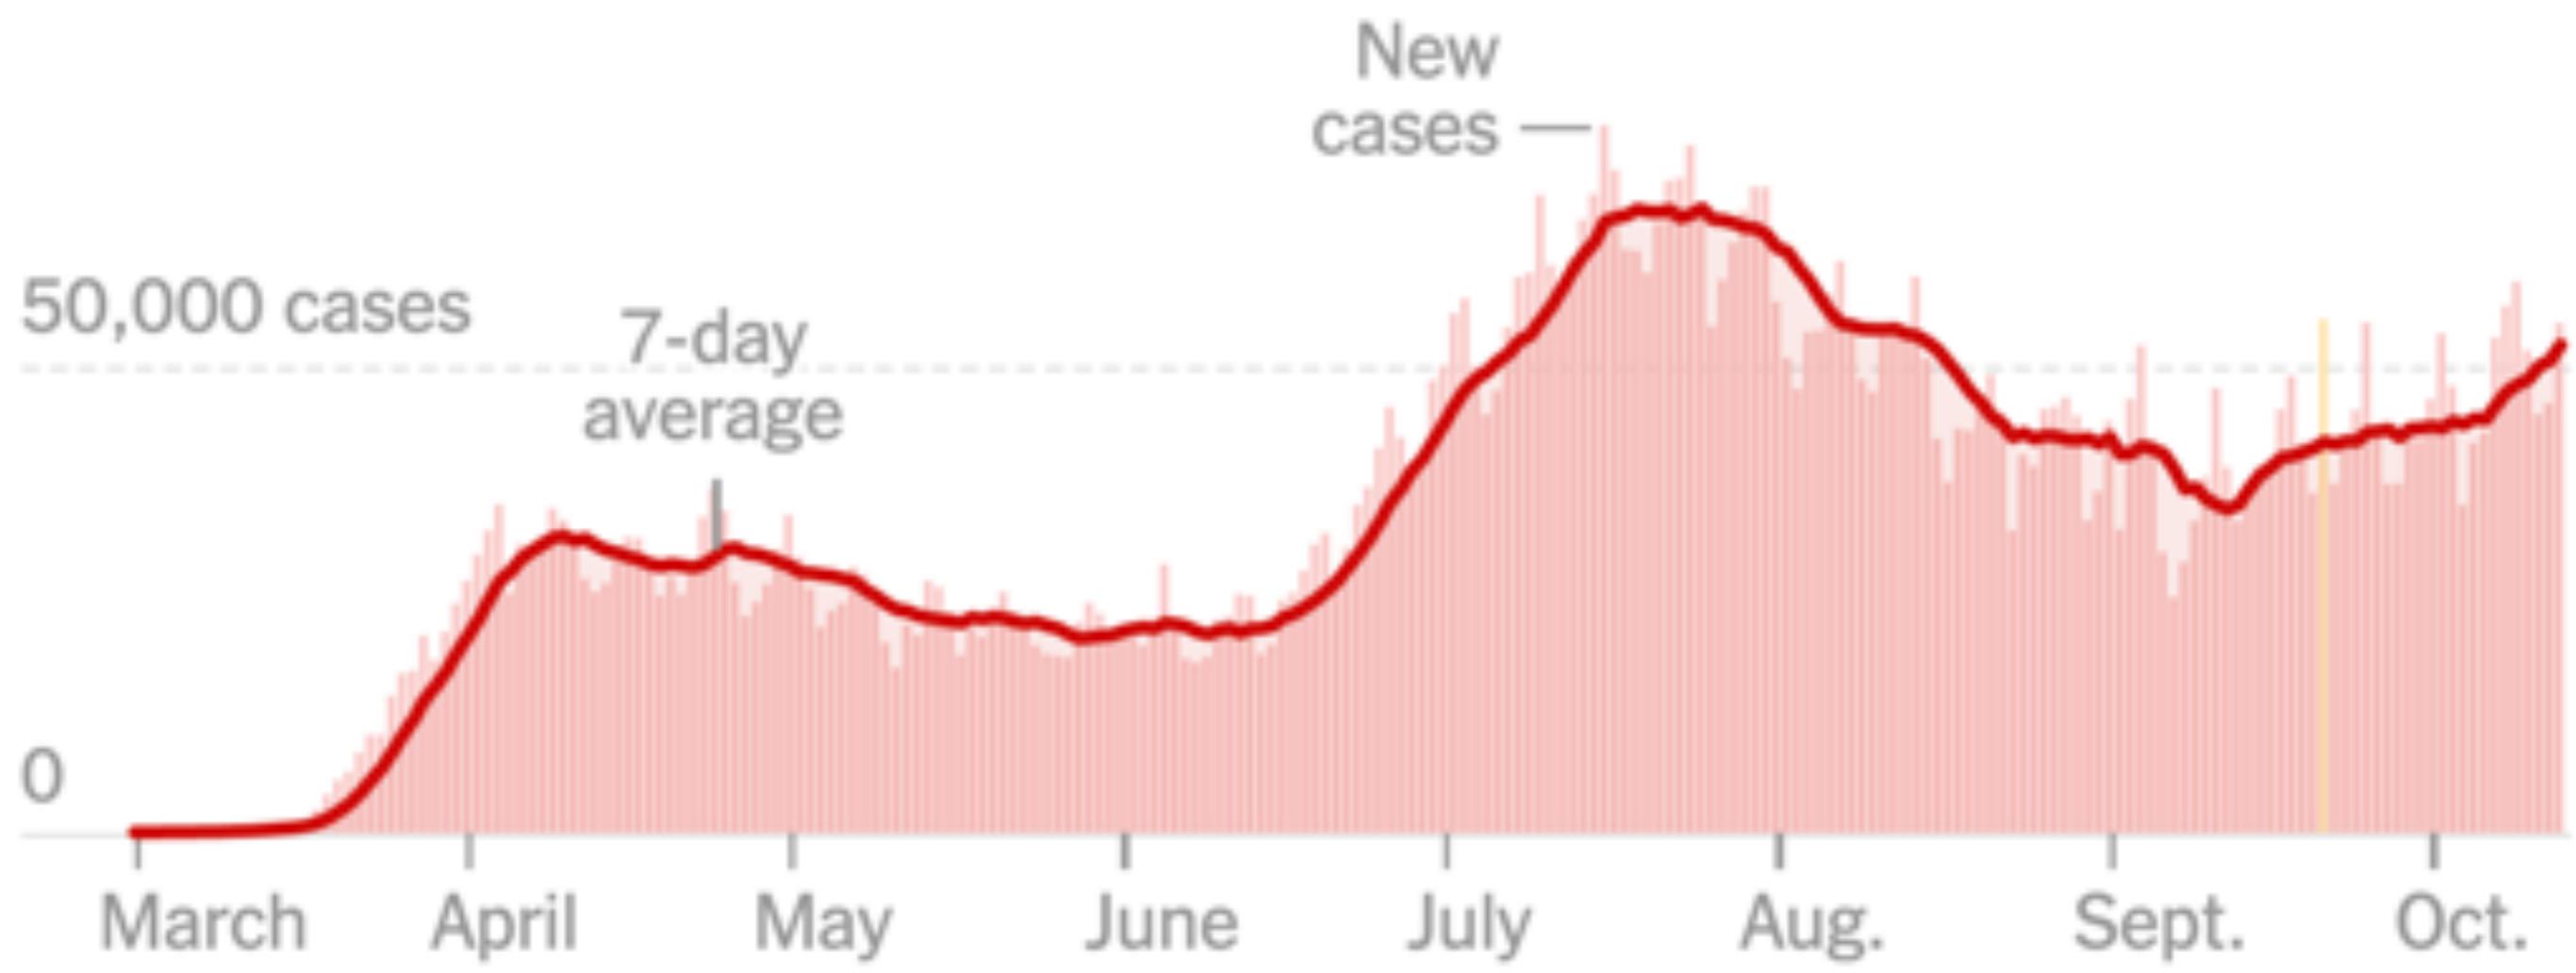
* underline denotes ensemble contributor

CDC Collaborators: Michael Johansson, Matthew Biggerstaff,
Joseph Walker, Velma Lopez, Rachel Slayton

Ensemble "advisory committee": Sebastian Funk, Tilmann
Gneiting, Anja Muhlemann, Ryan Tibshirani

COVID-19 Forecast Hub

Timeline



COVID-19 Forecast Hub

Timeline

Early 2020:
Nick enjoying
sabbatical in
Munich

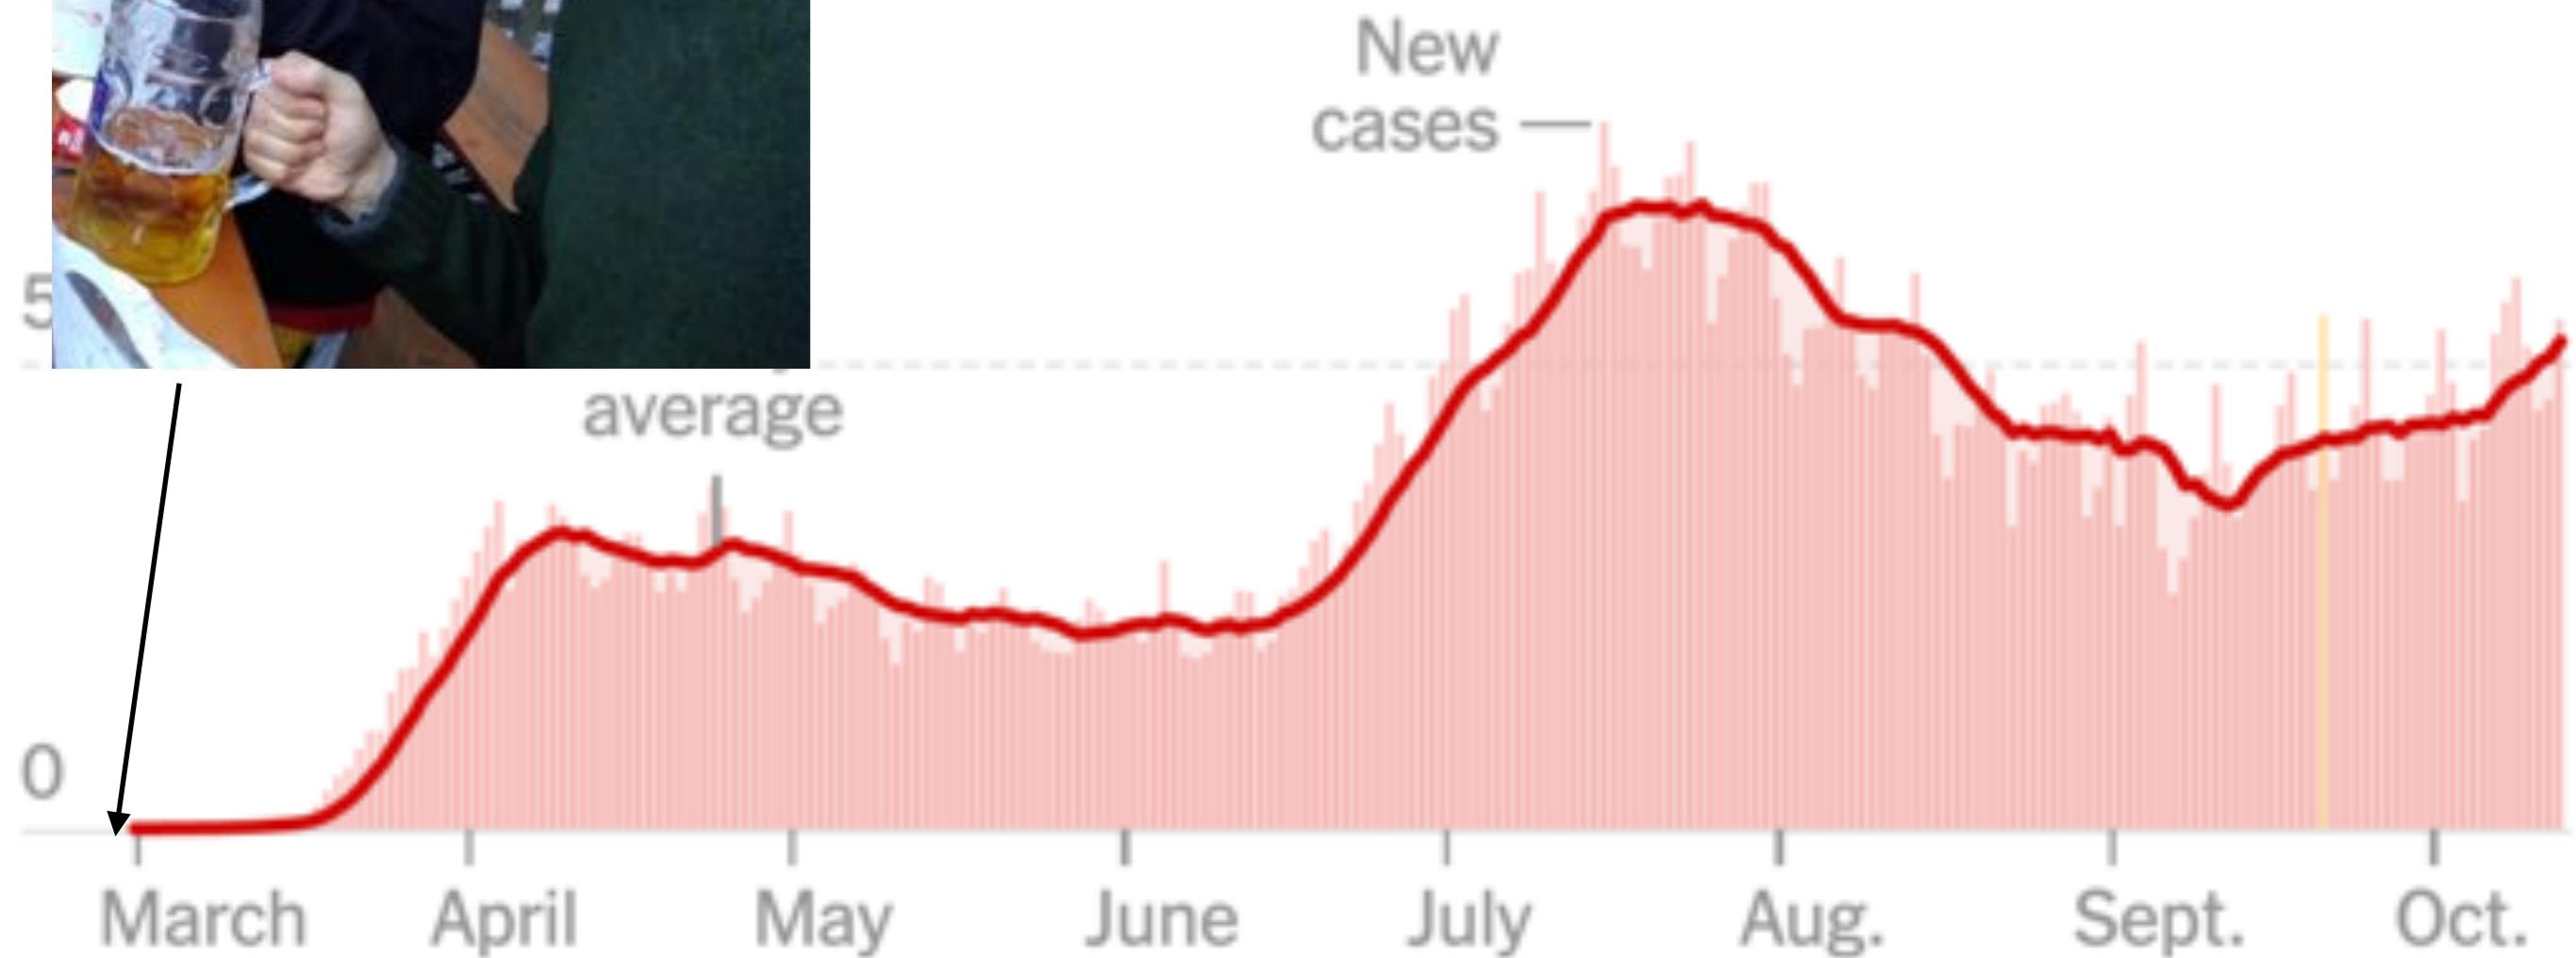


image credit: NY Times

COVID-19 Forecast Hub

Timeline

Early 2020:
Nick enjoying
sabbatical in
Munich



Jarad enjoying sabbatical
in Vancouver

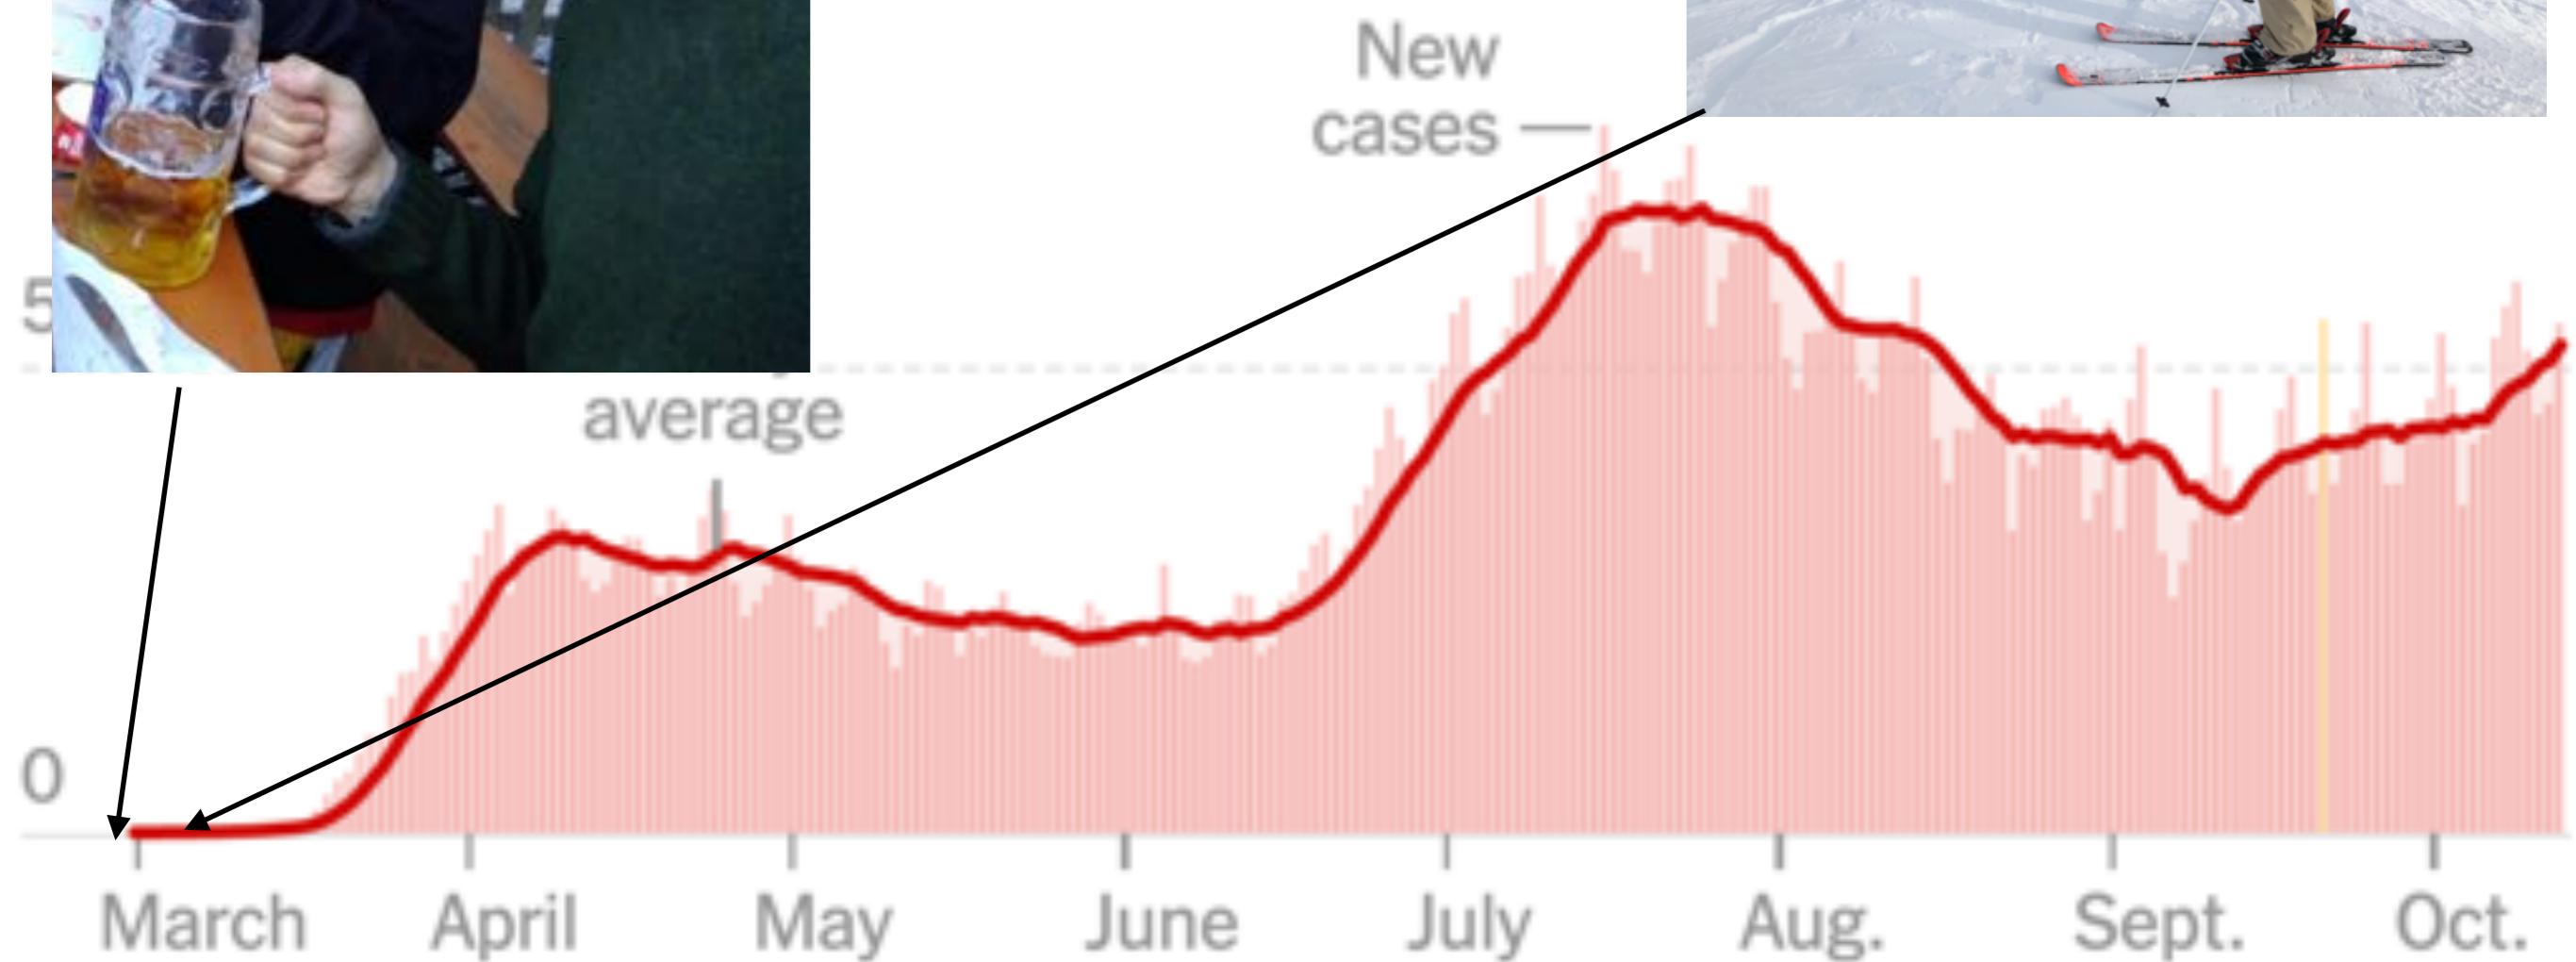


image credit: NY Times

COVID-19 Forecast Hub

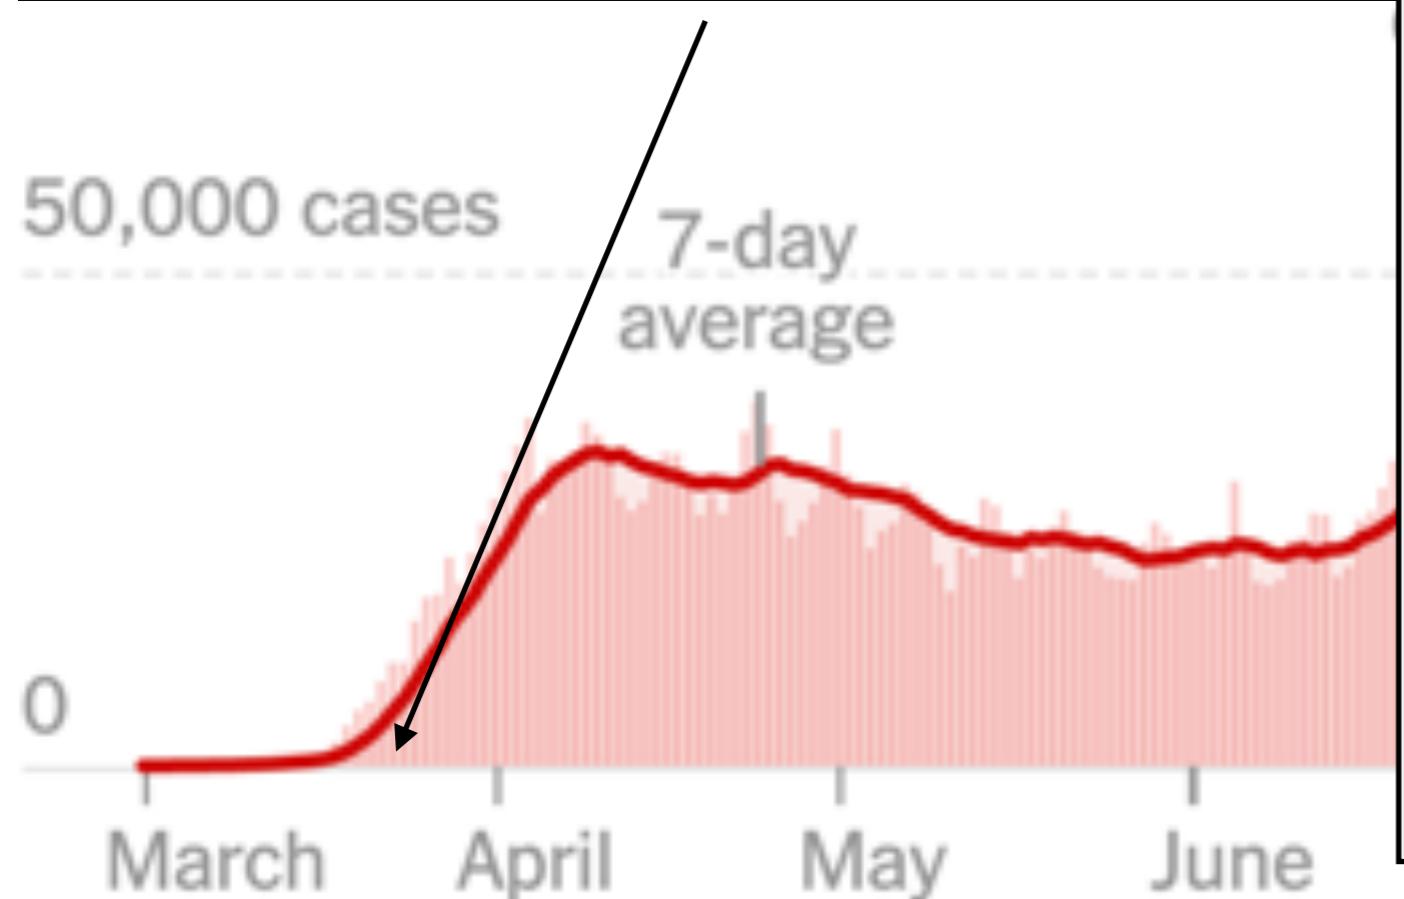
Timeline

FiveThirtyEight

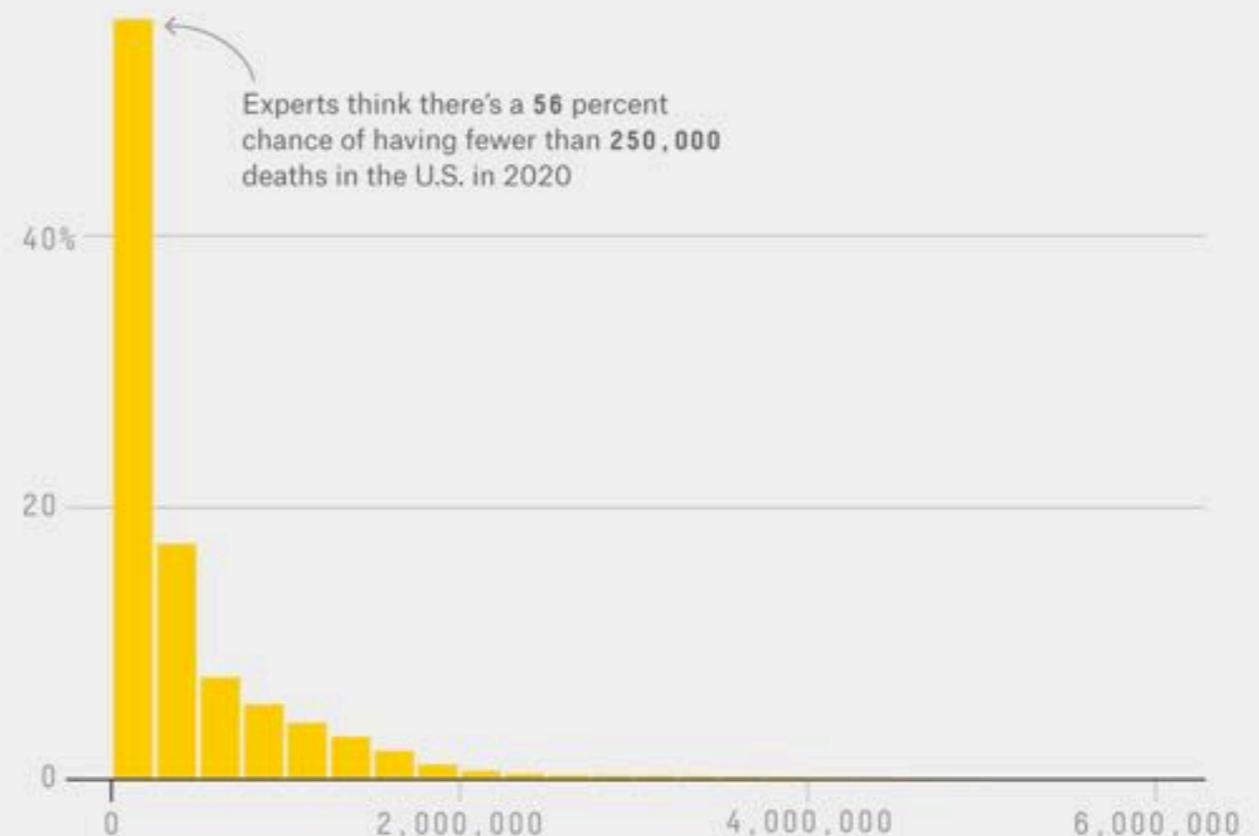
Politics Sports Science Podcasts Video

MAR. 20, 2020, AT 5:01 PM

Infectious Disease Experts Don't Know How Bad The Coronavirus Is Going To Get, Either



Around 200,000 U.S. deaths is most likely number
Probability that each number will be the total number of COVID-19 deaths in the U.S. in 2020, according to experts' estimates



FiveThirtyEight

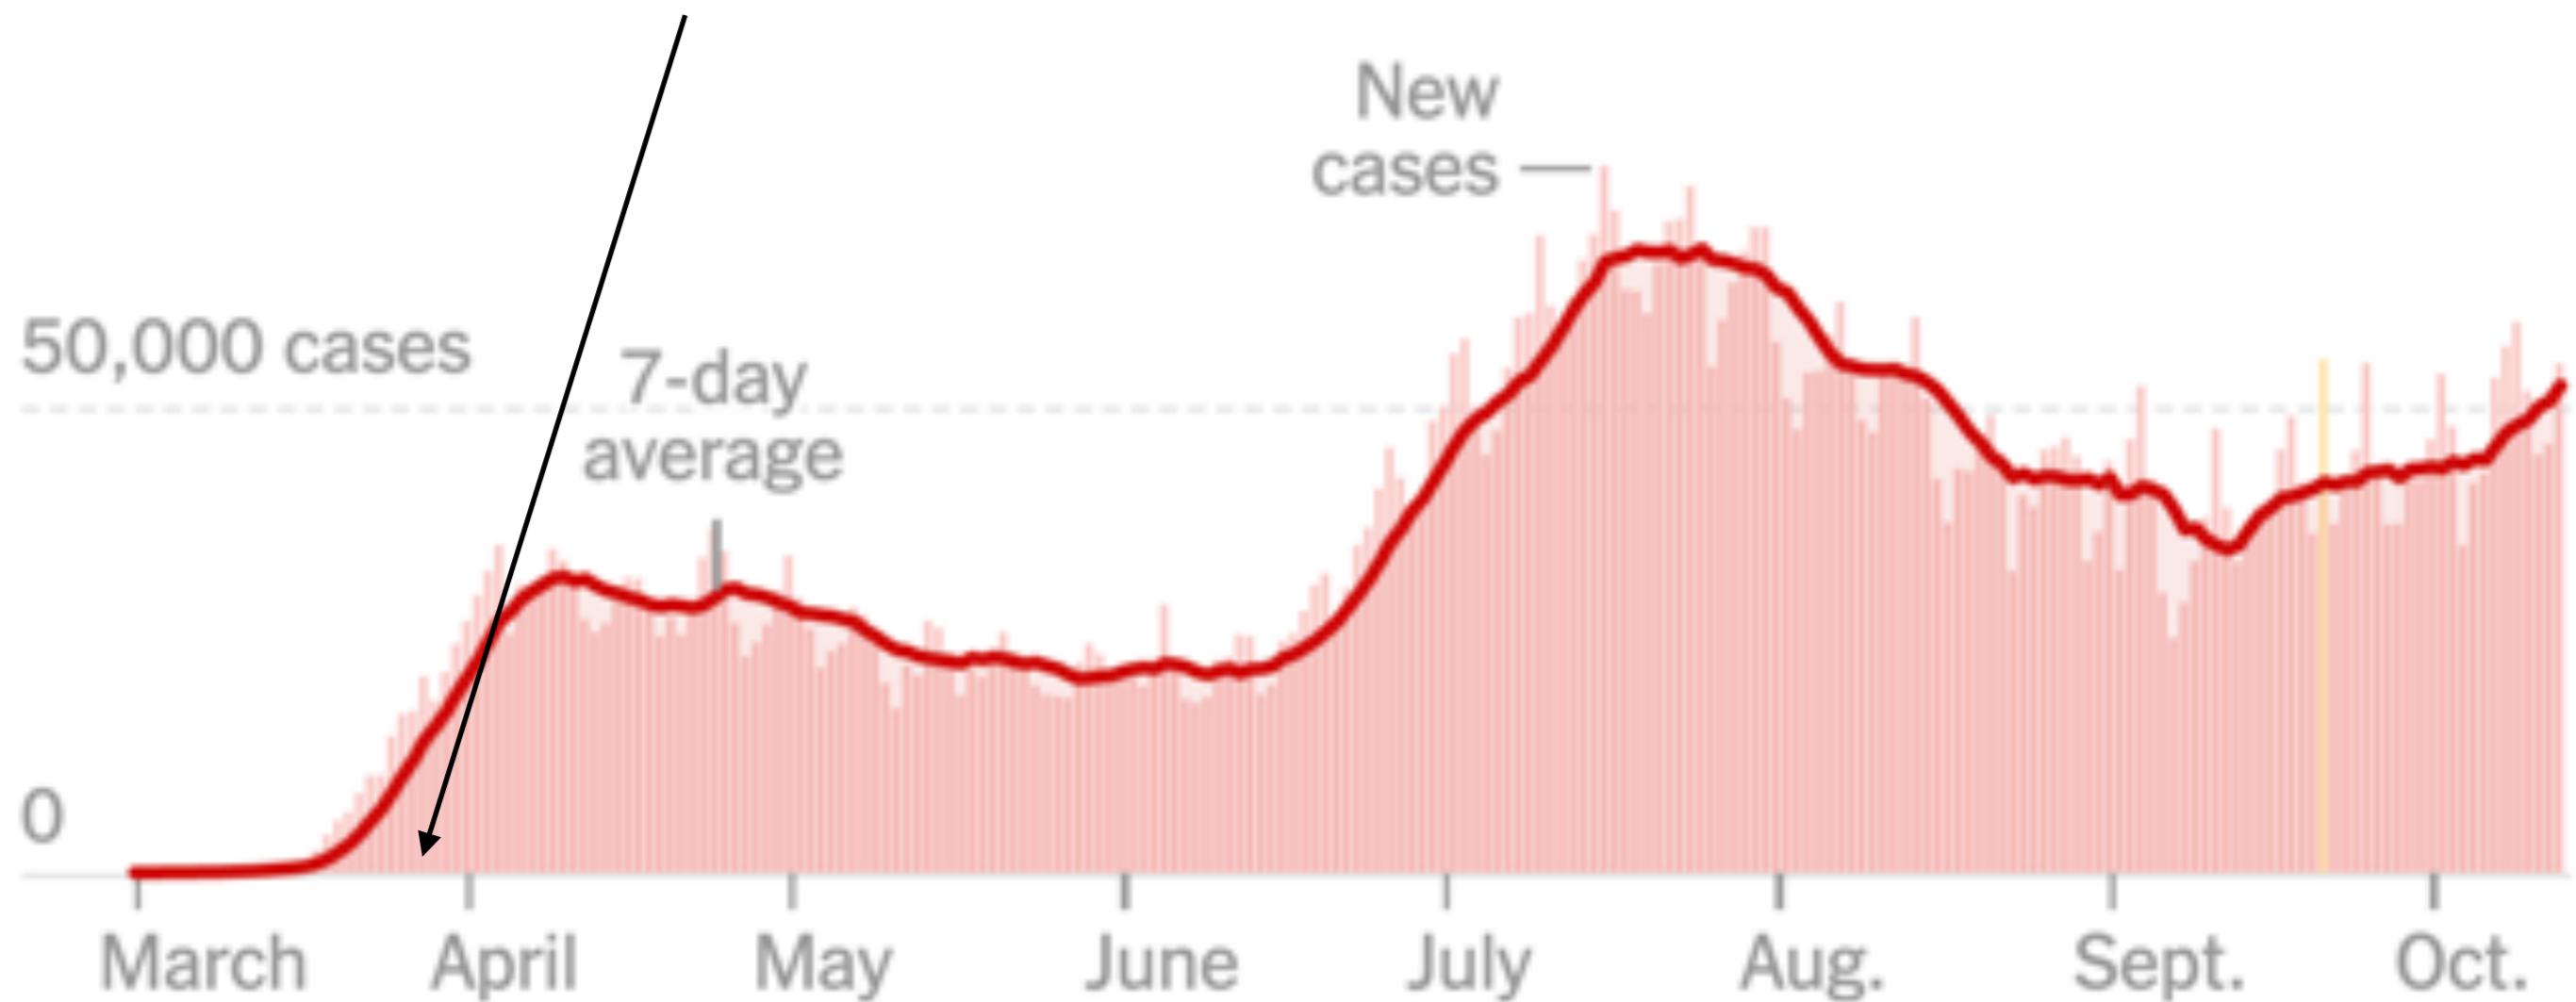
SOURCE: UNIVERSITY OF MASSACHUSETTS AMHERST

image credit: NY Times

COVID-19 Forecast Hub

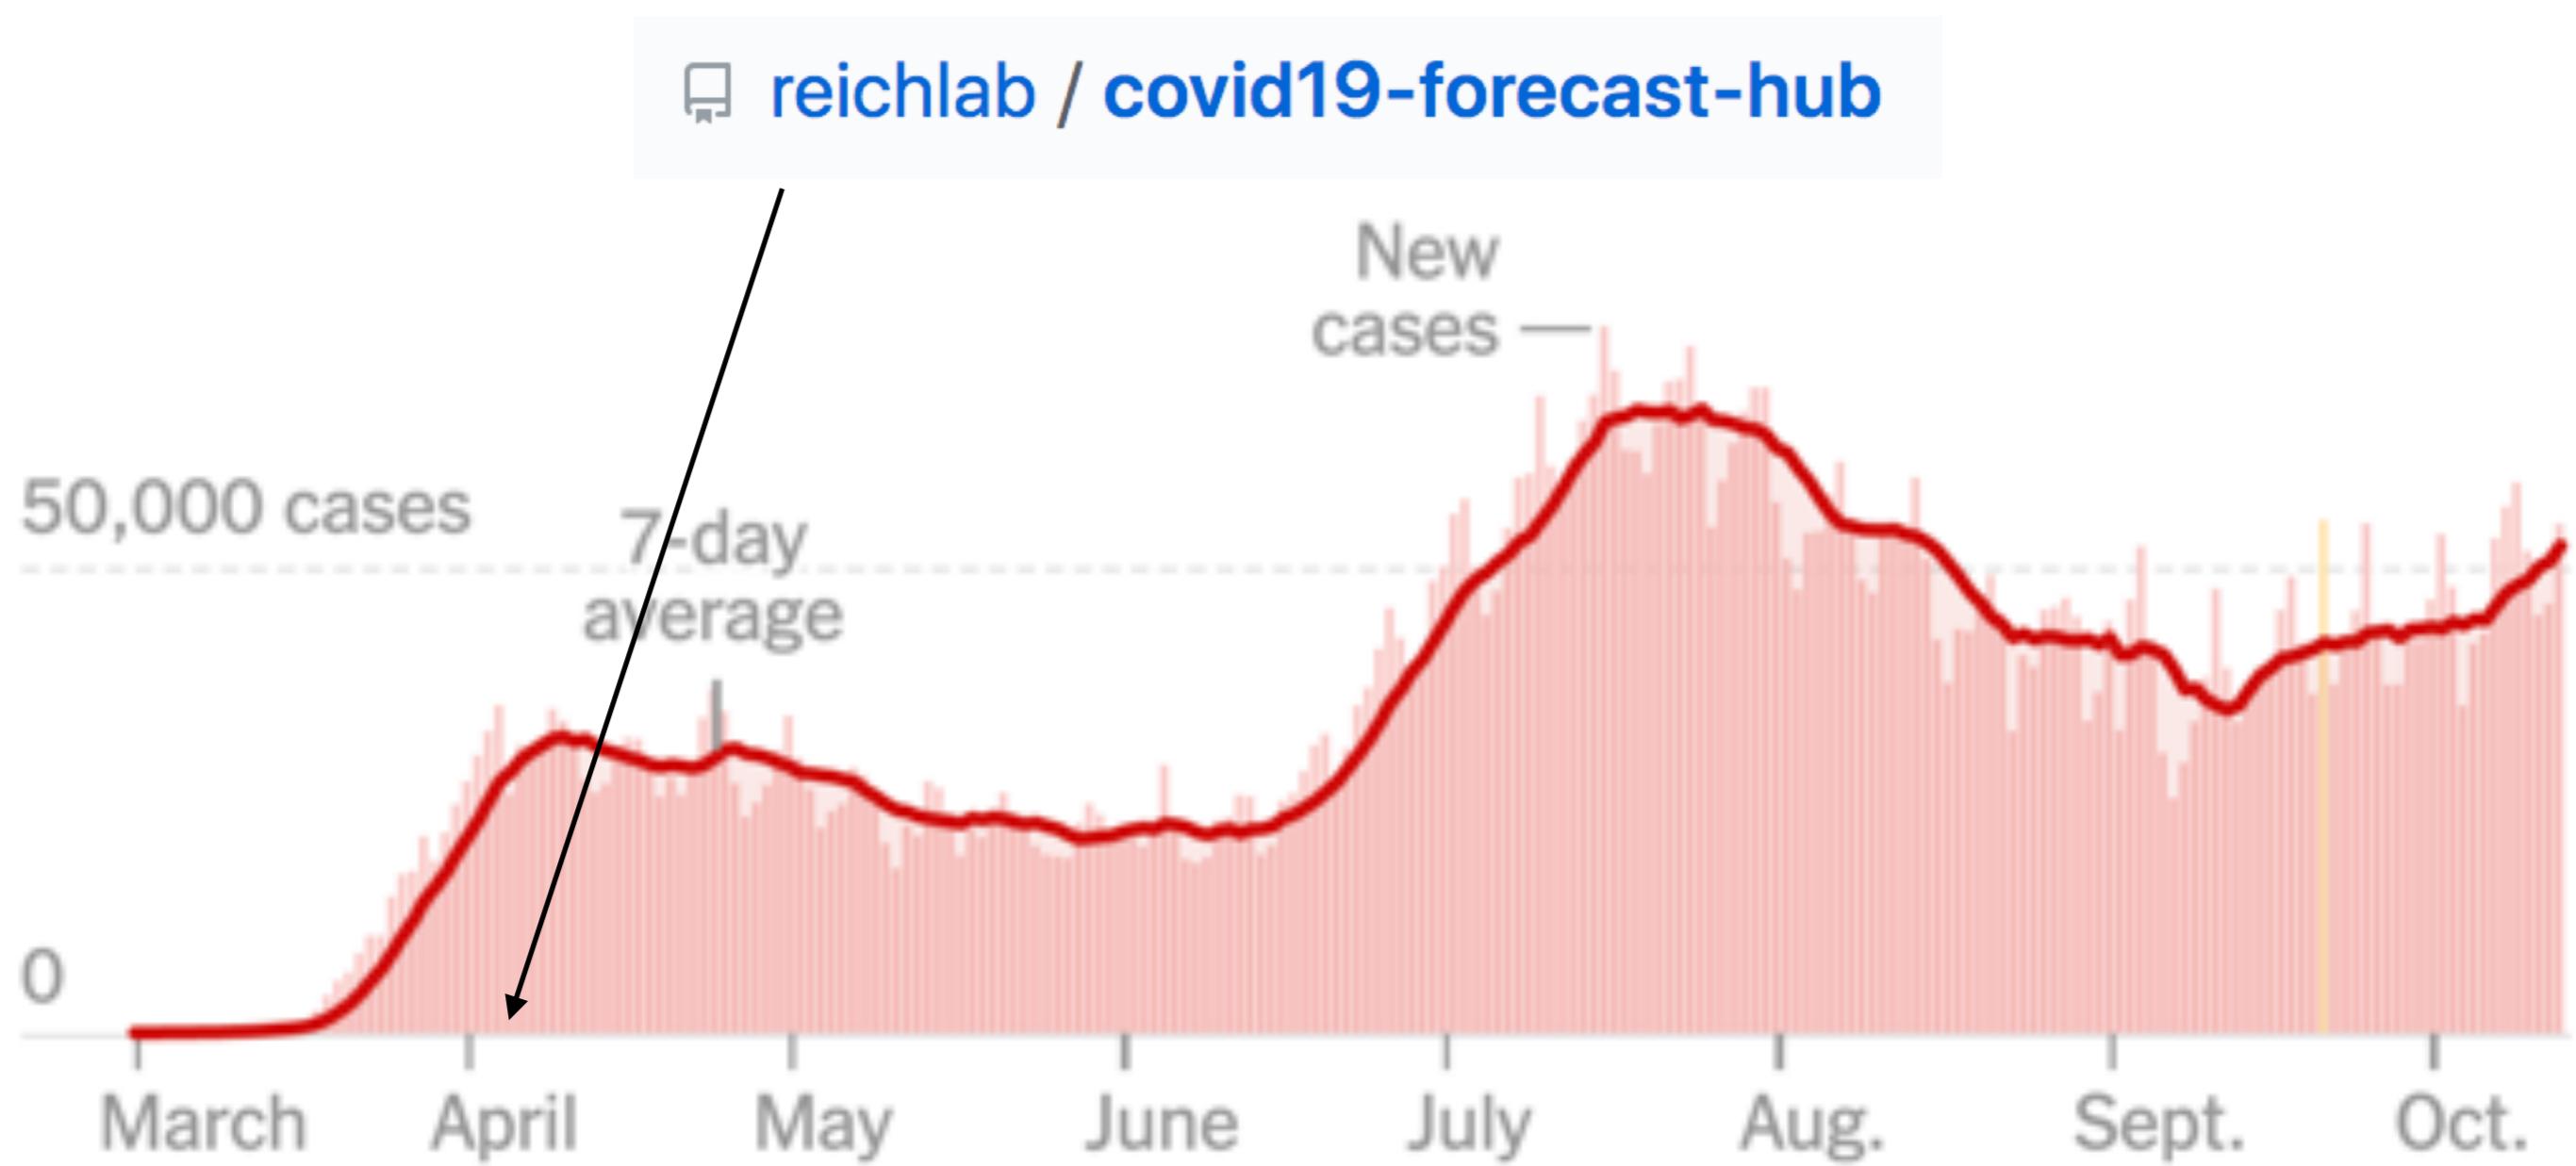
Timeline

Late March: First projections from Columbia University and IHME
Early April: First projections from Los Alamos Nat'l Labs



COVID-19 Forecast Hub Timeline

April 5, 2020: COVID-19 Forecast Hub GitHub repository created



COVID-19 Forecast Hub Timeline

April 13, 2020:
first Hub forecasts
displayed on CDC website.
3 teams.

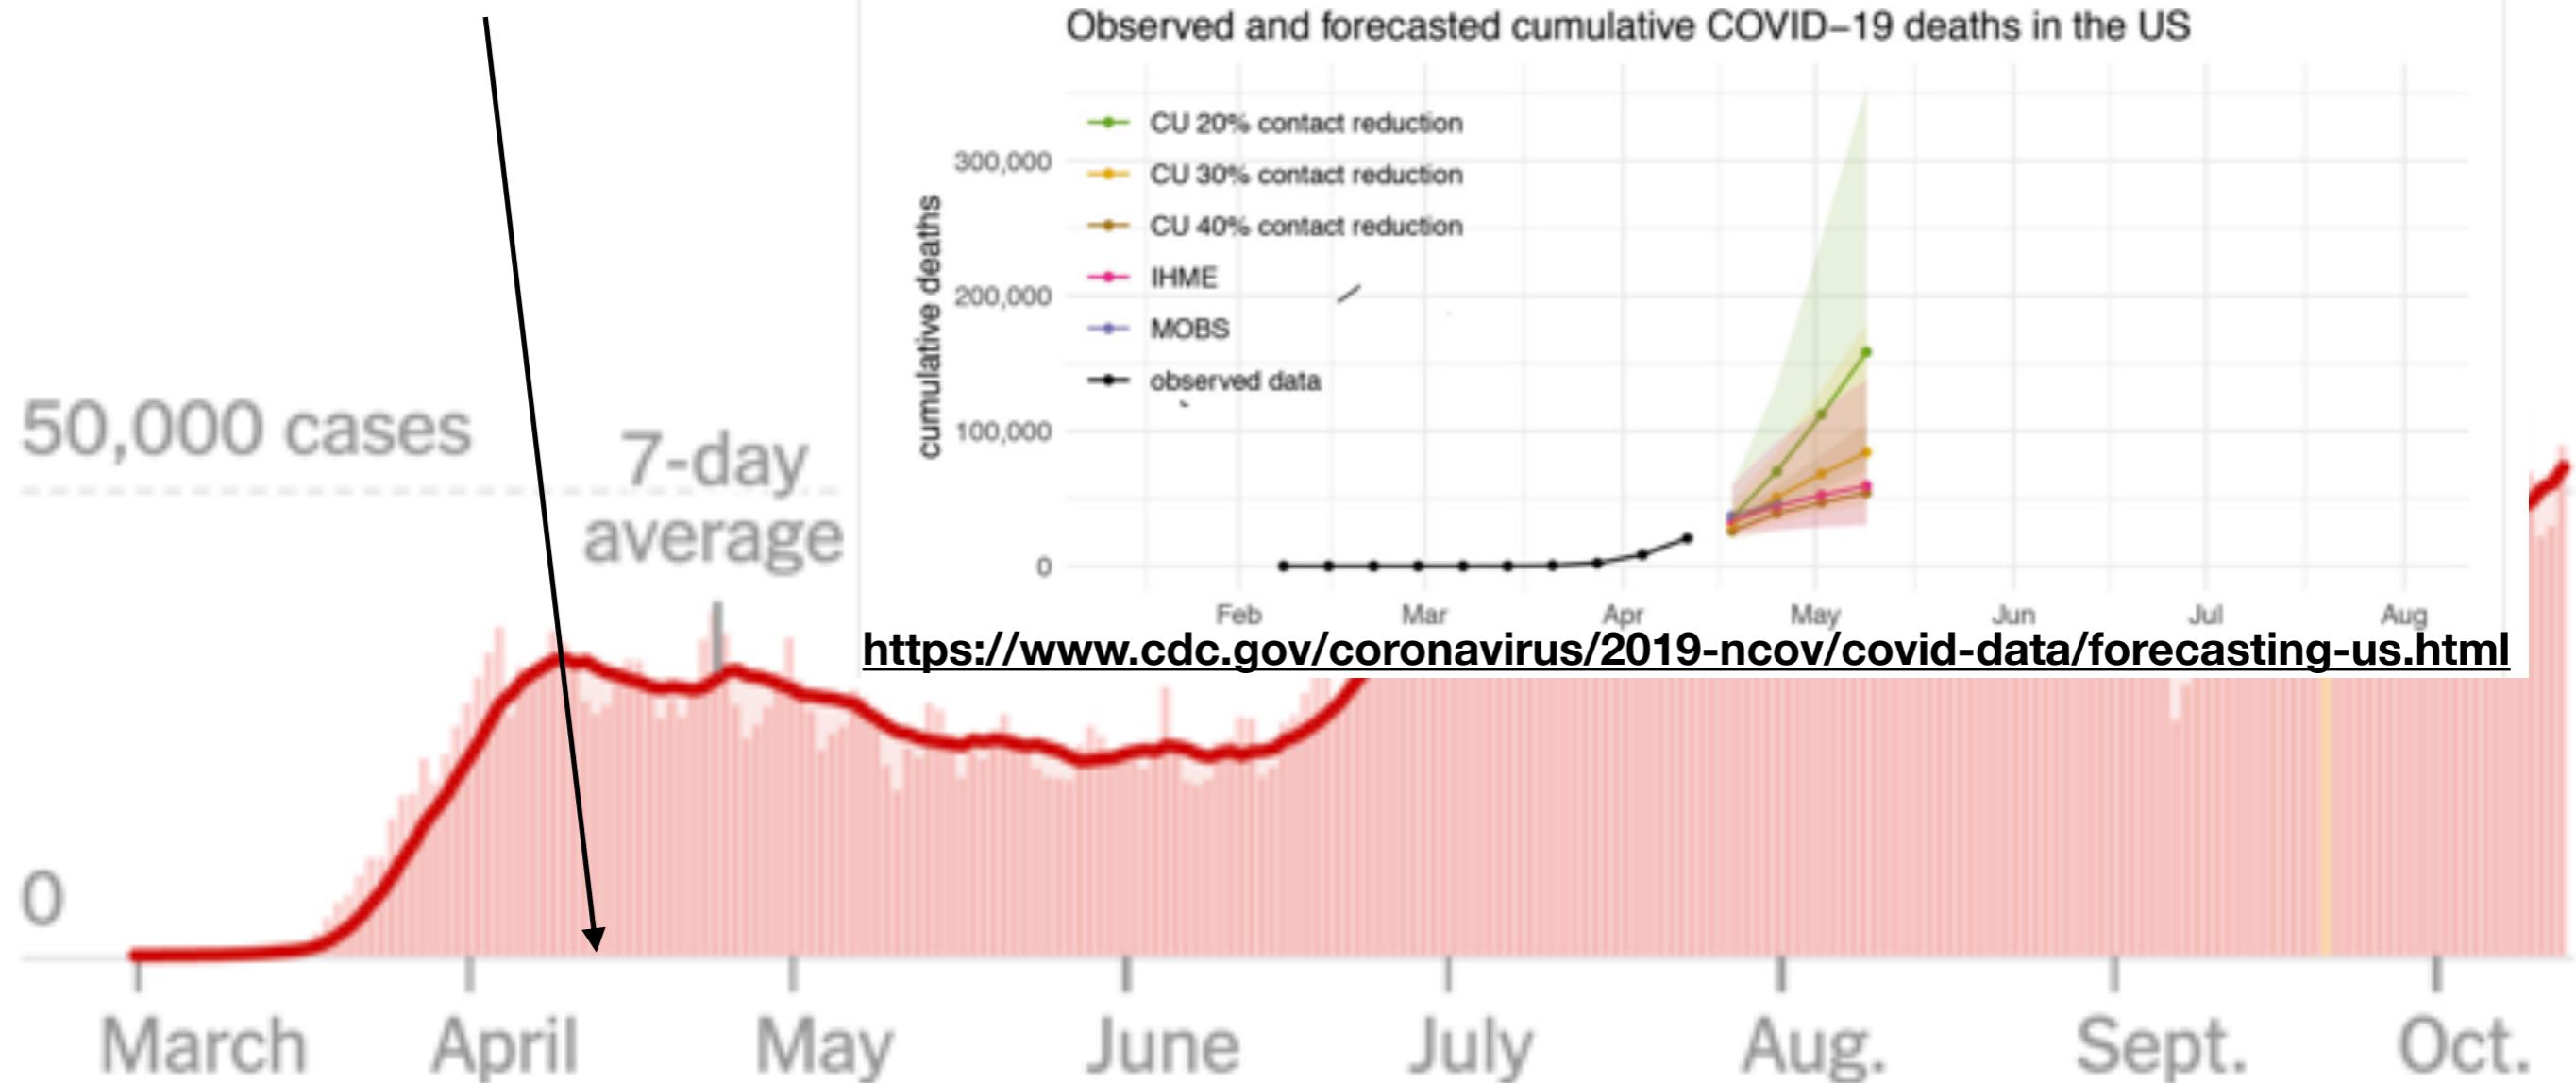
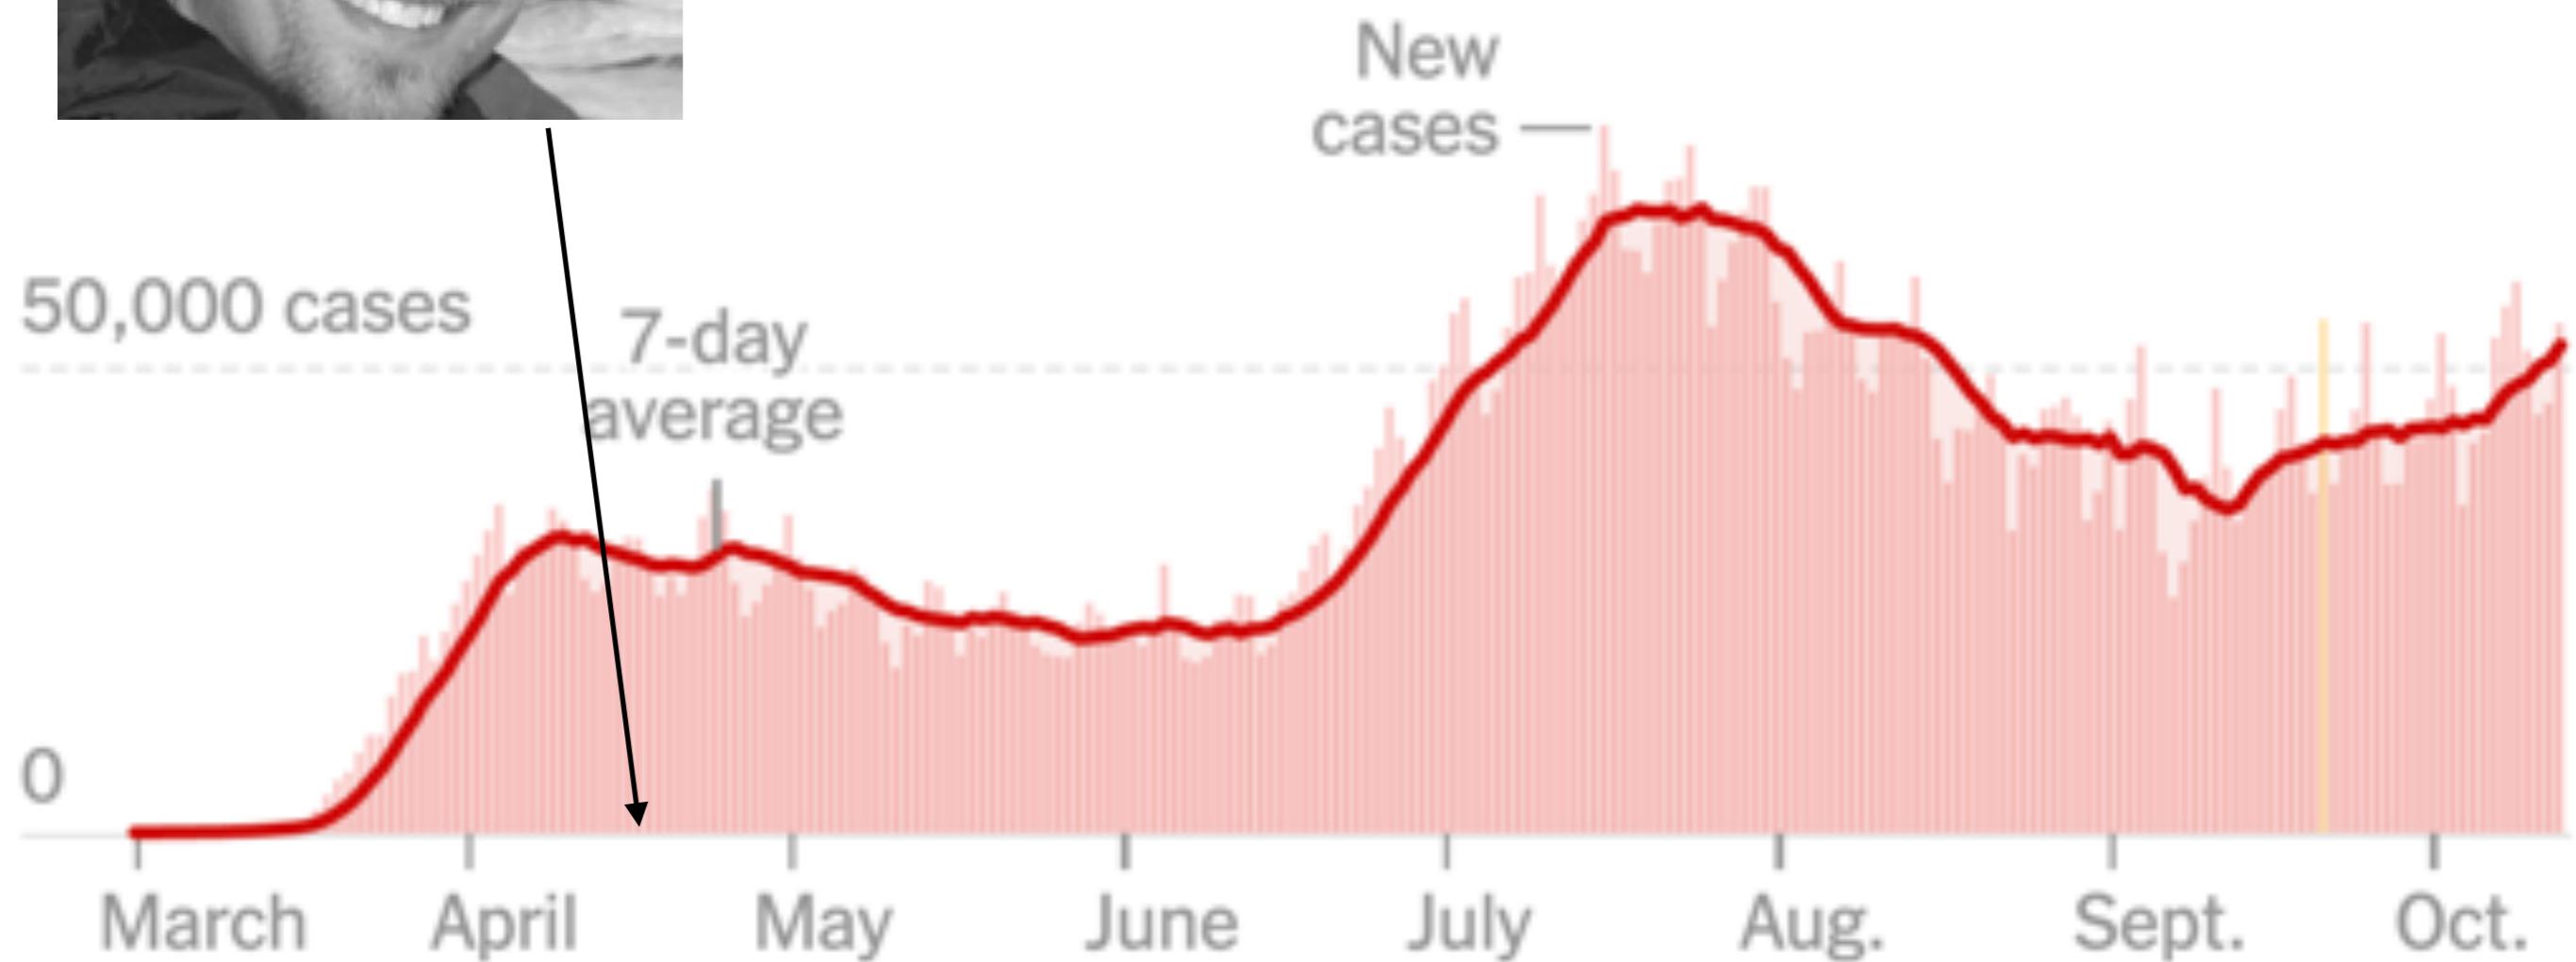


image credit: NY Times

COVID-19 Forecast Hub Timeline



**April 22, 2020:
Jarad Niemi joins
COVID-19
Forecast Hub**



COVID-19 Forecast Hub

Timeline

April 27, 2020:
first Hub ensemble
forecast on CDC
website.
6 teams.

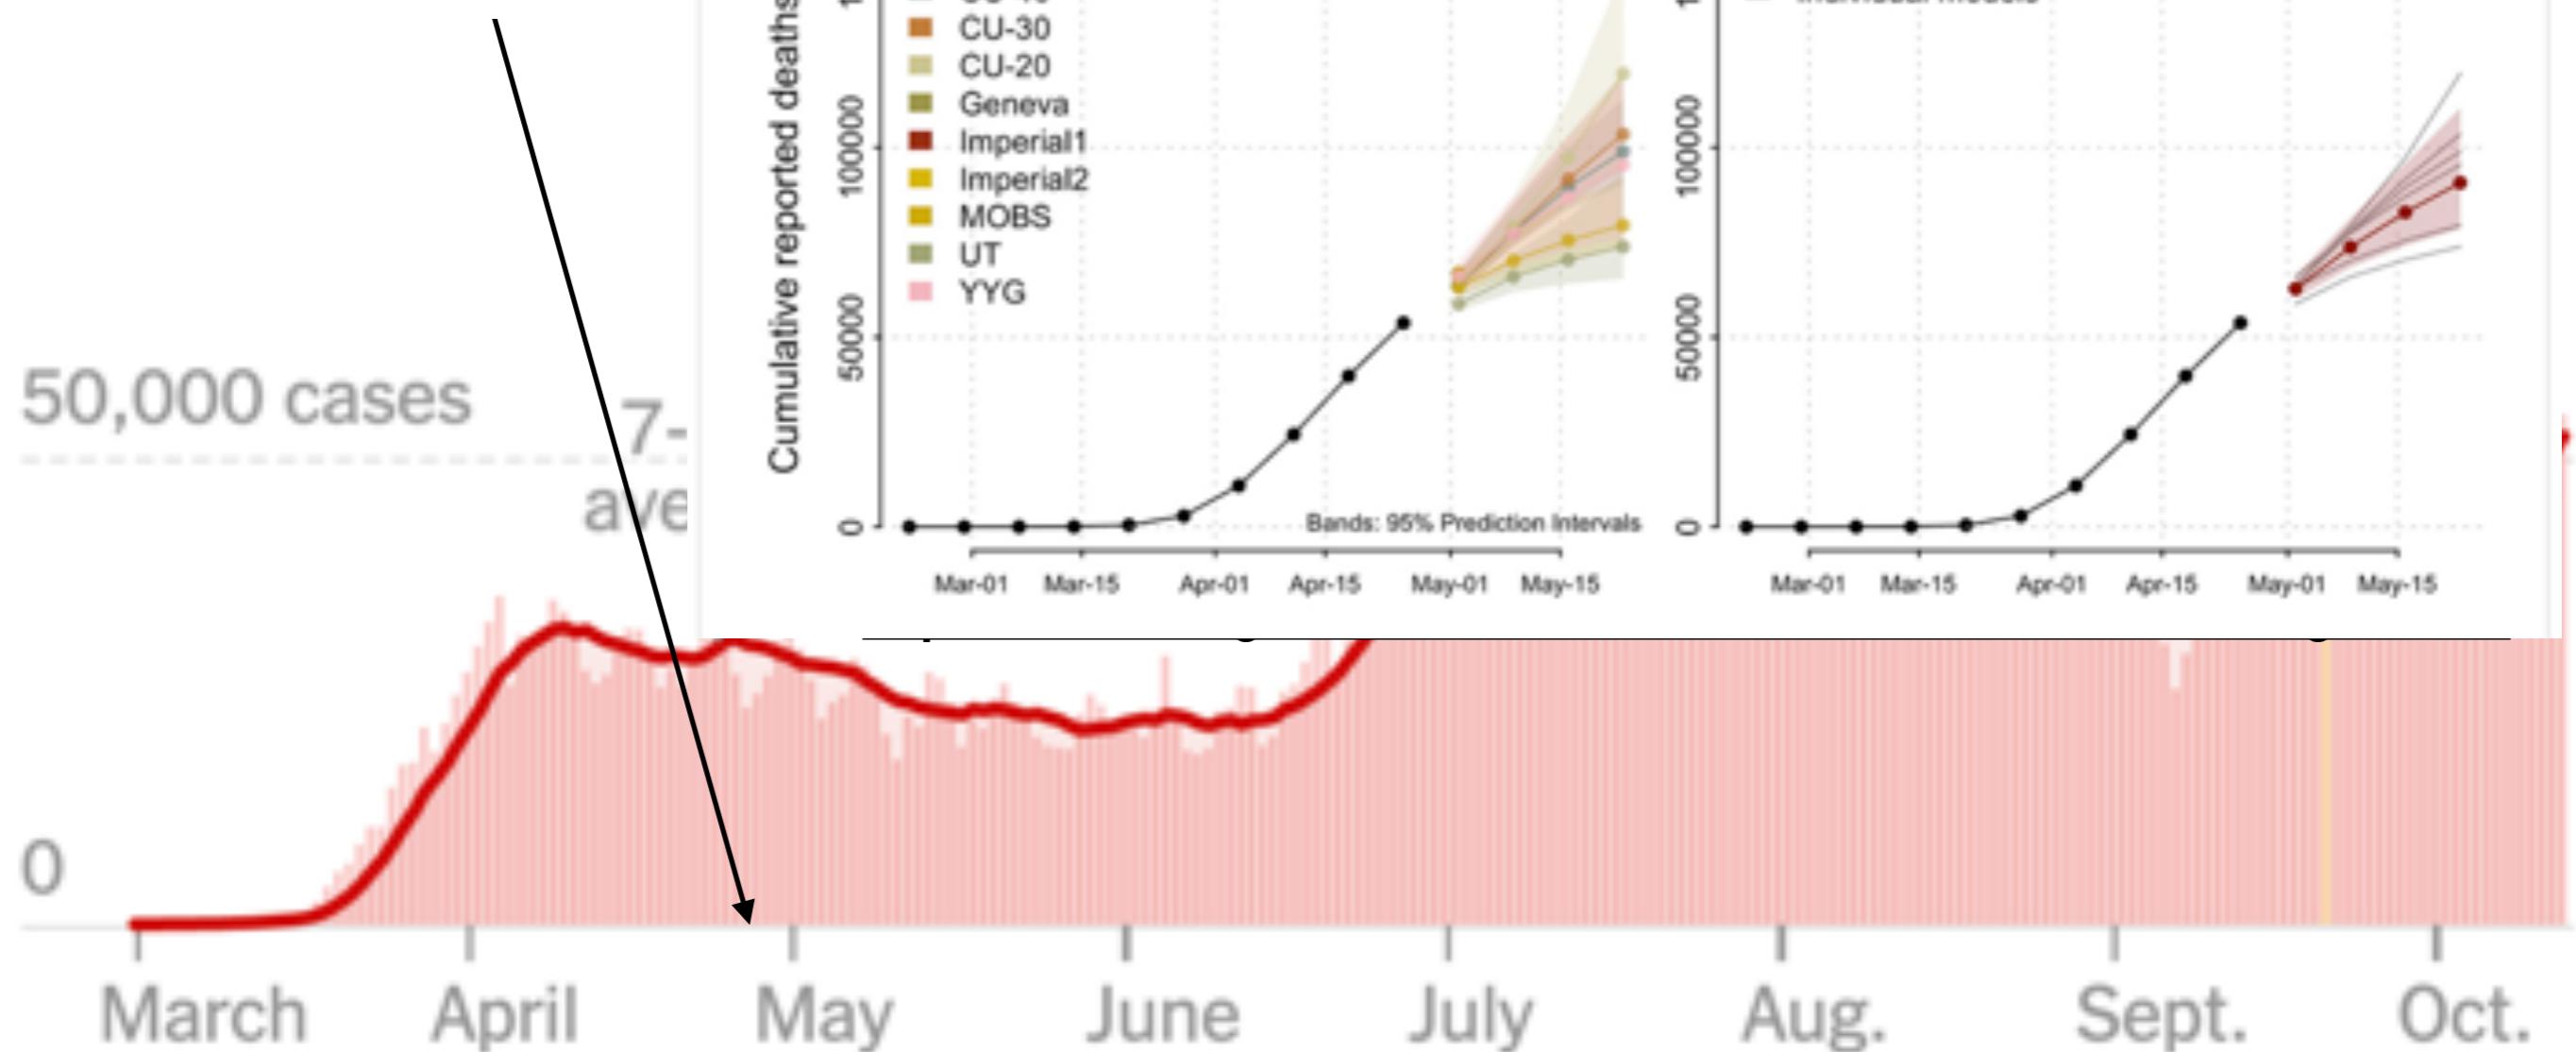


image credit: NY Times

COVID-19 Forecast Hub

Timeline

May 1, 2020

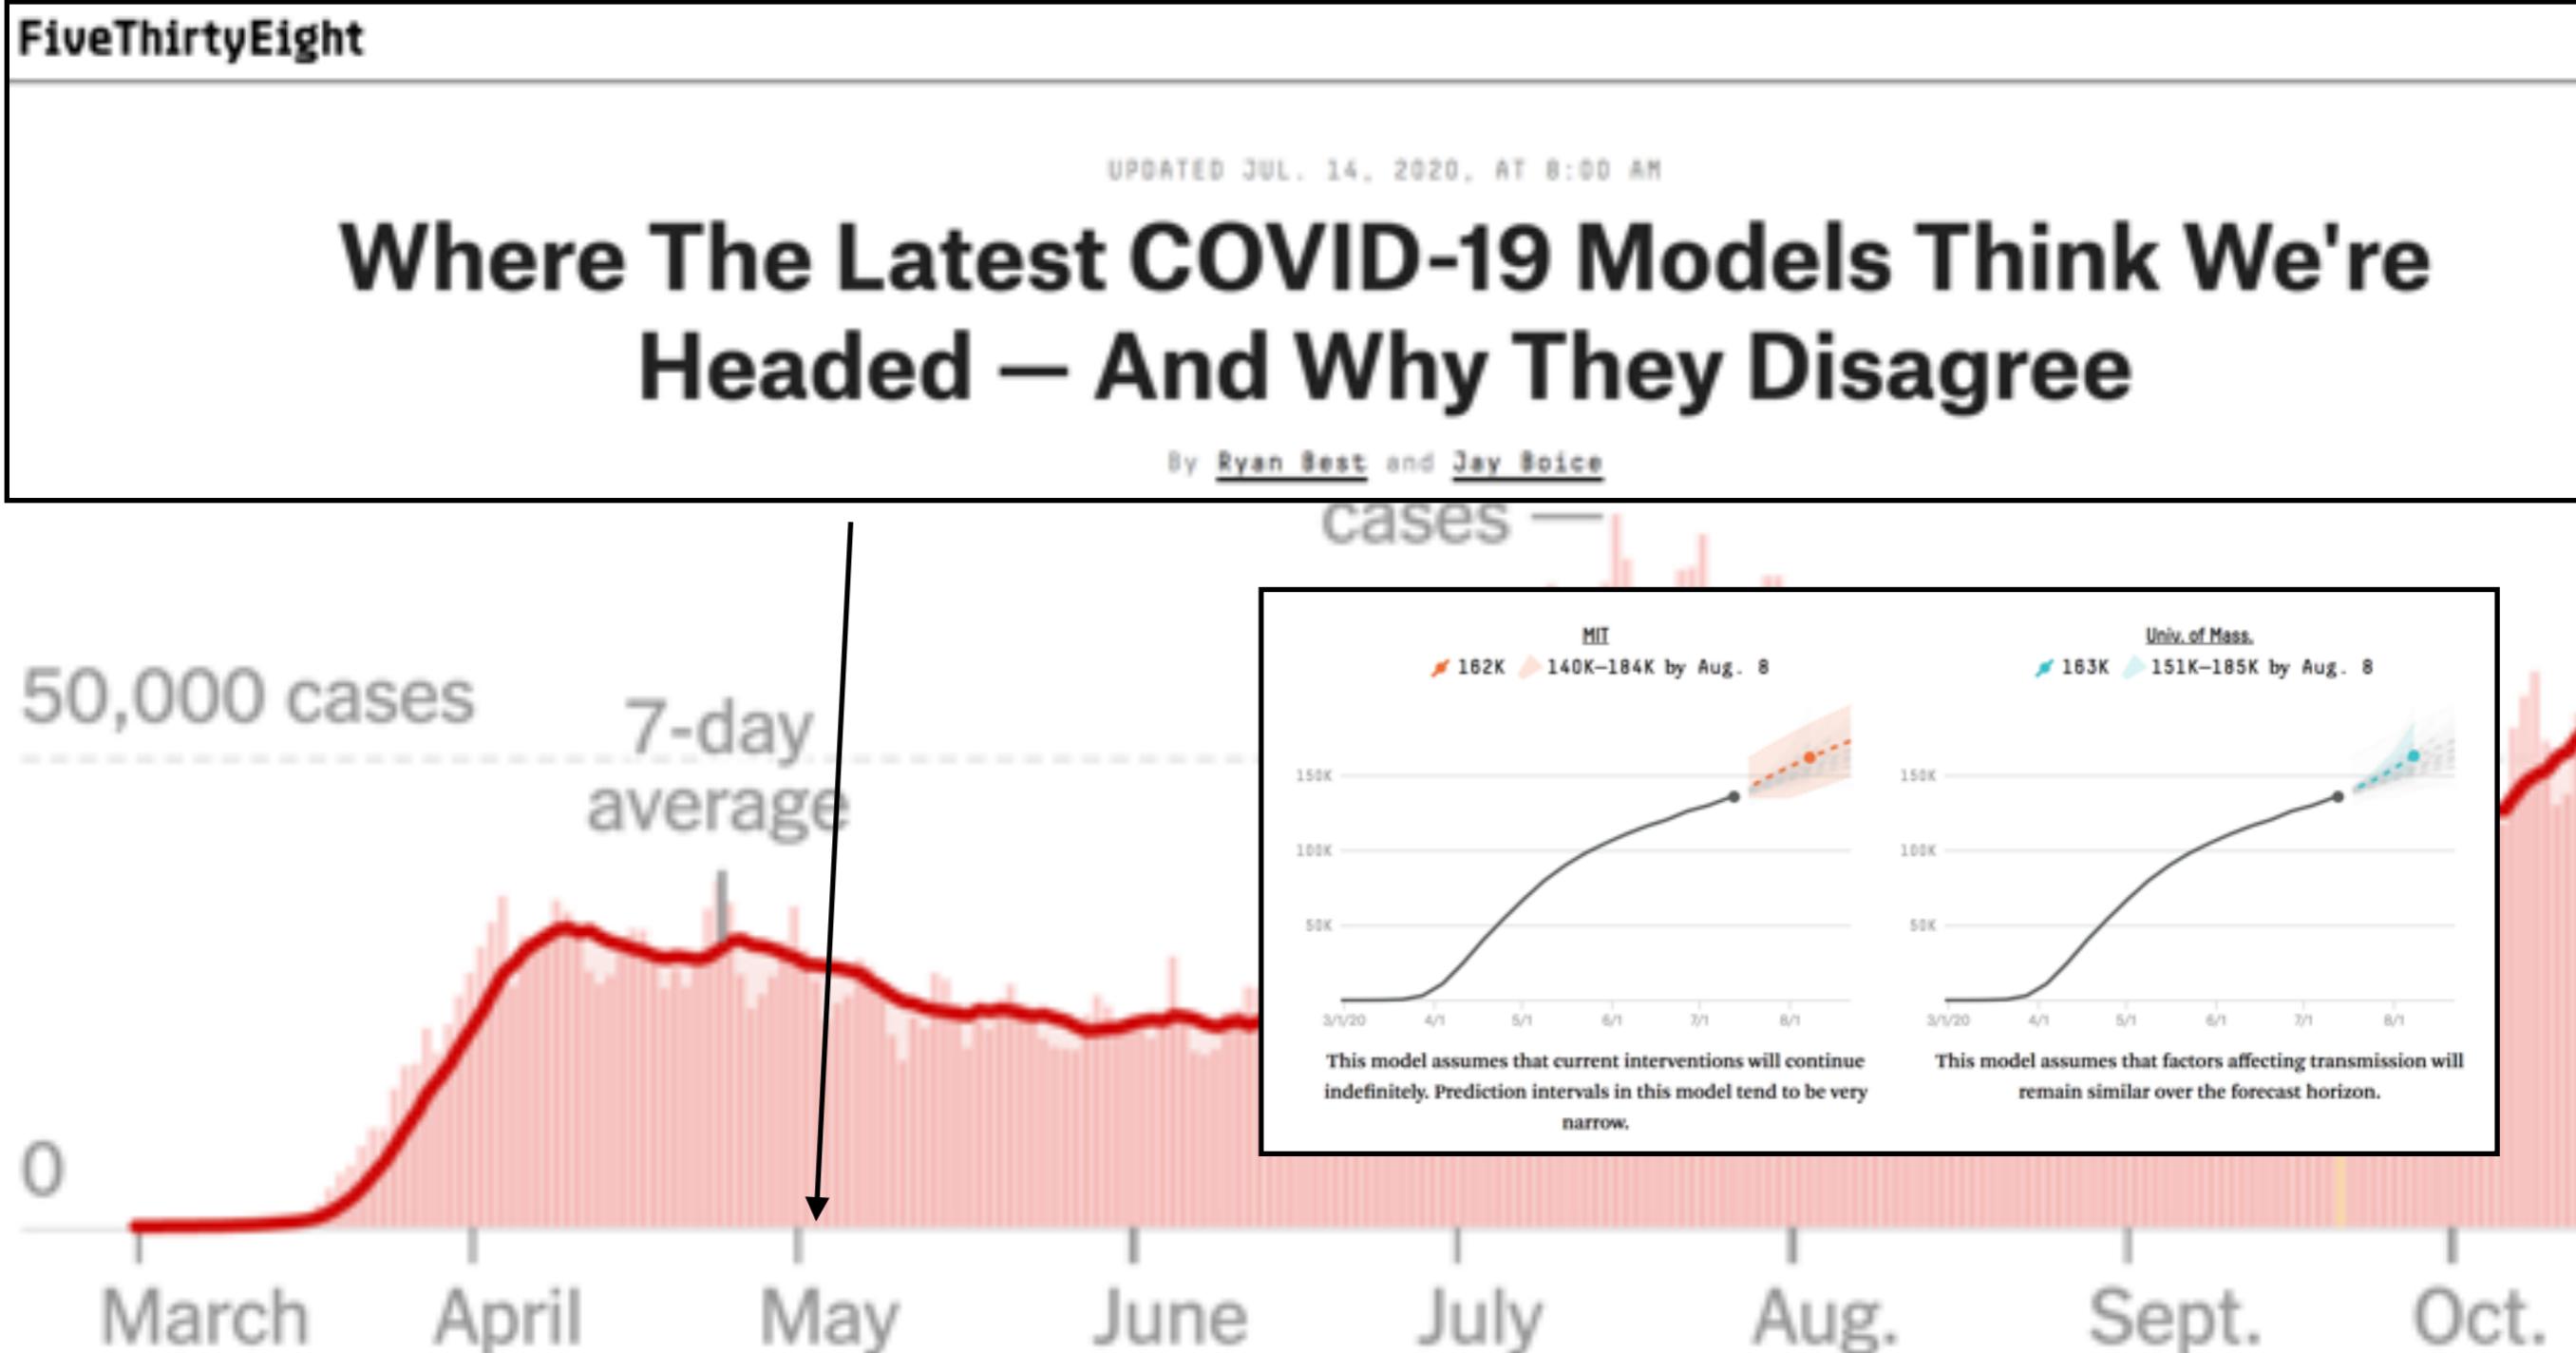


image credit: NY Times

COVID-19

May 15, 2020:
Director Redfield tweets out
results. CDC forecasting website
viewed by >1m visitors.
11 teams.

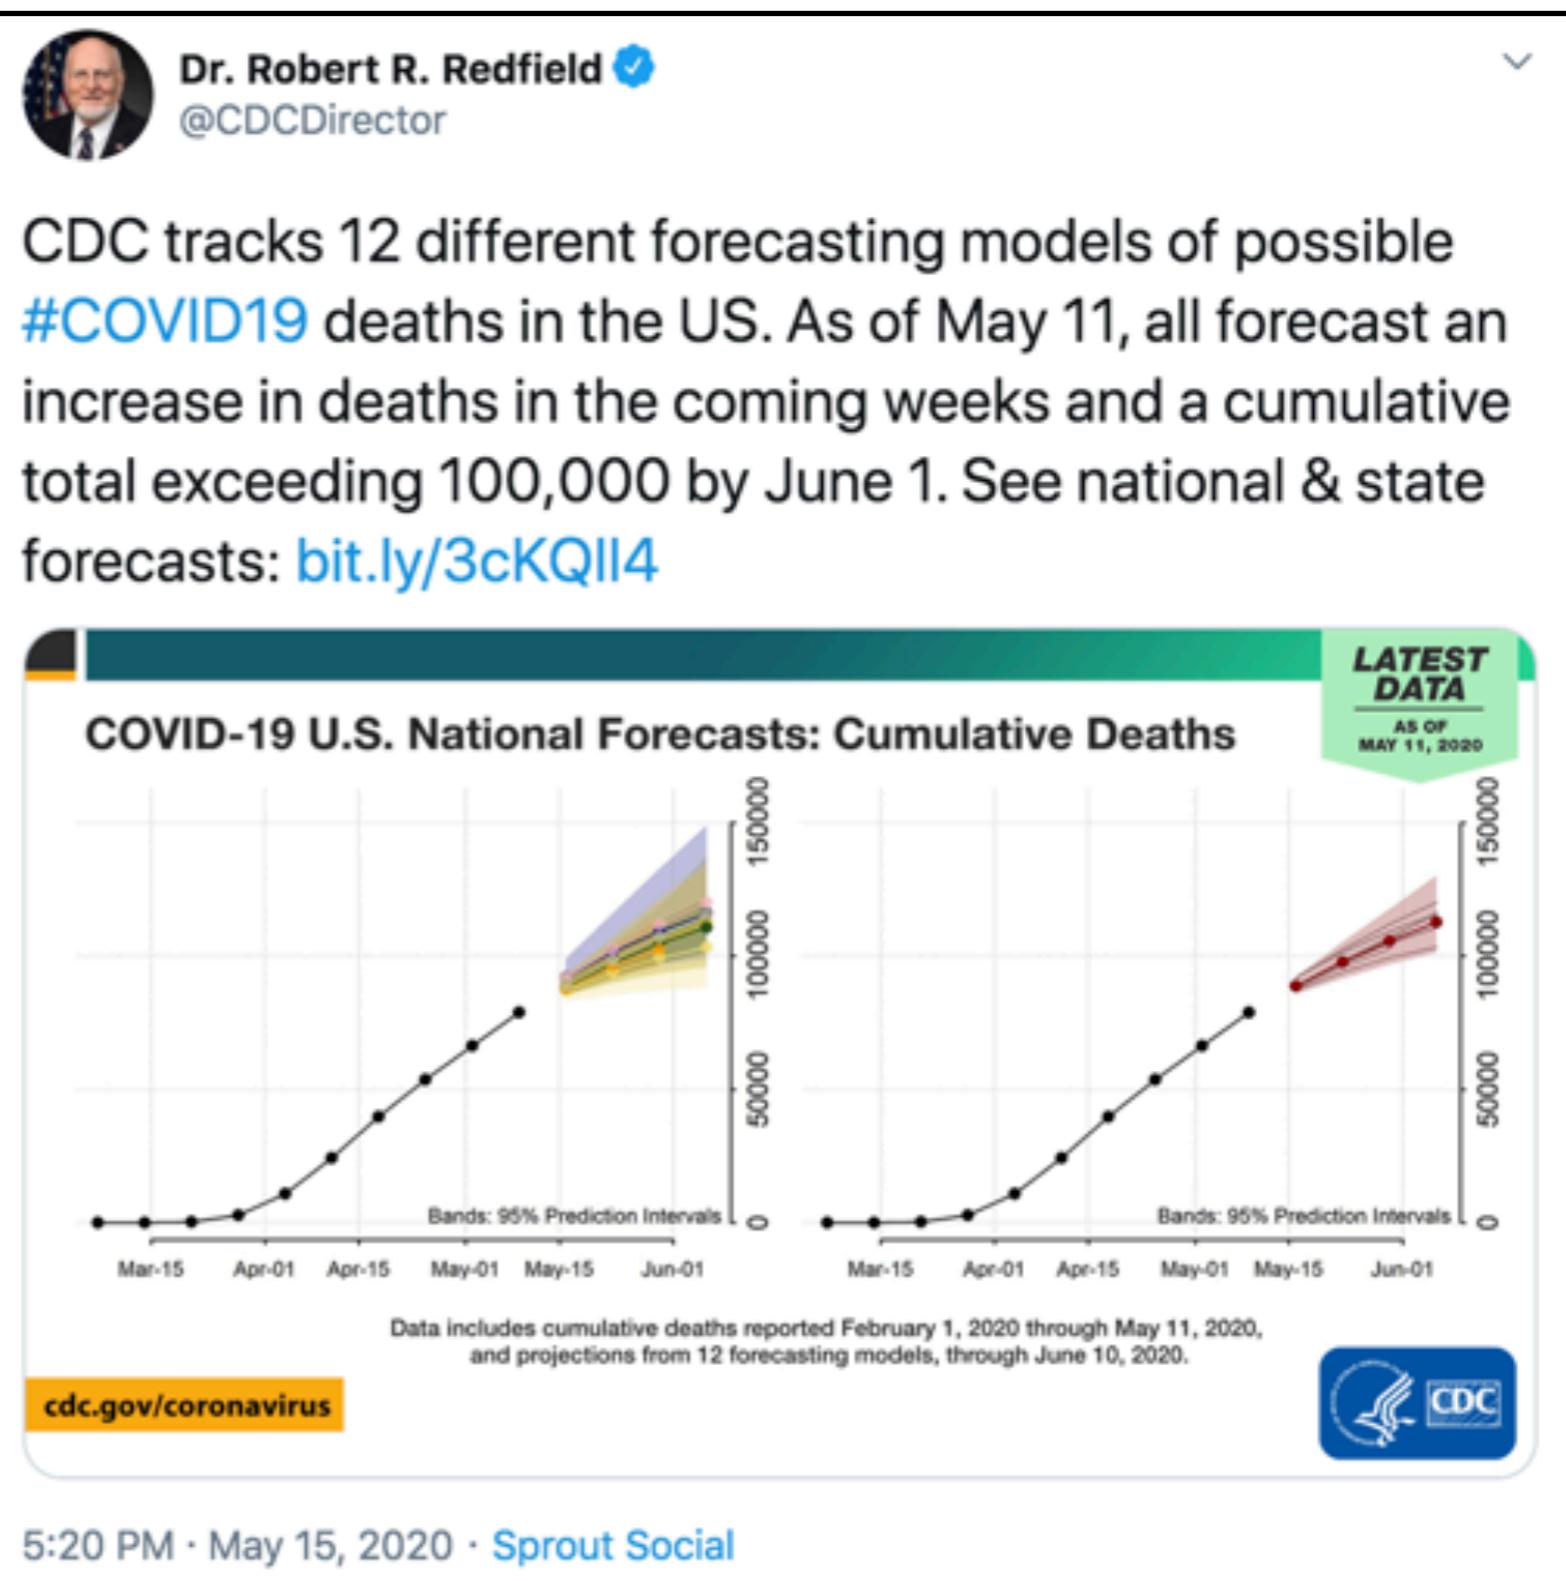
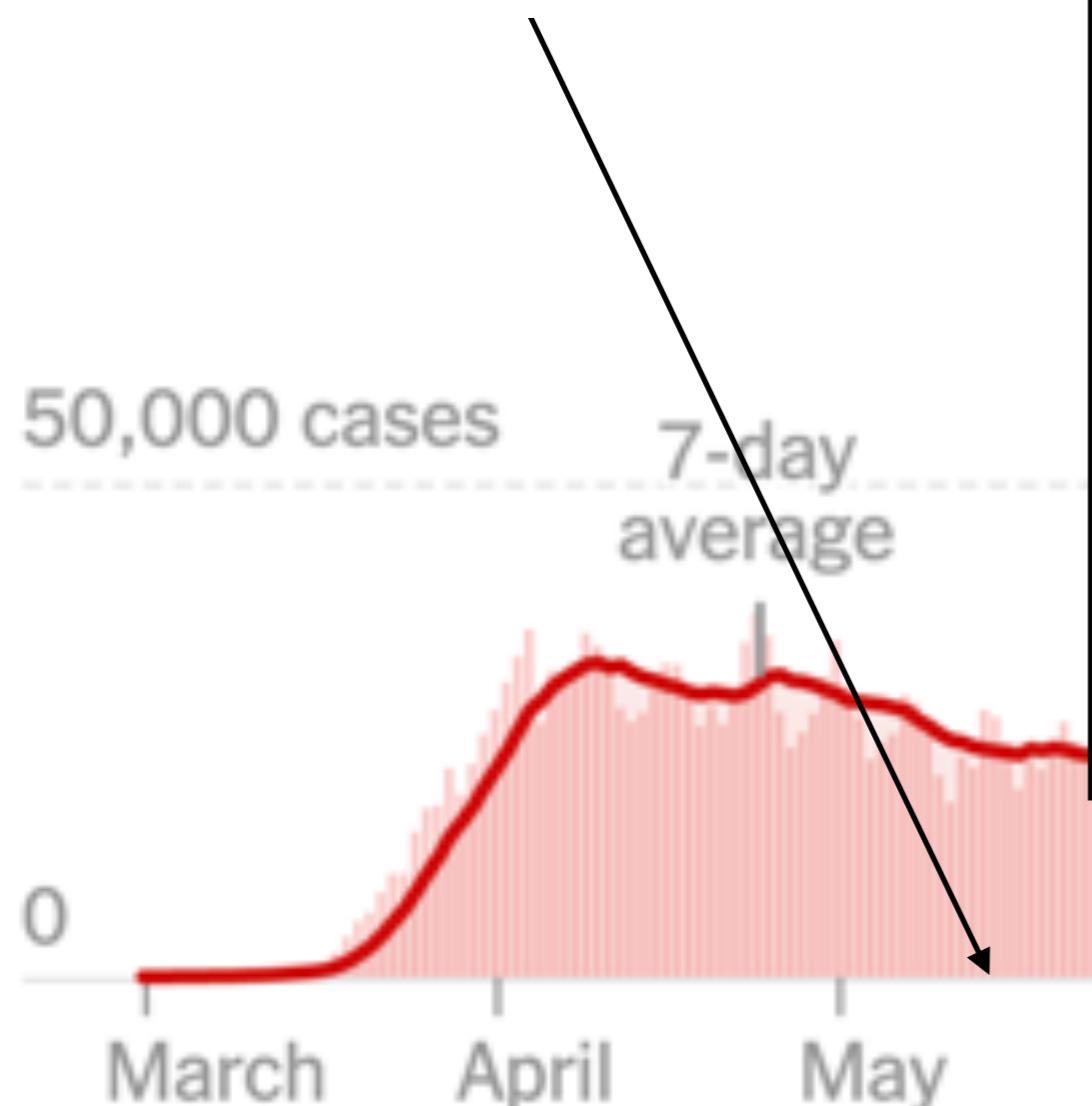


image credit: NY Times

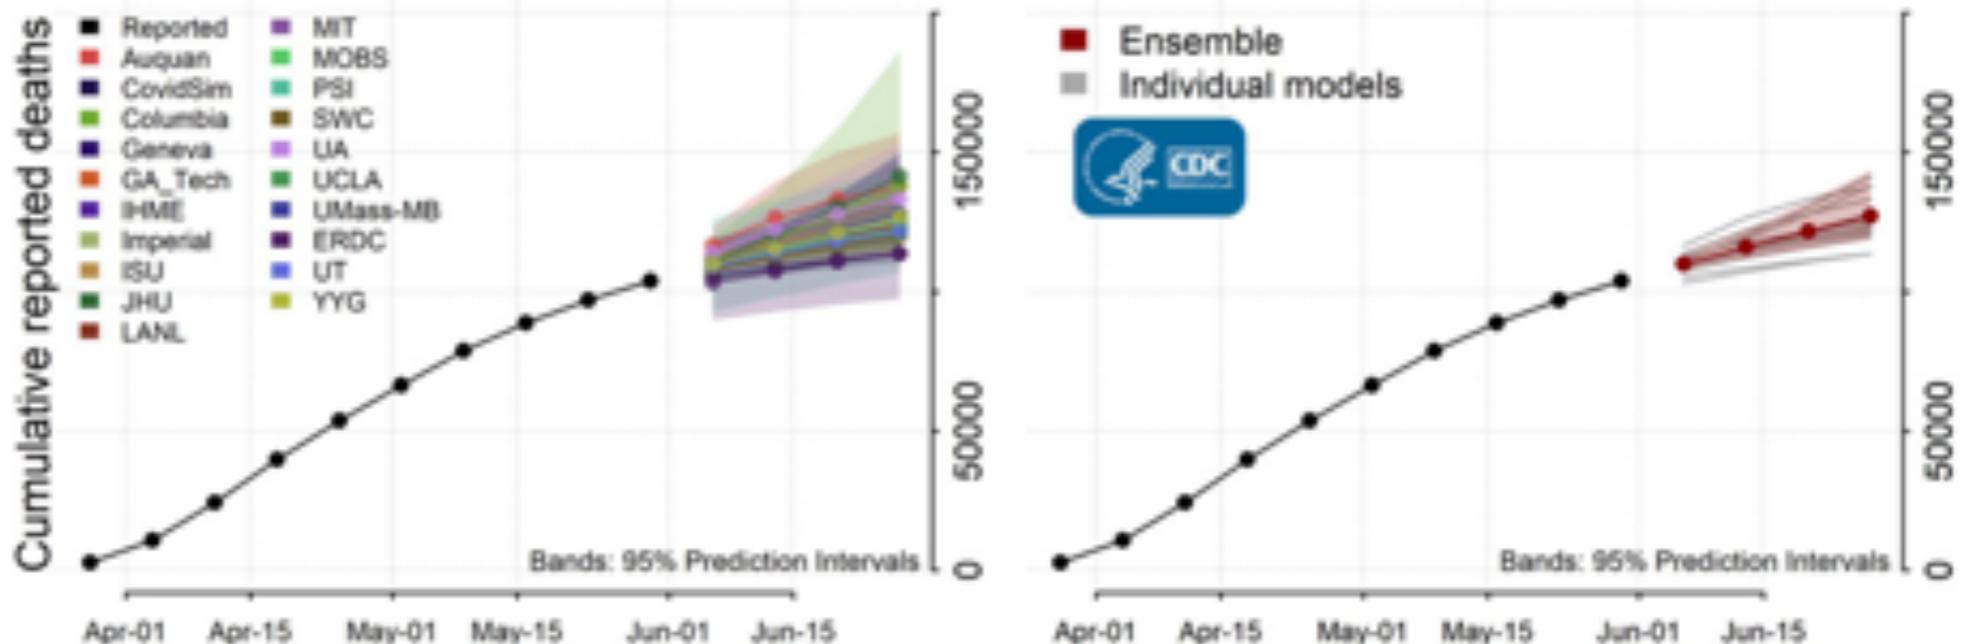
COVID-19 Forecast Hub

June 4, 2020

National Forecast

**June 4, 2020:
20 teams**

50,000 cases



average



image credit: NY Times

COVID-19 Forecast Hub Timeline

July 15, 2020:
24 teams

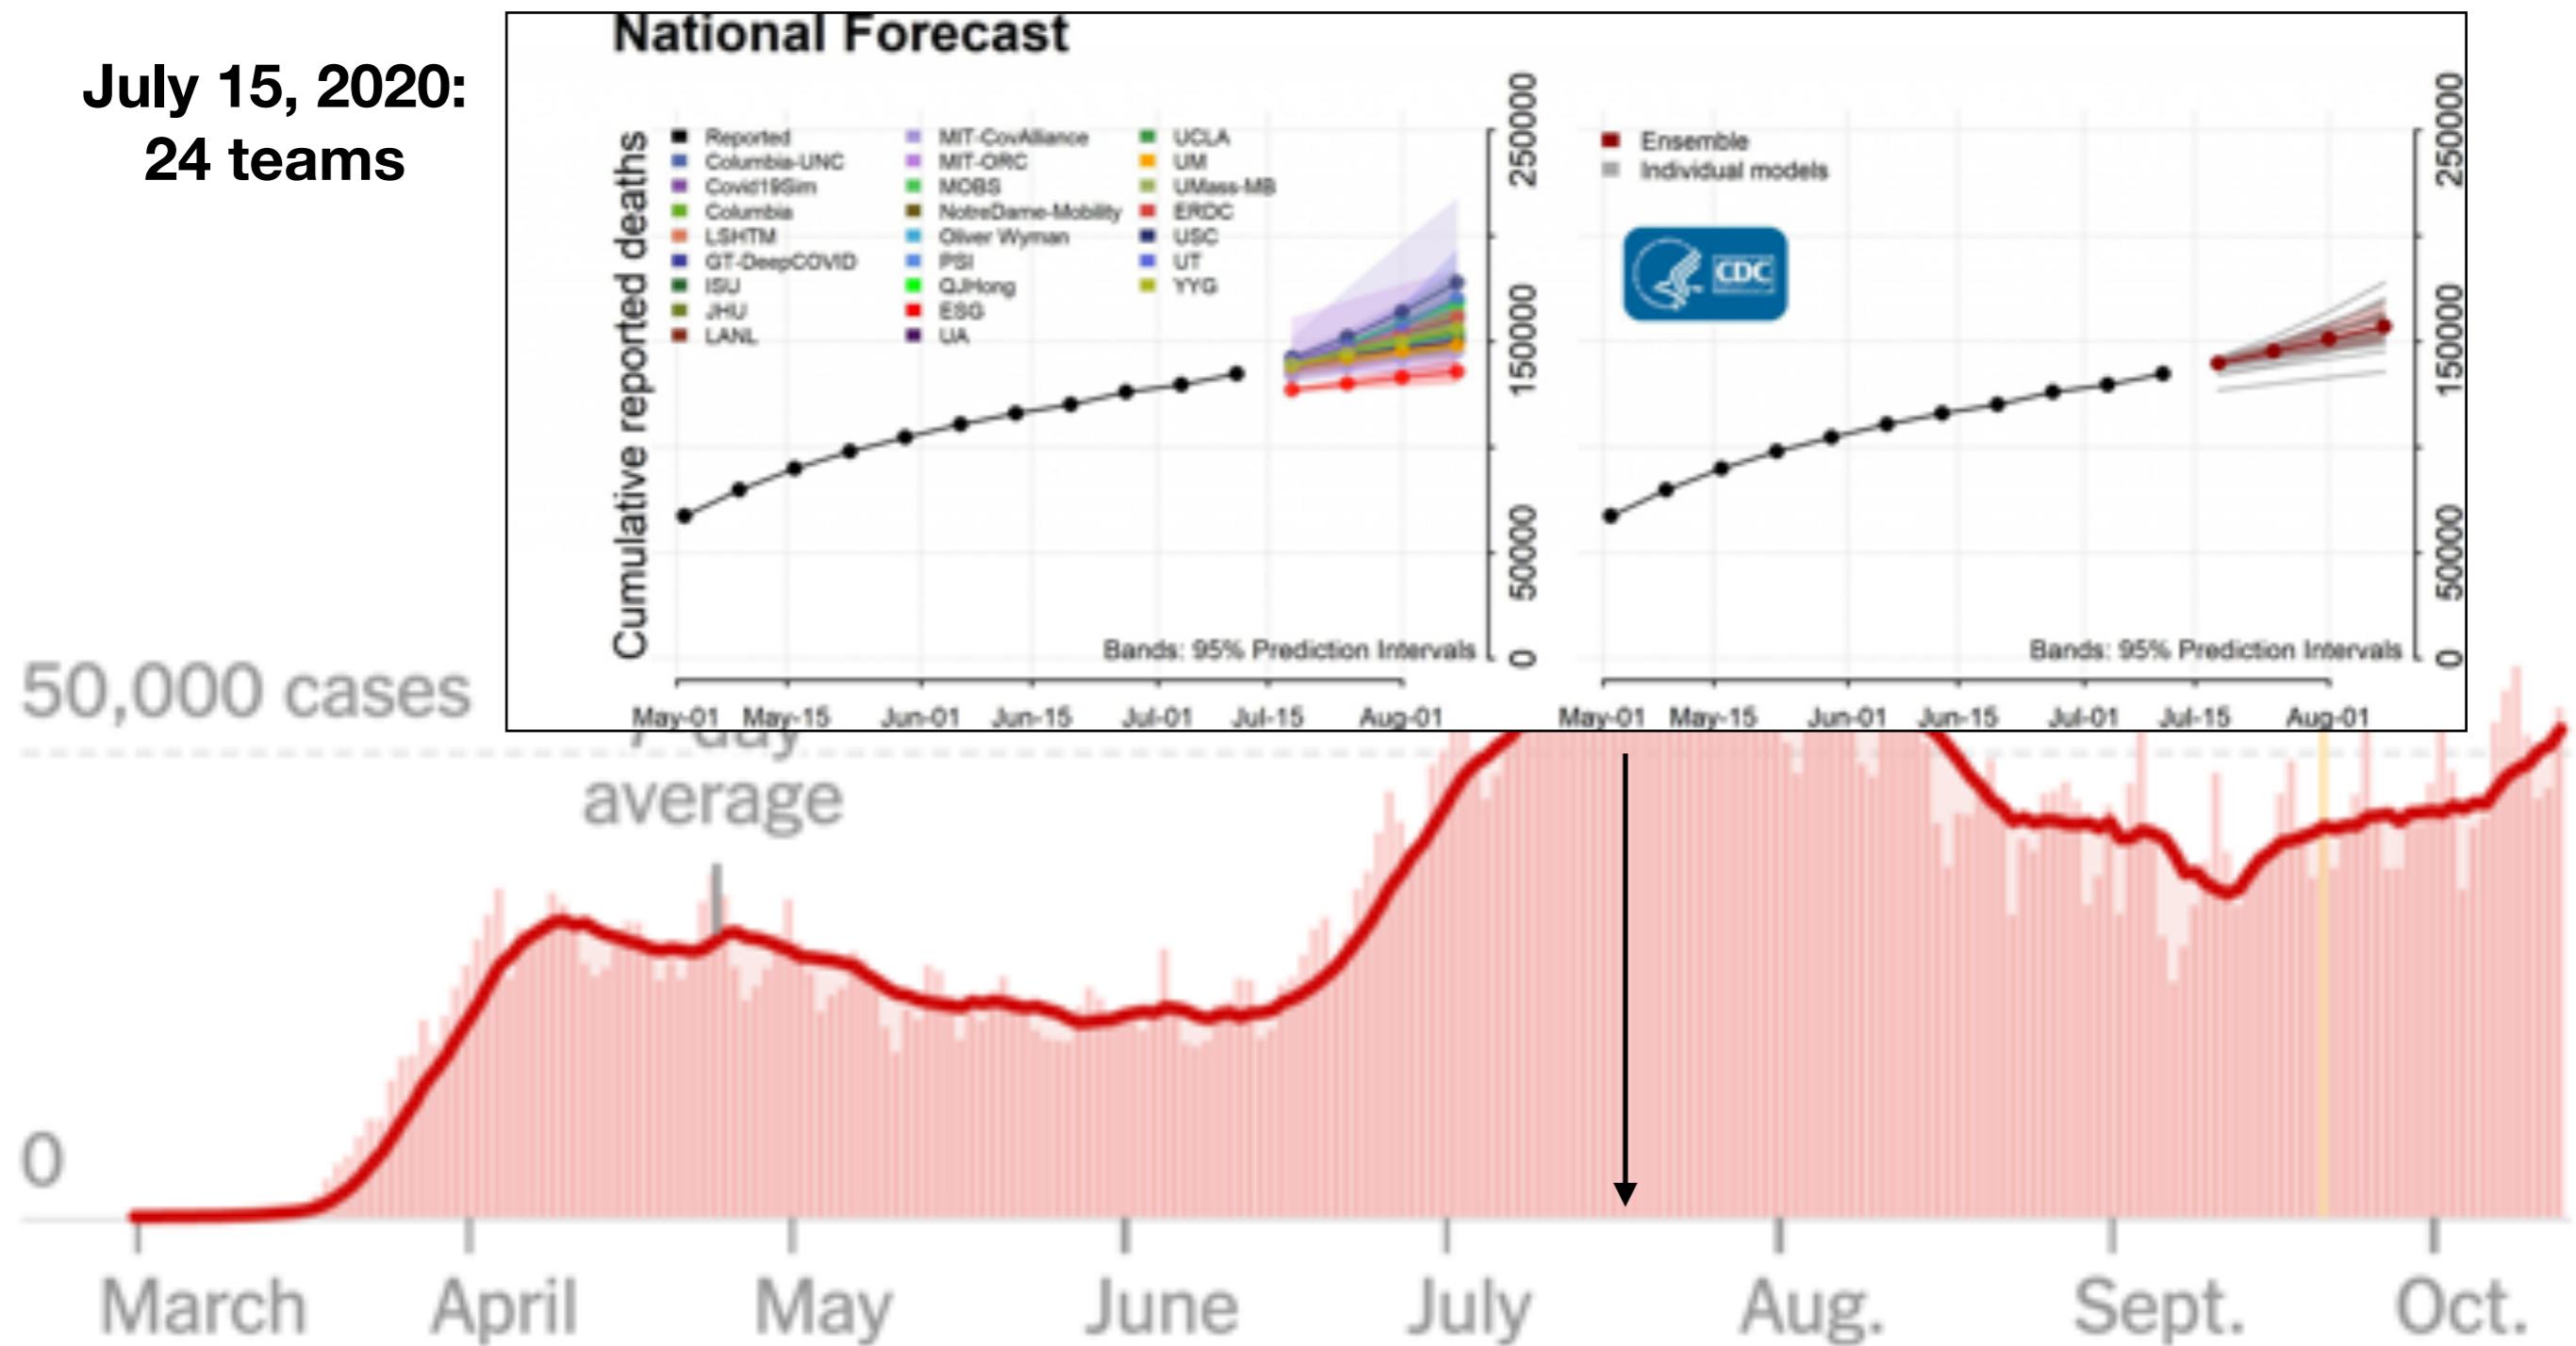
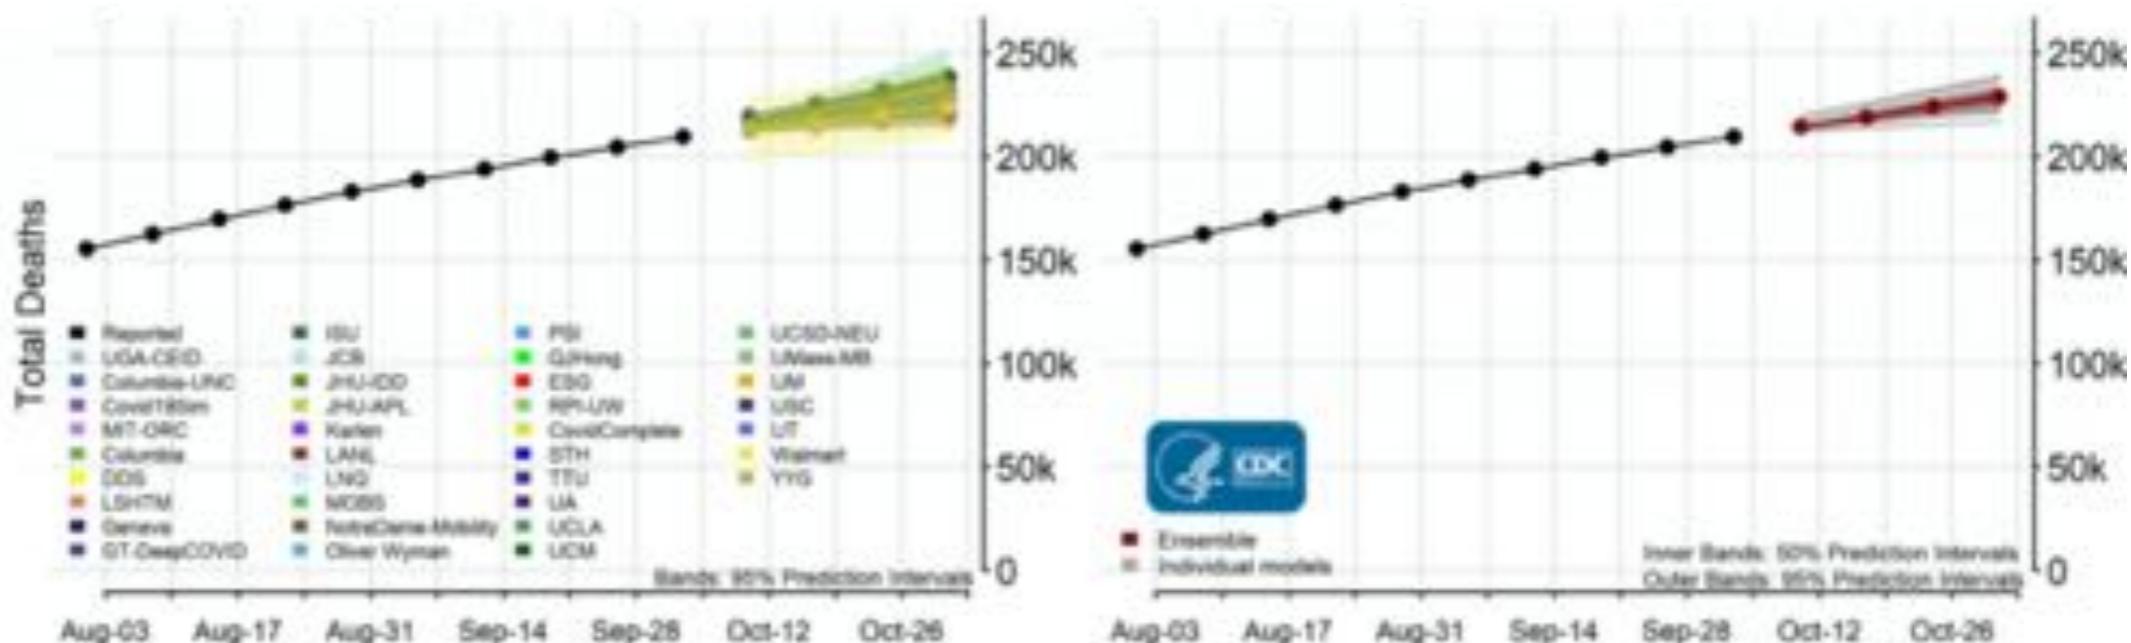


image credit: NY Times

COVID-19 Forecast Hub Timeline

Oct 15, 2020:
36 teams



50,000 cases

7-day average

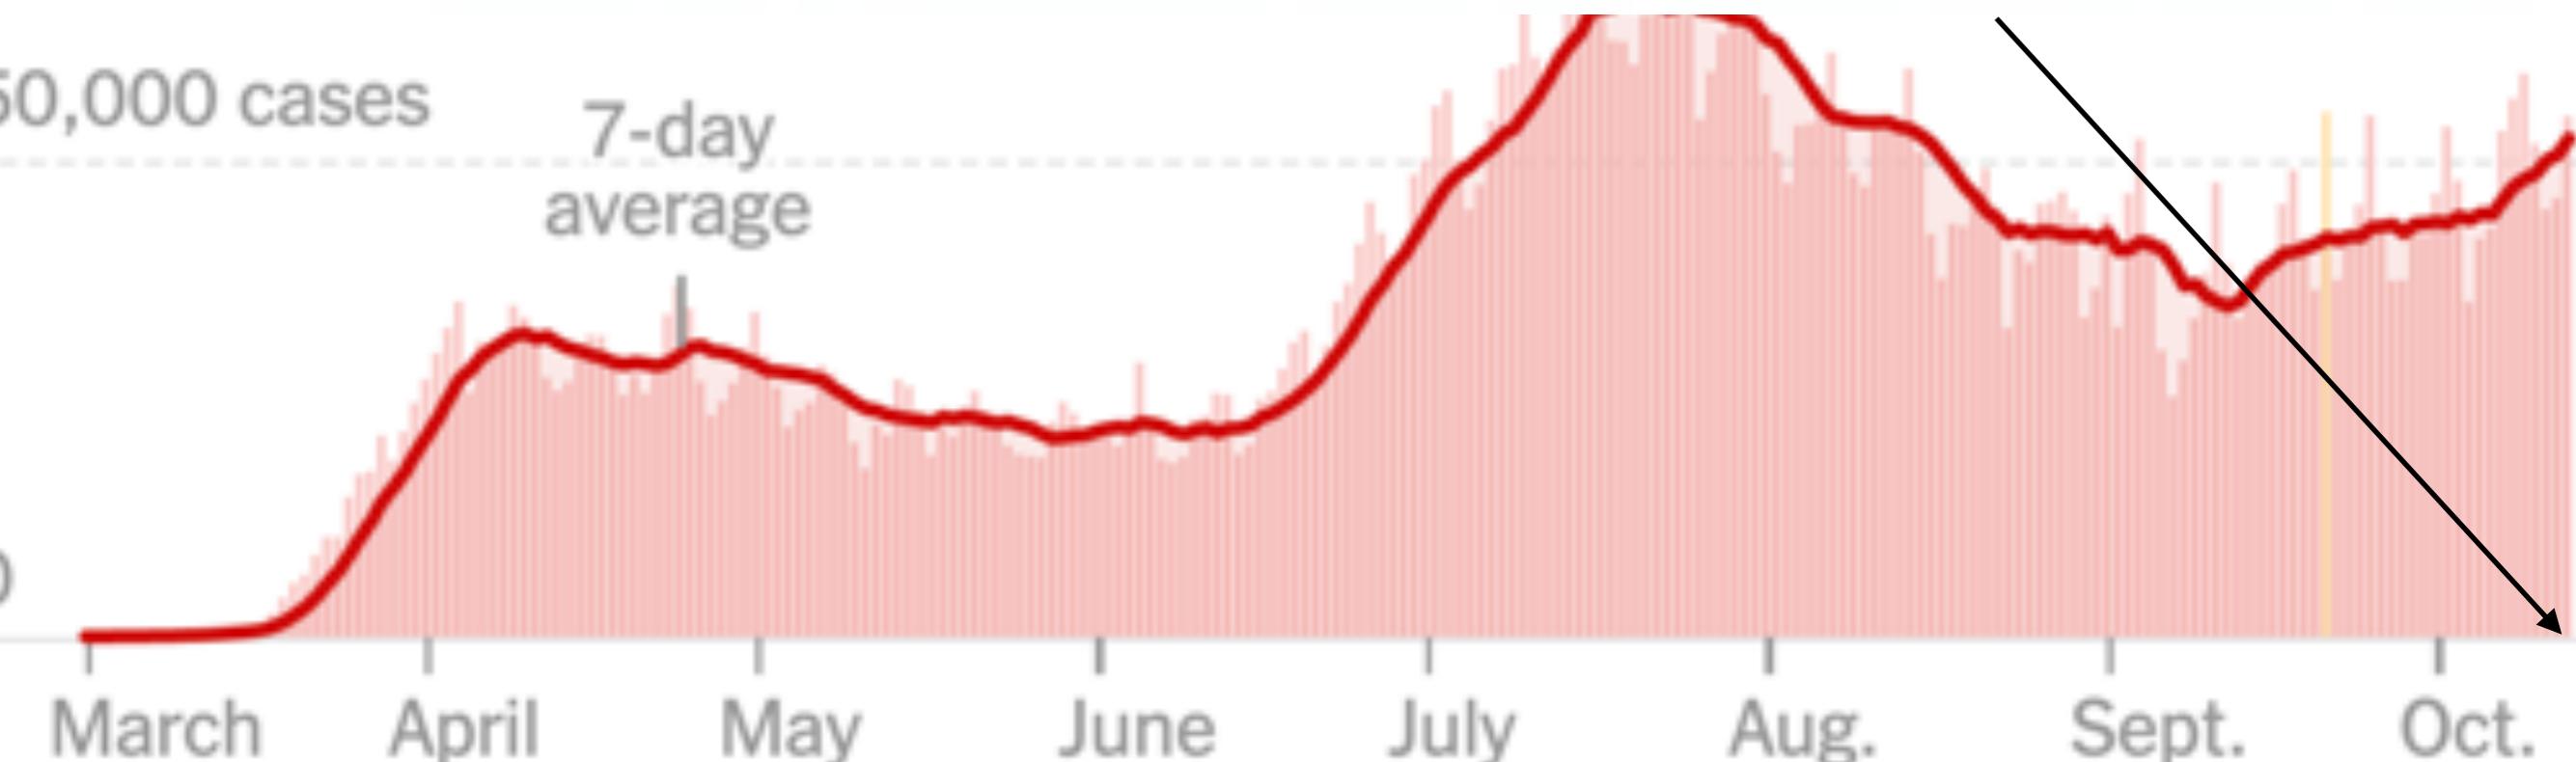


image credit: NY Times

How did we get here?

What are we doing?

How accurate are these models?



COVID-19
ForecastHub

How did we get here?

From a long-standing forecasting collaboration with CDC.

What are we doing?

We are building some simple ensemble models.

How accurate are these models?

A little bit more accurate than a naive baseline model, but not much.



COVID-19
ForecastHub

Forecasting Seasonal Flu

CDC FluSight challenges: U.S. national, regional, state forecasts



Target variable "weighted ILI":

The % of all outpatient visits with primary complaint of influenza-like illness (ILI), weighted by state population.



https://www.cdc.gov/flu/weekly/flowsight

Influenza (Flu)

Seasonal Influenza (Flu) > Flu Activity & Surveillance

FluSight: Flu Forecasting

Language: English (US)

Current Flu Forecasting:

As of January 13, forecasts indicate (nationally):

- Flu activity is likely to remain elevated through January.
- There's a 60% chance flu activity has peaked, but a 15% chance of a peak in late January and a 25% chance of a peak in February.

Unlike CDC's traditional influenza surveillance systems, which measure influenza activity after it has occurred, flu forecasting offers the possibility to look into the future and better plan ahead, potentially reducing the impact of flu.

SEARCH

CDC A-Z INDEX

Seasonal Influenza (Flu)

- About Flu
- Flu Season
- Prevent Flu
- Symptoms & Diagnosis
- Treatment
- Schools, Businesses & Travelers
- Flu Activity & Surveillance
- CDC's WHO Collaborating Center
- Situation Update: Summary of Weekly



These numbers are from the ensemble forecast.



Influenza (Flu)

Seasonal Influenza (Flu)

About Flu

Flu Season

Prevent Flu

Symptoms & Diagnosis

Treatment

Schools, Businesses & Travelers

Flu Activity & Surveillance

CDC's WHO Collaborating Center

Situation Update: Summary of Weekly

Seasonal Influenza (Flu) > Flu Activity & Surveillance

FluSight: Flu Forecasting



Unlike CDC's traditional influenza surveillance system, which only tracks influenza activity after it has occurred, flu forecasting provides the possibility to look into the future and better plan ahead, reducing the impact of flu.

Current Flu Forecasting:

As of January 13, forecasts indicate (nationally):

- Flu activity is likely to remain elevated through January.
- There's a 60% chance flu activity has peaked, but a 15% chance of a peak in late January and a 25% chance of a peak in February.

Article | Open Access | Published: 24 January 2019

Collaborative efforts to forecast seasonal influenza in the United States, 2015–2016

Craig J. McGowan, Matthew Biggerstaff , Michael Johansson, Karyn M. Apfeldorf, Michal Ben-Nun, Logan Brooks, Matteo Convertino, Madhav Erraguntla, David C. Farrow, John Freeze, Saurav Ghosh, Sangwon Hyun, Sasikiran Kandula, Joceline Lega, Yang Liu, Nicholas Michaud, Haruka Morita, Jarad Niemi, Naren Ramakrishnan, Evan L. Ray, Nicholas G. Reich, Pete Riley, Jeffrey Shaman, Ryan Tibshirani, Alessandro Vespignani, Qian Zhang, Carrie Reed & The Influenza Forecasting Working Group -Show fewer authors

From the abstract:

“Higher forecast skill was associated with team participation in previous influenza forecasting challenges and utilization of ensemble forecasting techniques. The mean ensemble consistently performed well and outperformed historical trend predictions.”

RESEARCH ARTICLE

Accuracy of real-time multi-model ensemble forecasts for seasonal influenza in the U.S.

Nicholas G. Reich^{1*}, Craig J. McGowan², Teresa K. Yamana³, Abhinav Tushar⁴, Evan L. Ray⁵, Dave Osthuis⁶, Sasikiran Kandula³, Logan C. Brooks⁷, Willow Crawford-Crudell⁸, Graham Casey Gibson¹, Evan Moore¹, Rebecca Silva⁹, Matthew Biggerstaff², Michael A. Johansson¹⁰, Roni Rosenfeld¹¹, Jeffrey Shaman³

"This approach creates a weighted average of predictive densities where the weight for each component is determined by maximizing overall ensemble accuracy over past seasons. In the 2017/2018 influenza season, one of the largest seasonal outbreaks in the last 15 years, this multi-model ensemble performed better on average than all individual component models and placed second overall in the CDC challenge. ... **This project shows that collaborative efforts between research teams to develop ensemble forecasting approaches can bring measurable improvements in forecast accuracy and important reductions in the variability of performance from year to year.**"

Truth Data

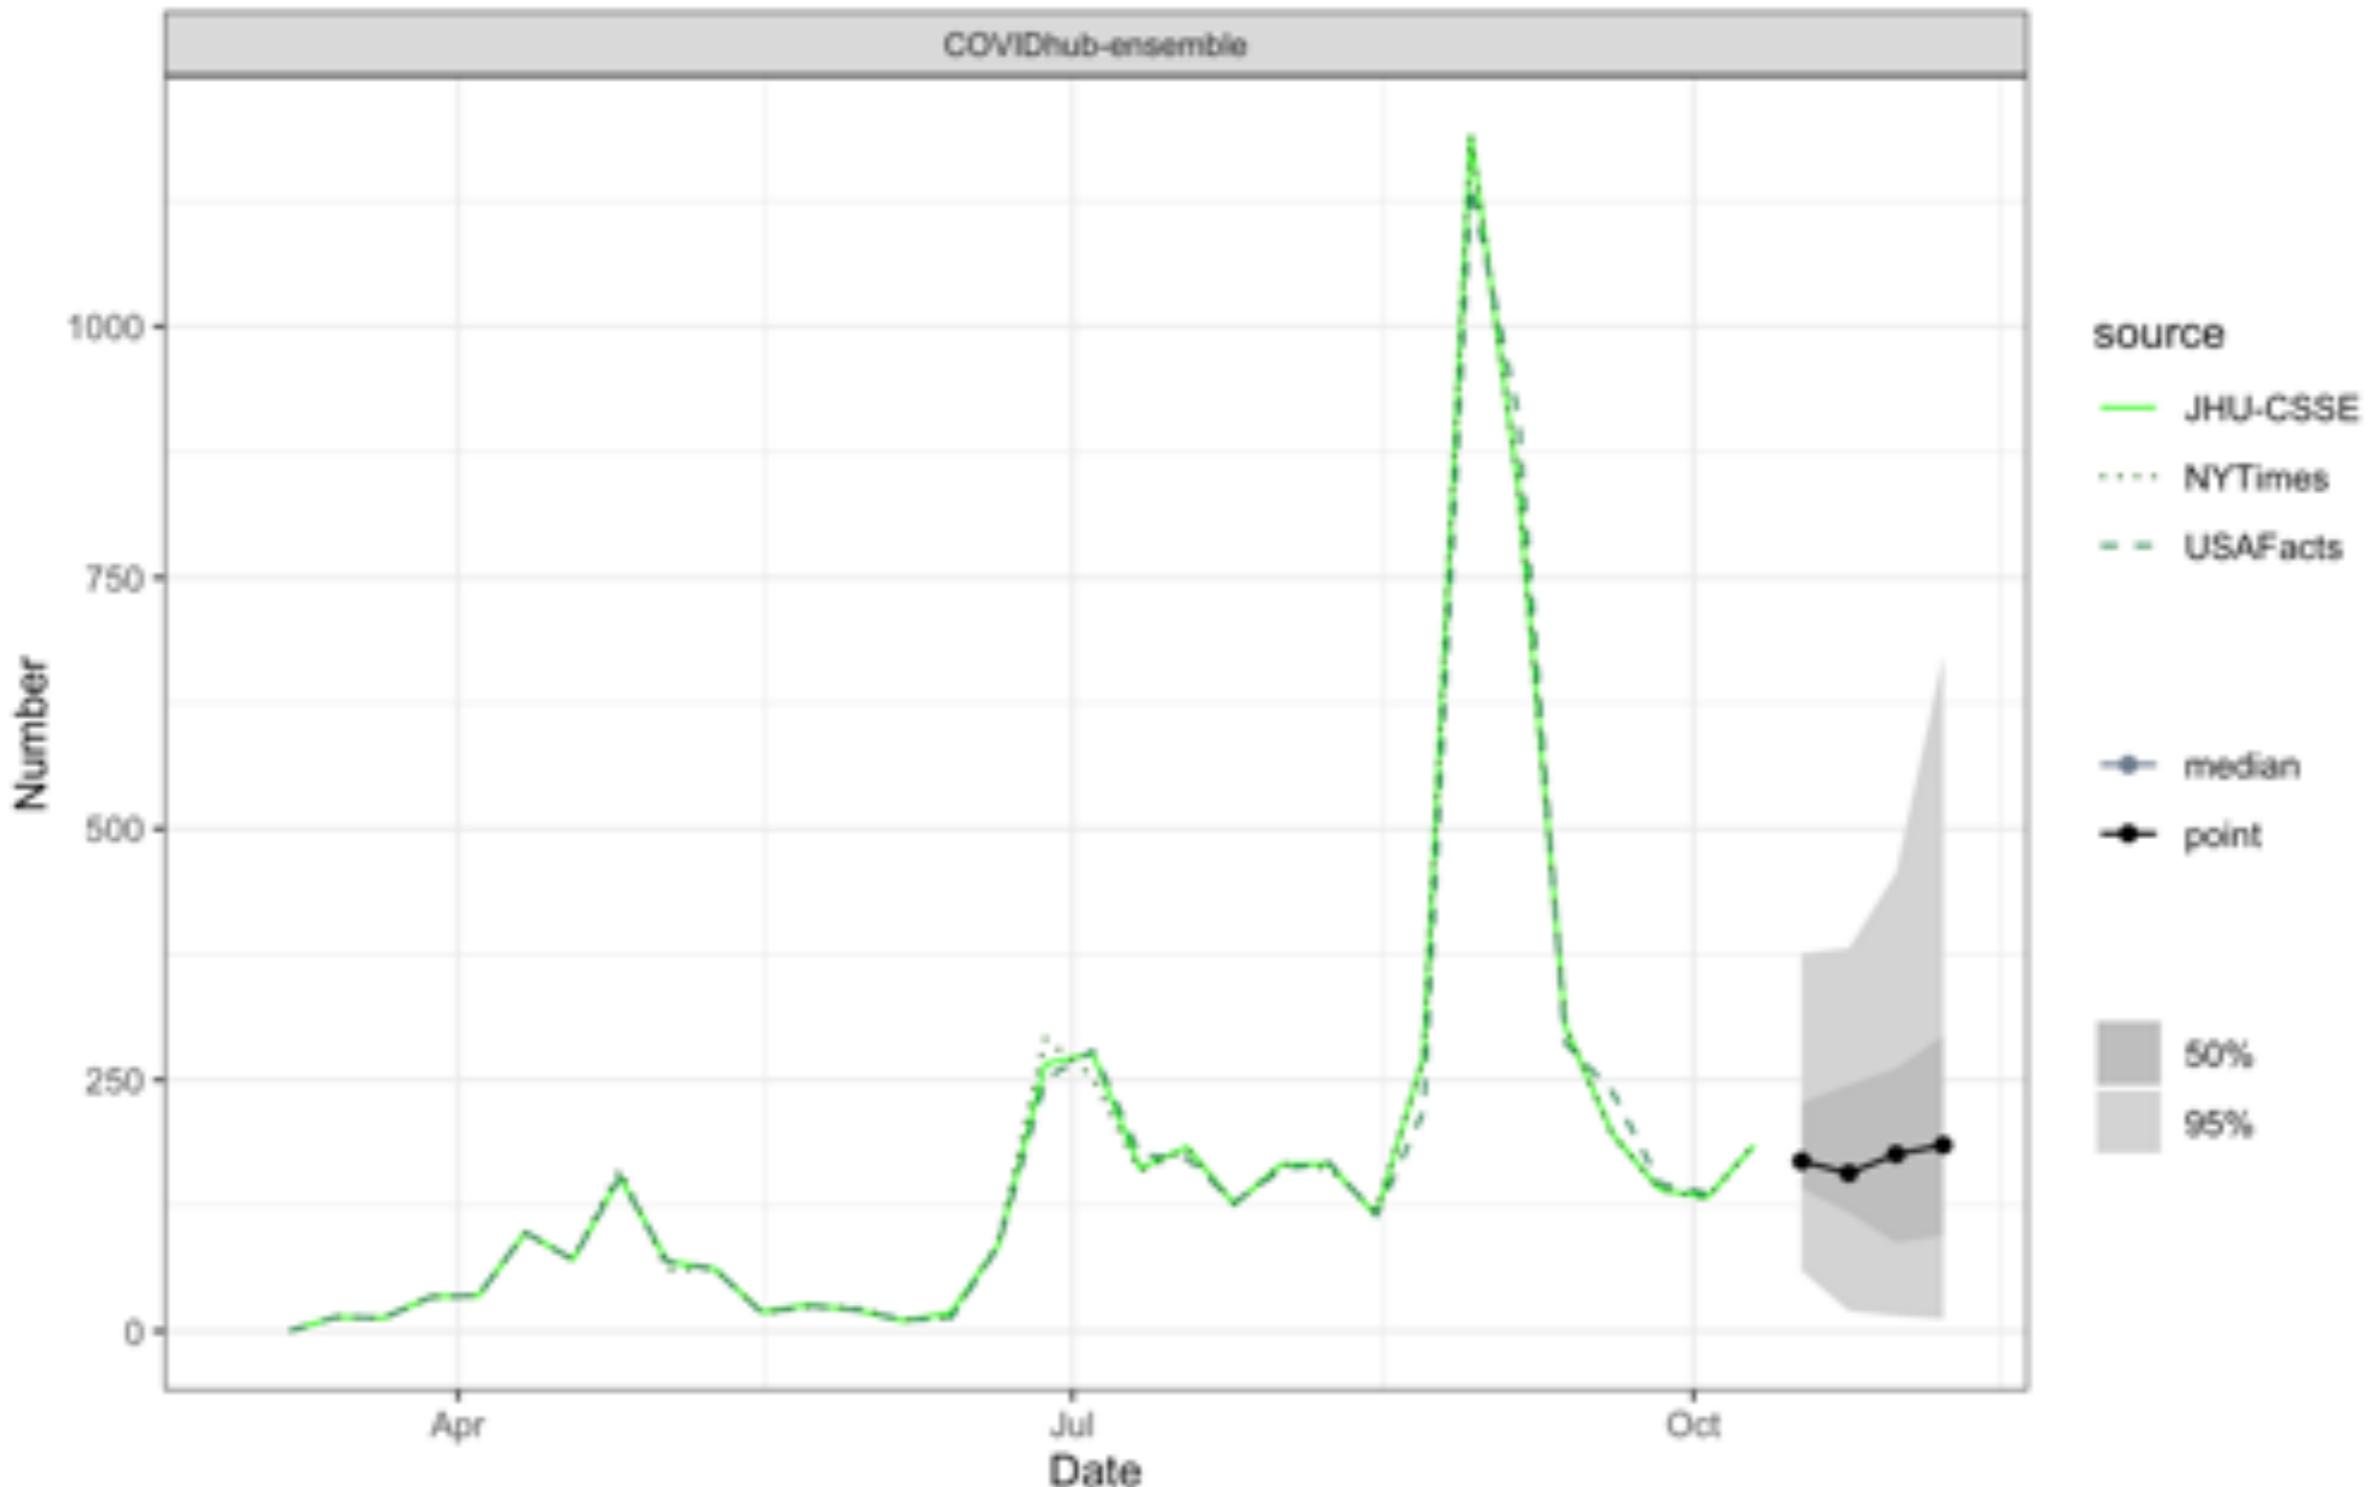


COVID-19
ForecastHub

Johnson County, Iowa

Truth discrepancy

Forecast date: 2020-10-12



ILI Backfill

CDC FluSight Network Collaborative Ensemble

WEEK 7 (2018)
HHS Region 7

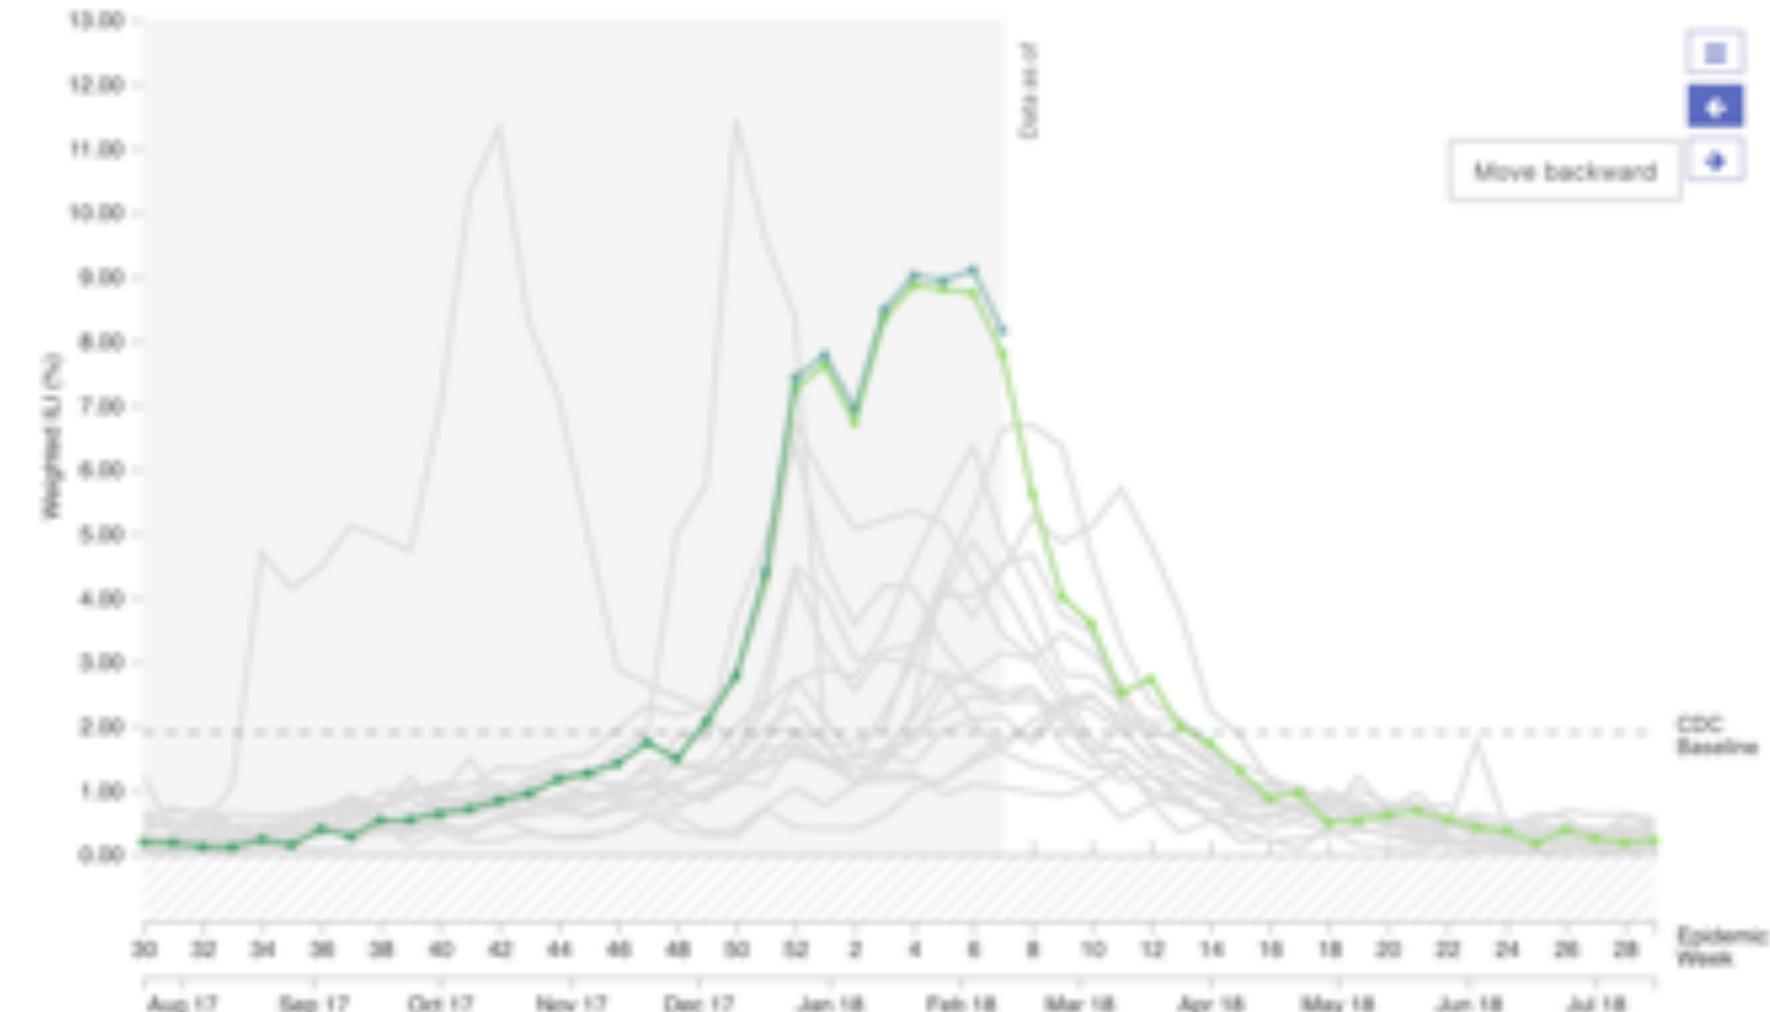
SEASON
2017-2018

Weighted ILI (%)

Absolute Relative



Time Chart Distribution Chart Scores



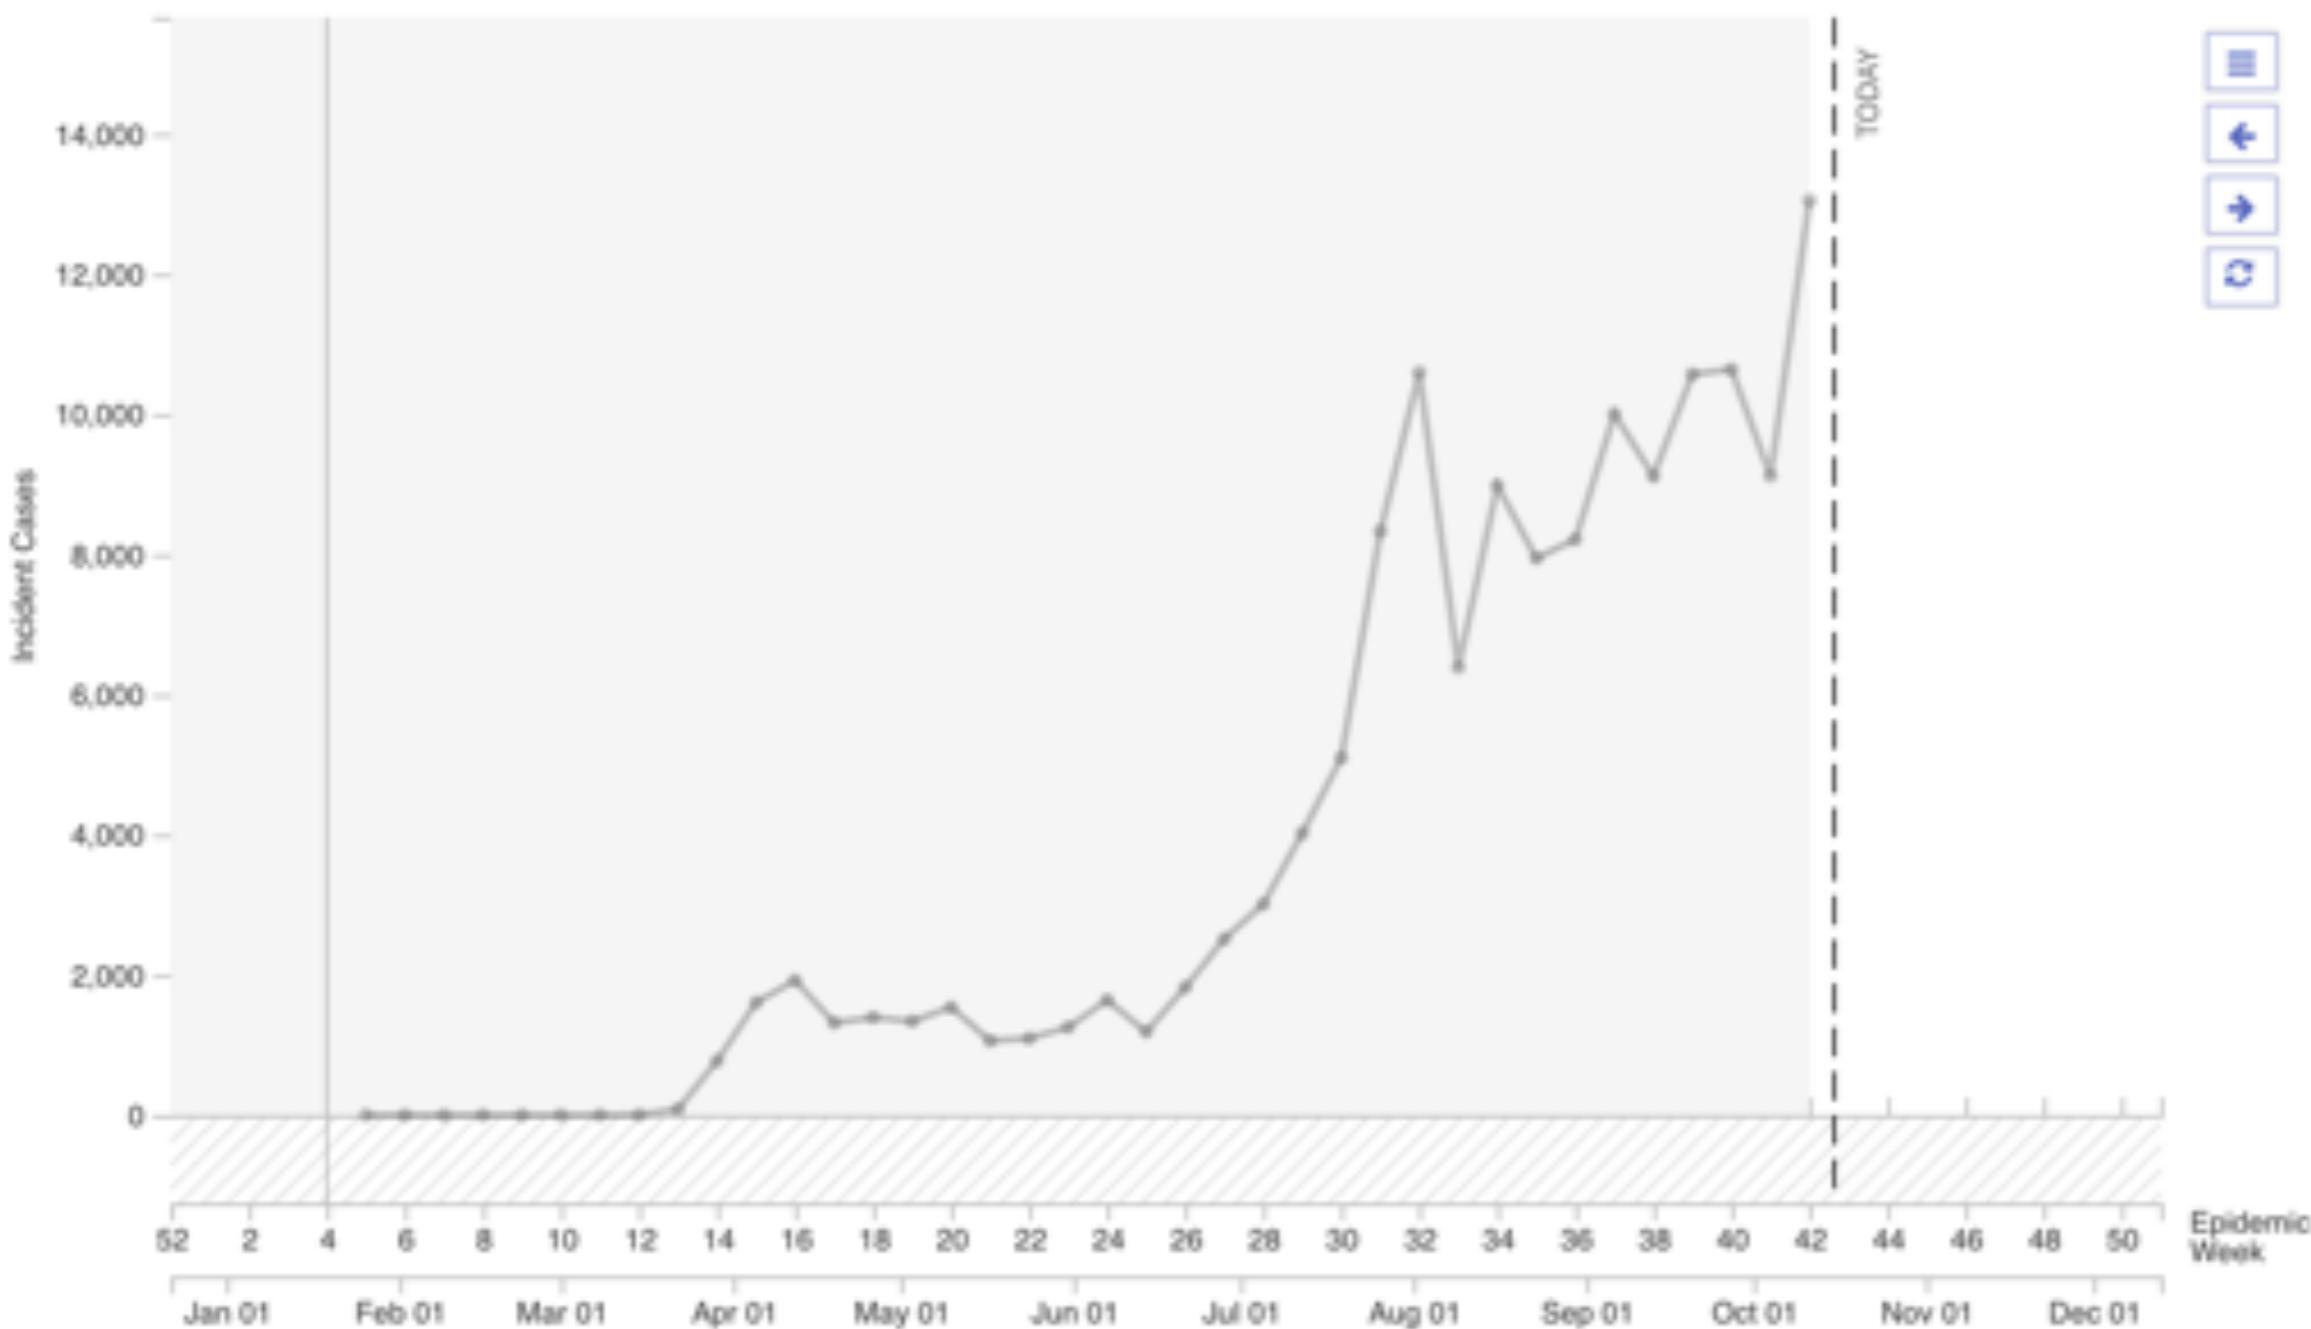
Data last updated on Mon, 16 Mar 2020 13:45:26 GMT.

Visualizations use D3, see the supported browsers here. The source is licensed MIT.

Missouri COVID-19 Backfill

github.com/CSSEGISandData/COVID-19/tree/master/csse_covid_19_data

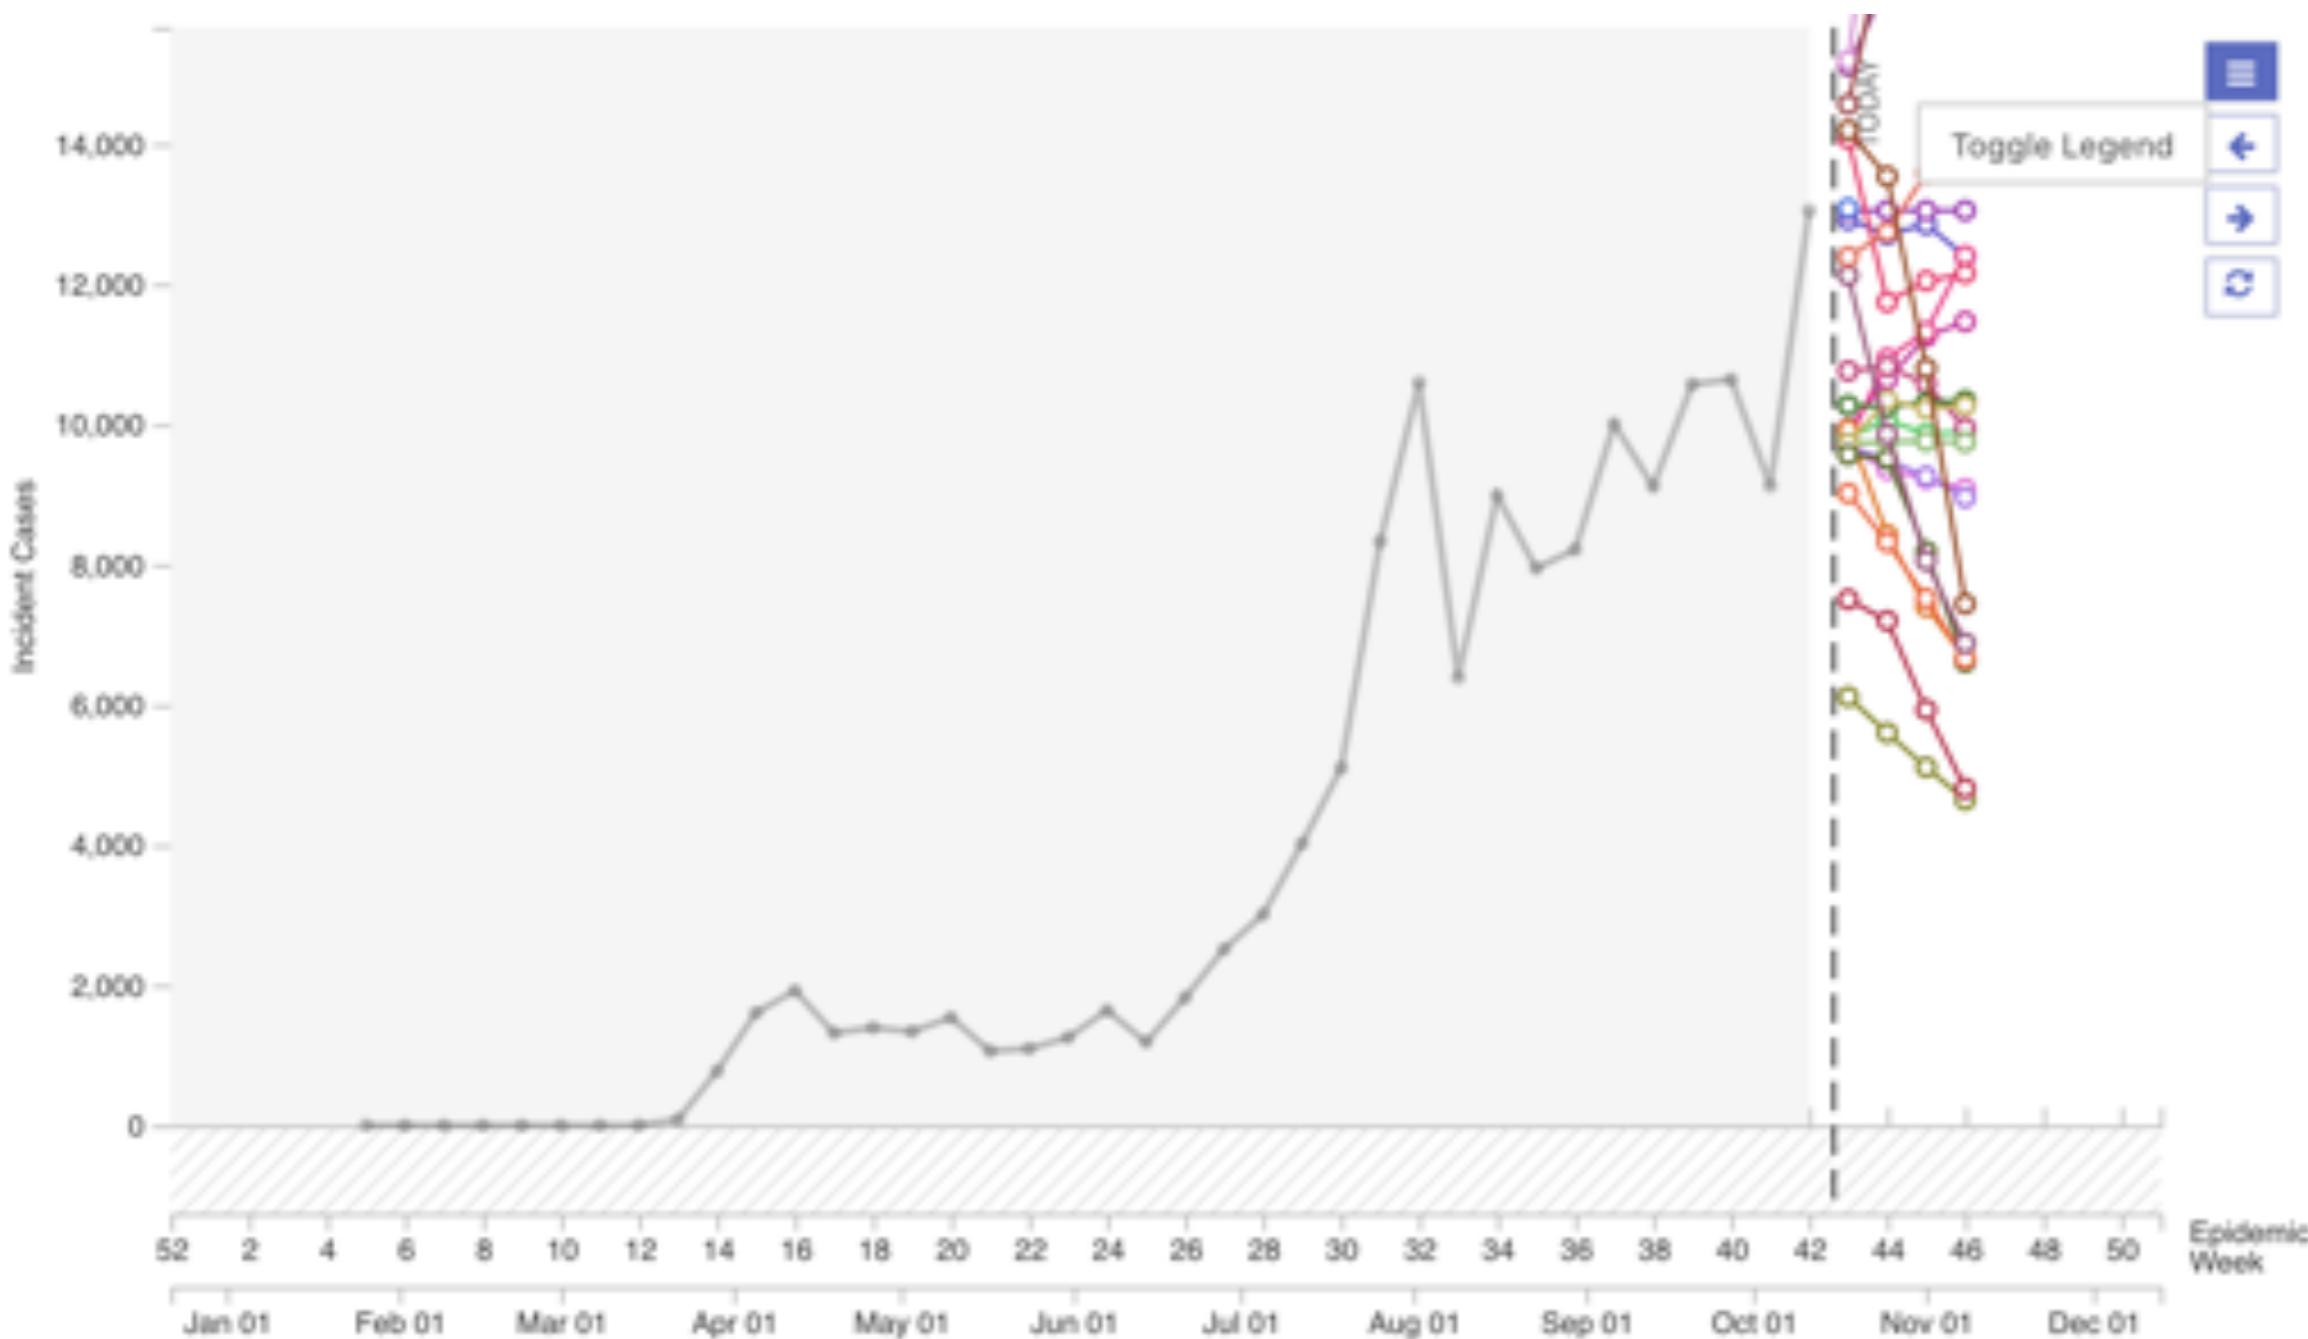
October 12, spike in cases in the State of Missouri is due to a database error. We are monitoring the dashboard and will redistribute if the error is fixed. [News report](#)



Missouri COVID-19 Forecasts

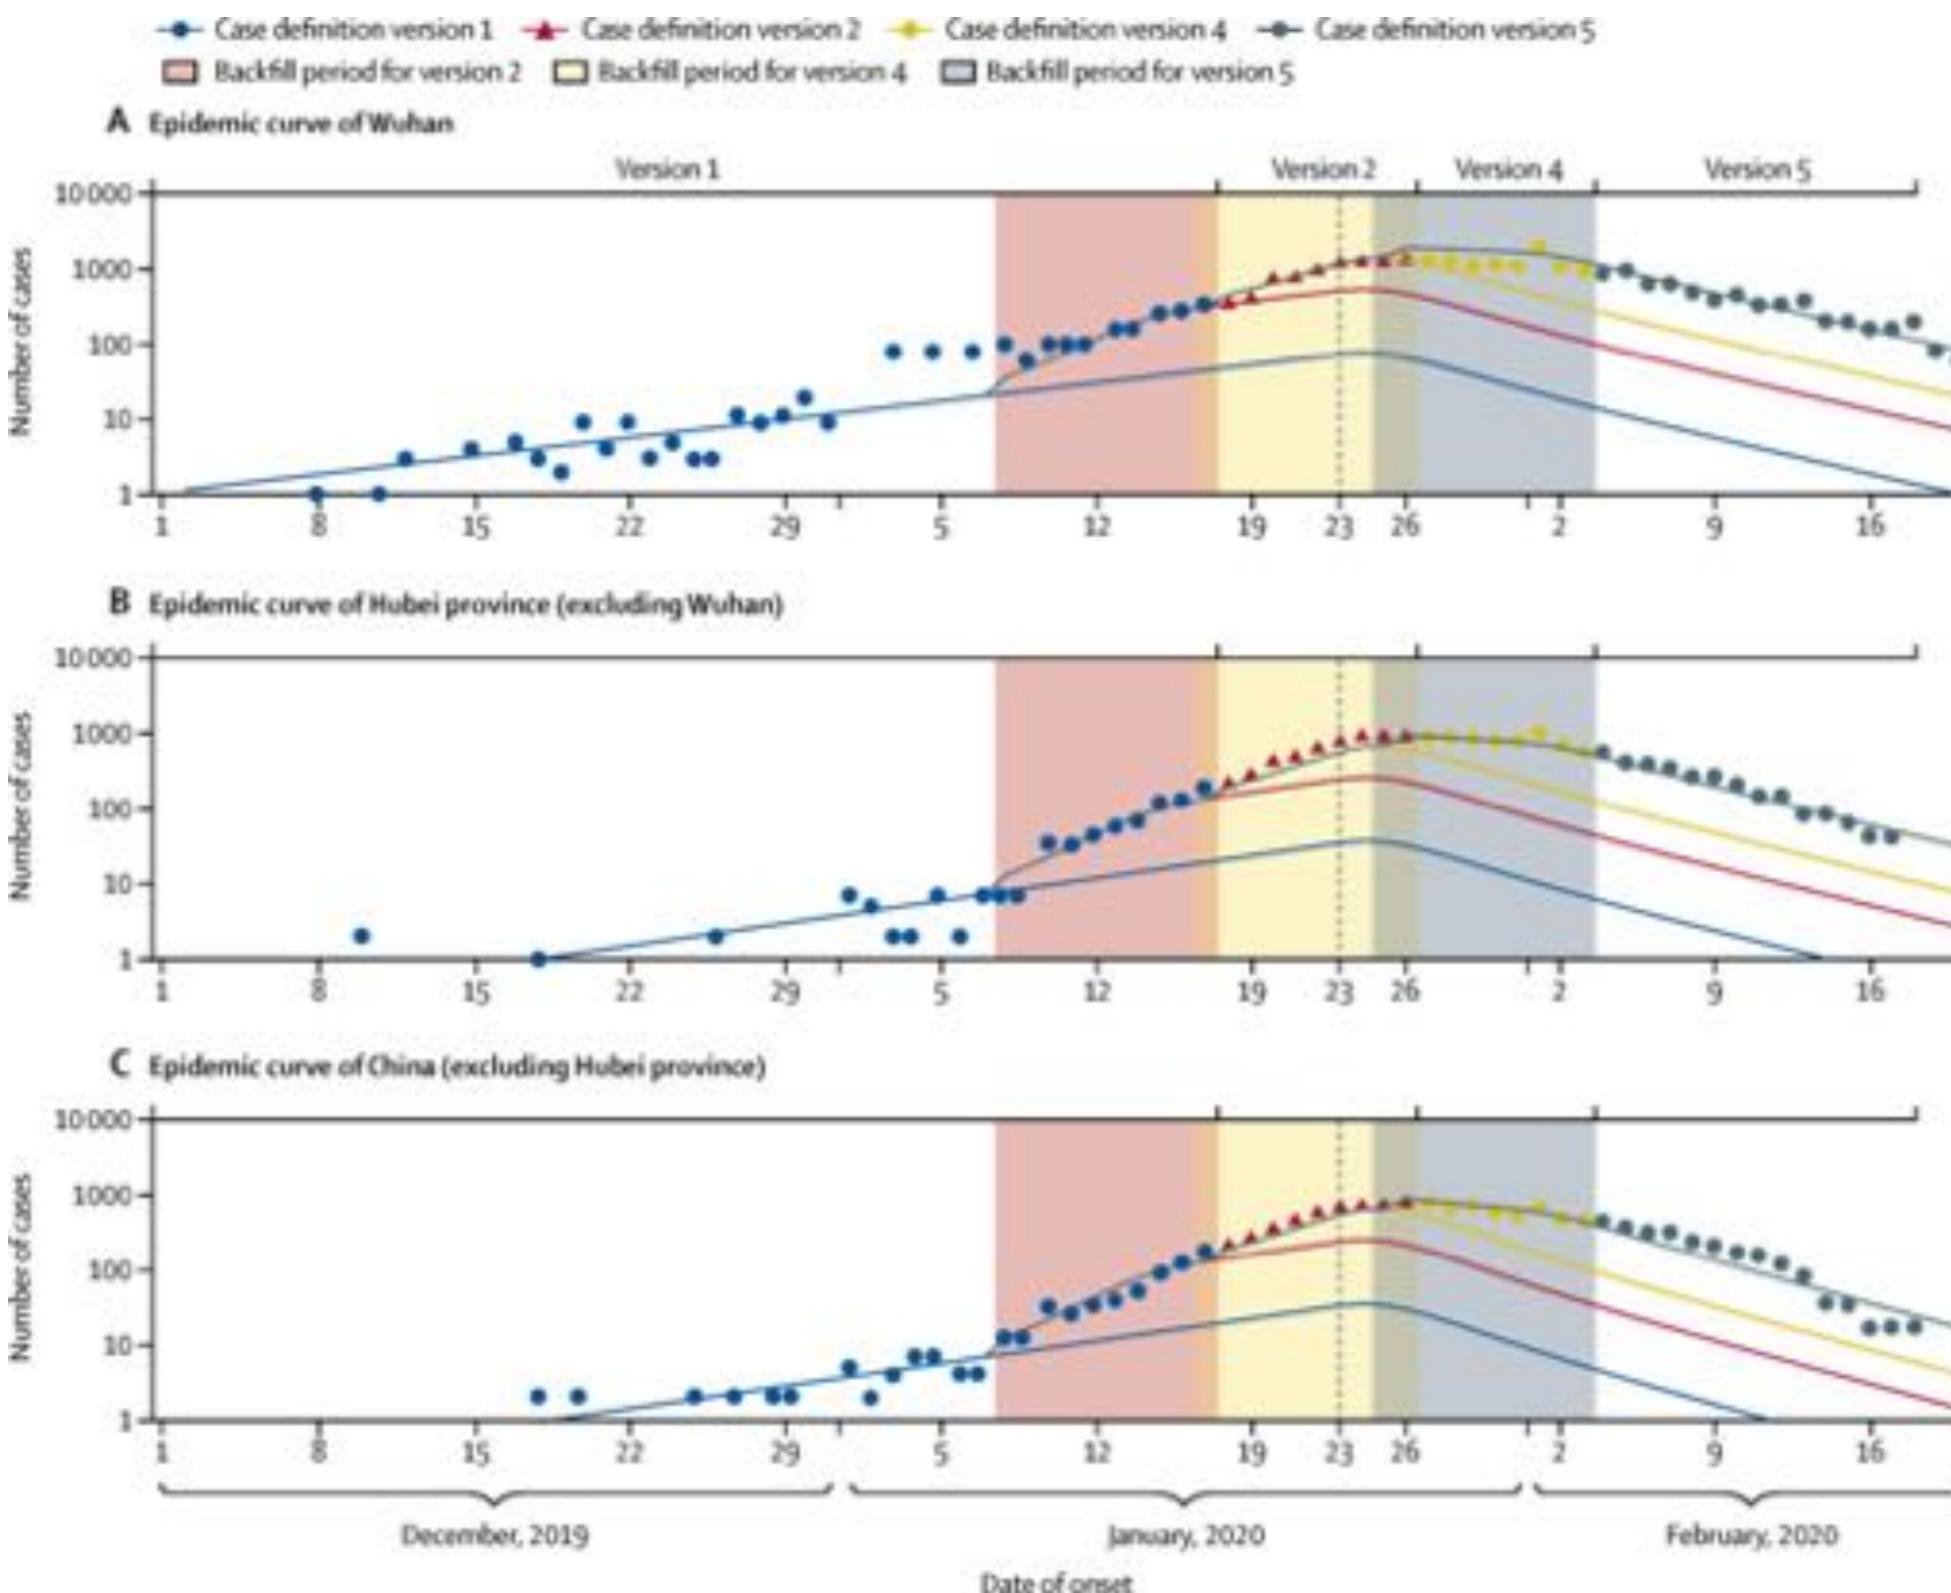
github.com/CSSEGISandData/COVID-19/tree/master/csse_covid_19_data

October 12, spike in cases in the State of Missouri is due to a database error. We are monitoring the dashboard and will redistribute if the error is fixed. [News report](#)



China COVID-19 Backfill

[www.thelancet.com/journals/lanpub/article/PIIS2468-2667\(20\)30089-X/fulltext](http://www.thelancet.com/journals/lanpub/article/PIIS2468-2667(20)30089-X/fulltext)



COVID-19 Backfill

COVID-19 backfill questions:

- Should models be required to model backfill?
 - No:
 - Just a data artifact, not interesting
 - Yes:
 - Check for non-negative incident, but backfill causes negative values
 - Evaluate models based on observed data
- Should models be required to target JHU CSSE data?
 - No:
 - JHU CSSE data choice was arbitrary
 - The Hub can model an “offset”
 - Yes:
 - If you want your model to be evaluated fairly, yes.

Forecast collection



COVID-19
ForecastHub

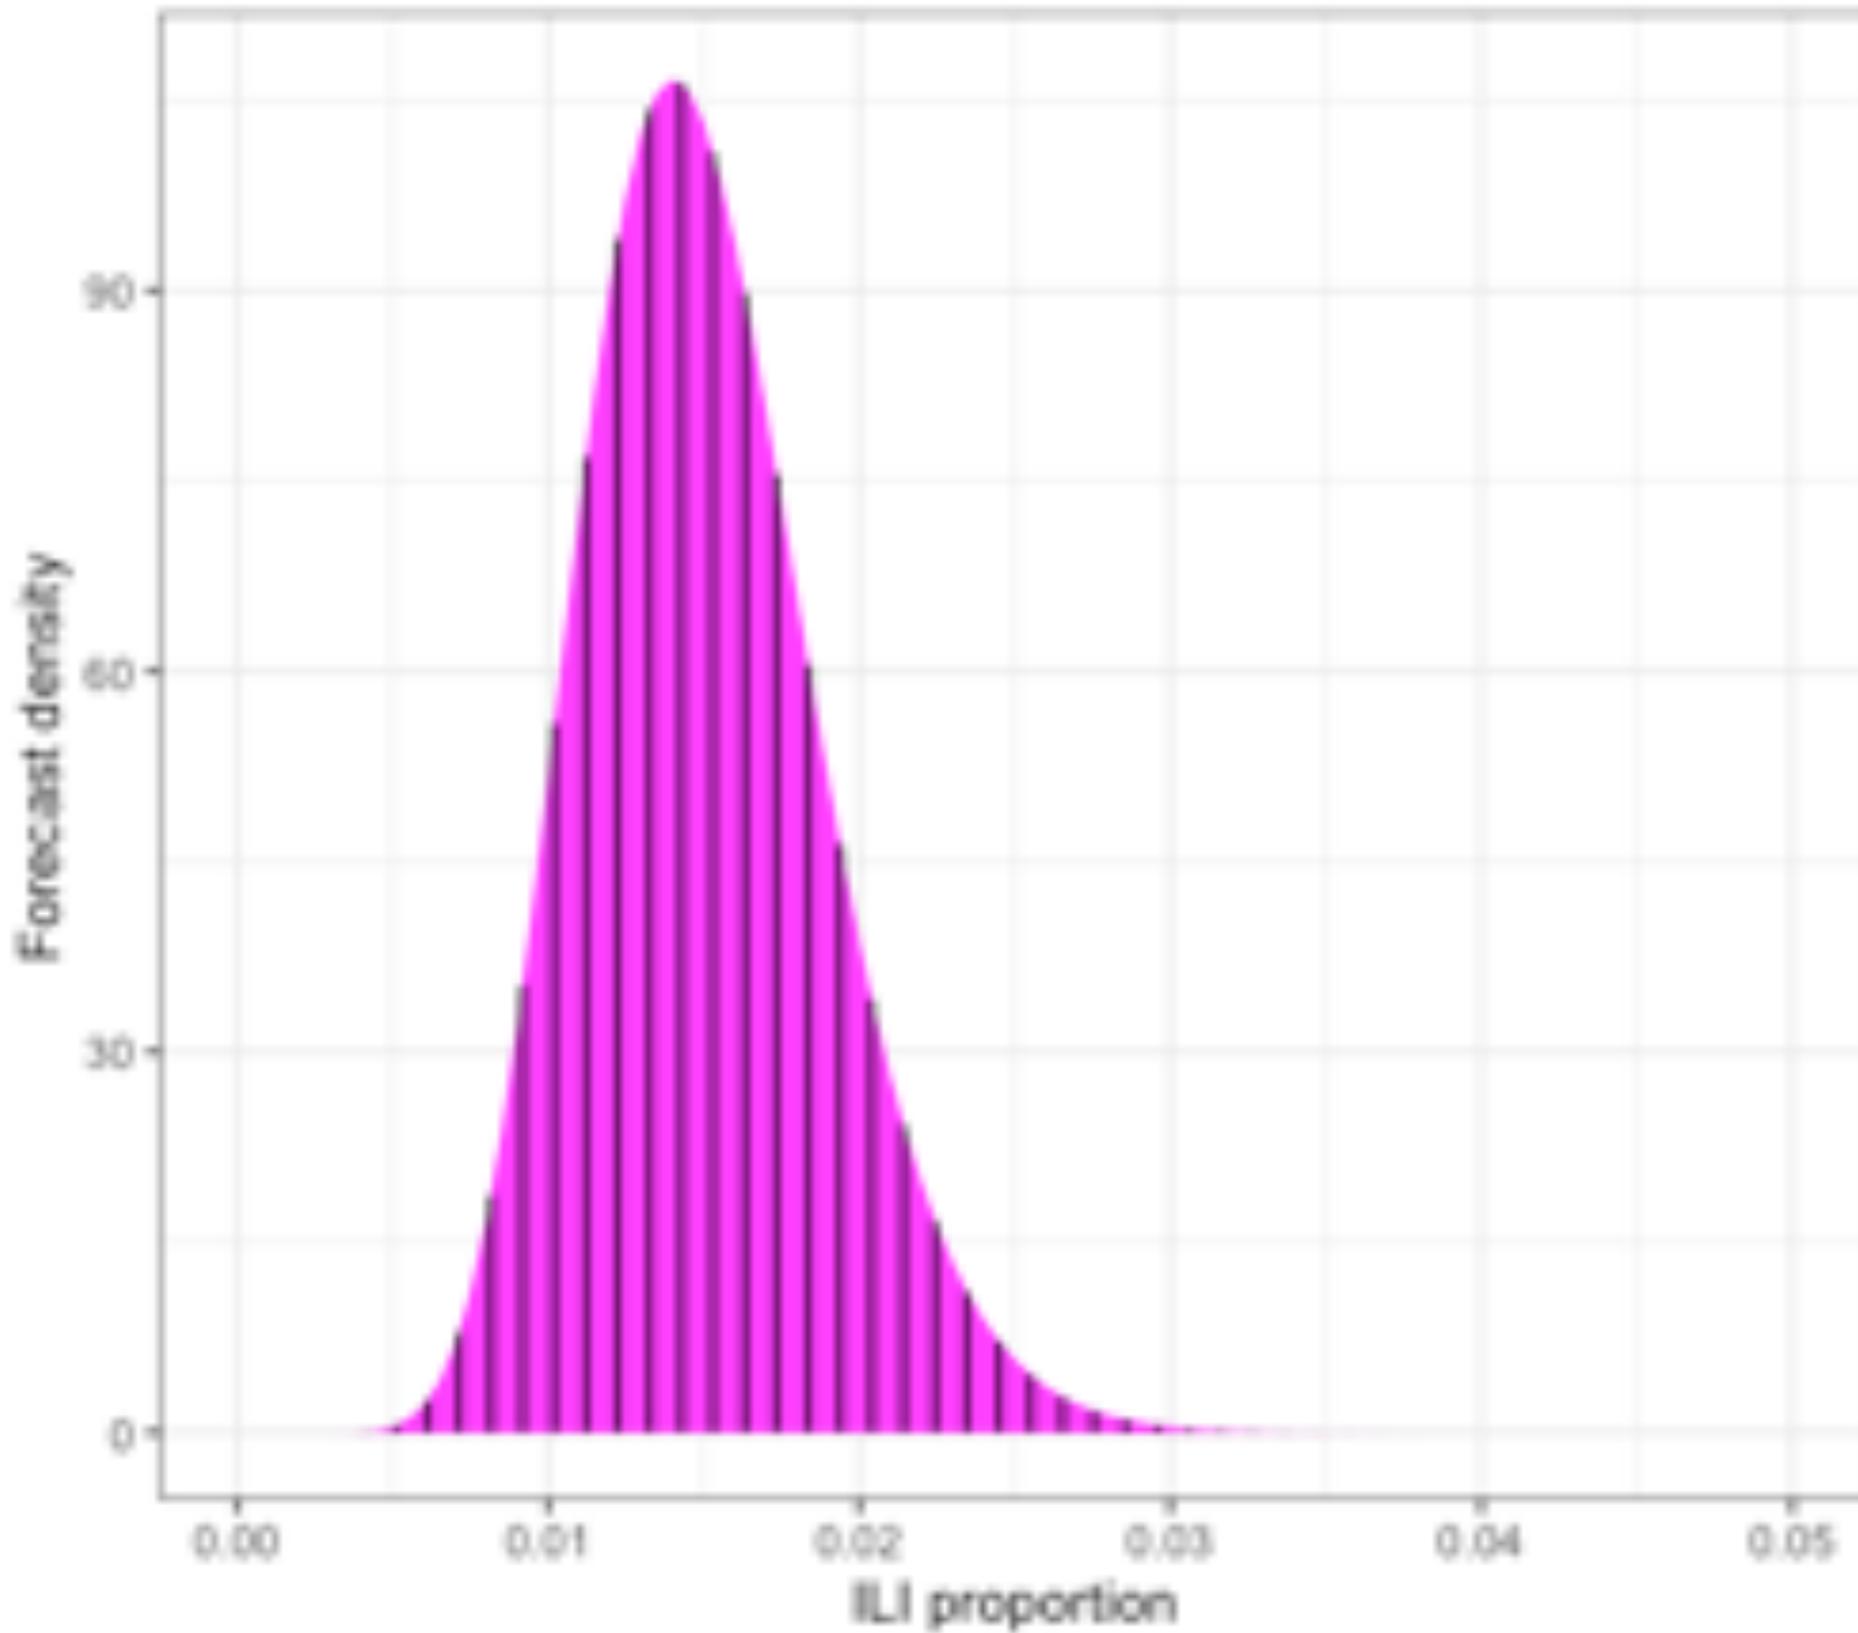
Standardization

github.com/reichlab/covid19-forecast-hub/wiki/Forecast-Checks

- checks `value` must be an int or float and non-negative
- validates date alignment as documented in the issue [add additional validations](#)
- validates quantiles and values (i.e., at the prediction level):
 - checks that entries in `value` must be non-decreasing as quantiles increase
 - checks that elements in the `quantile` are unique
- validates quantiles as a group:
 - there must be exactly one point prediction for each `location/target` pair
- validates if the prediction `value` for a location is at least less than the population of that location.
 - this check is run for all forecast submissions for all targets (in/cum deaths/cases).
 - the population truth data is present in [the locations.csv file](#).
 - To check which predictions are violating, check the logs in the Github Actions build do your PR and it should be printed.

Probabilistic ILI Forecasts

Example influenza forecast density with bins shown



Anti-probabilistic forecasts

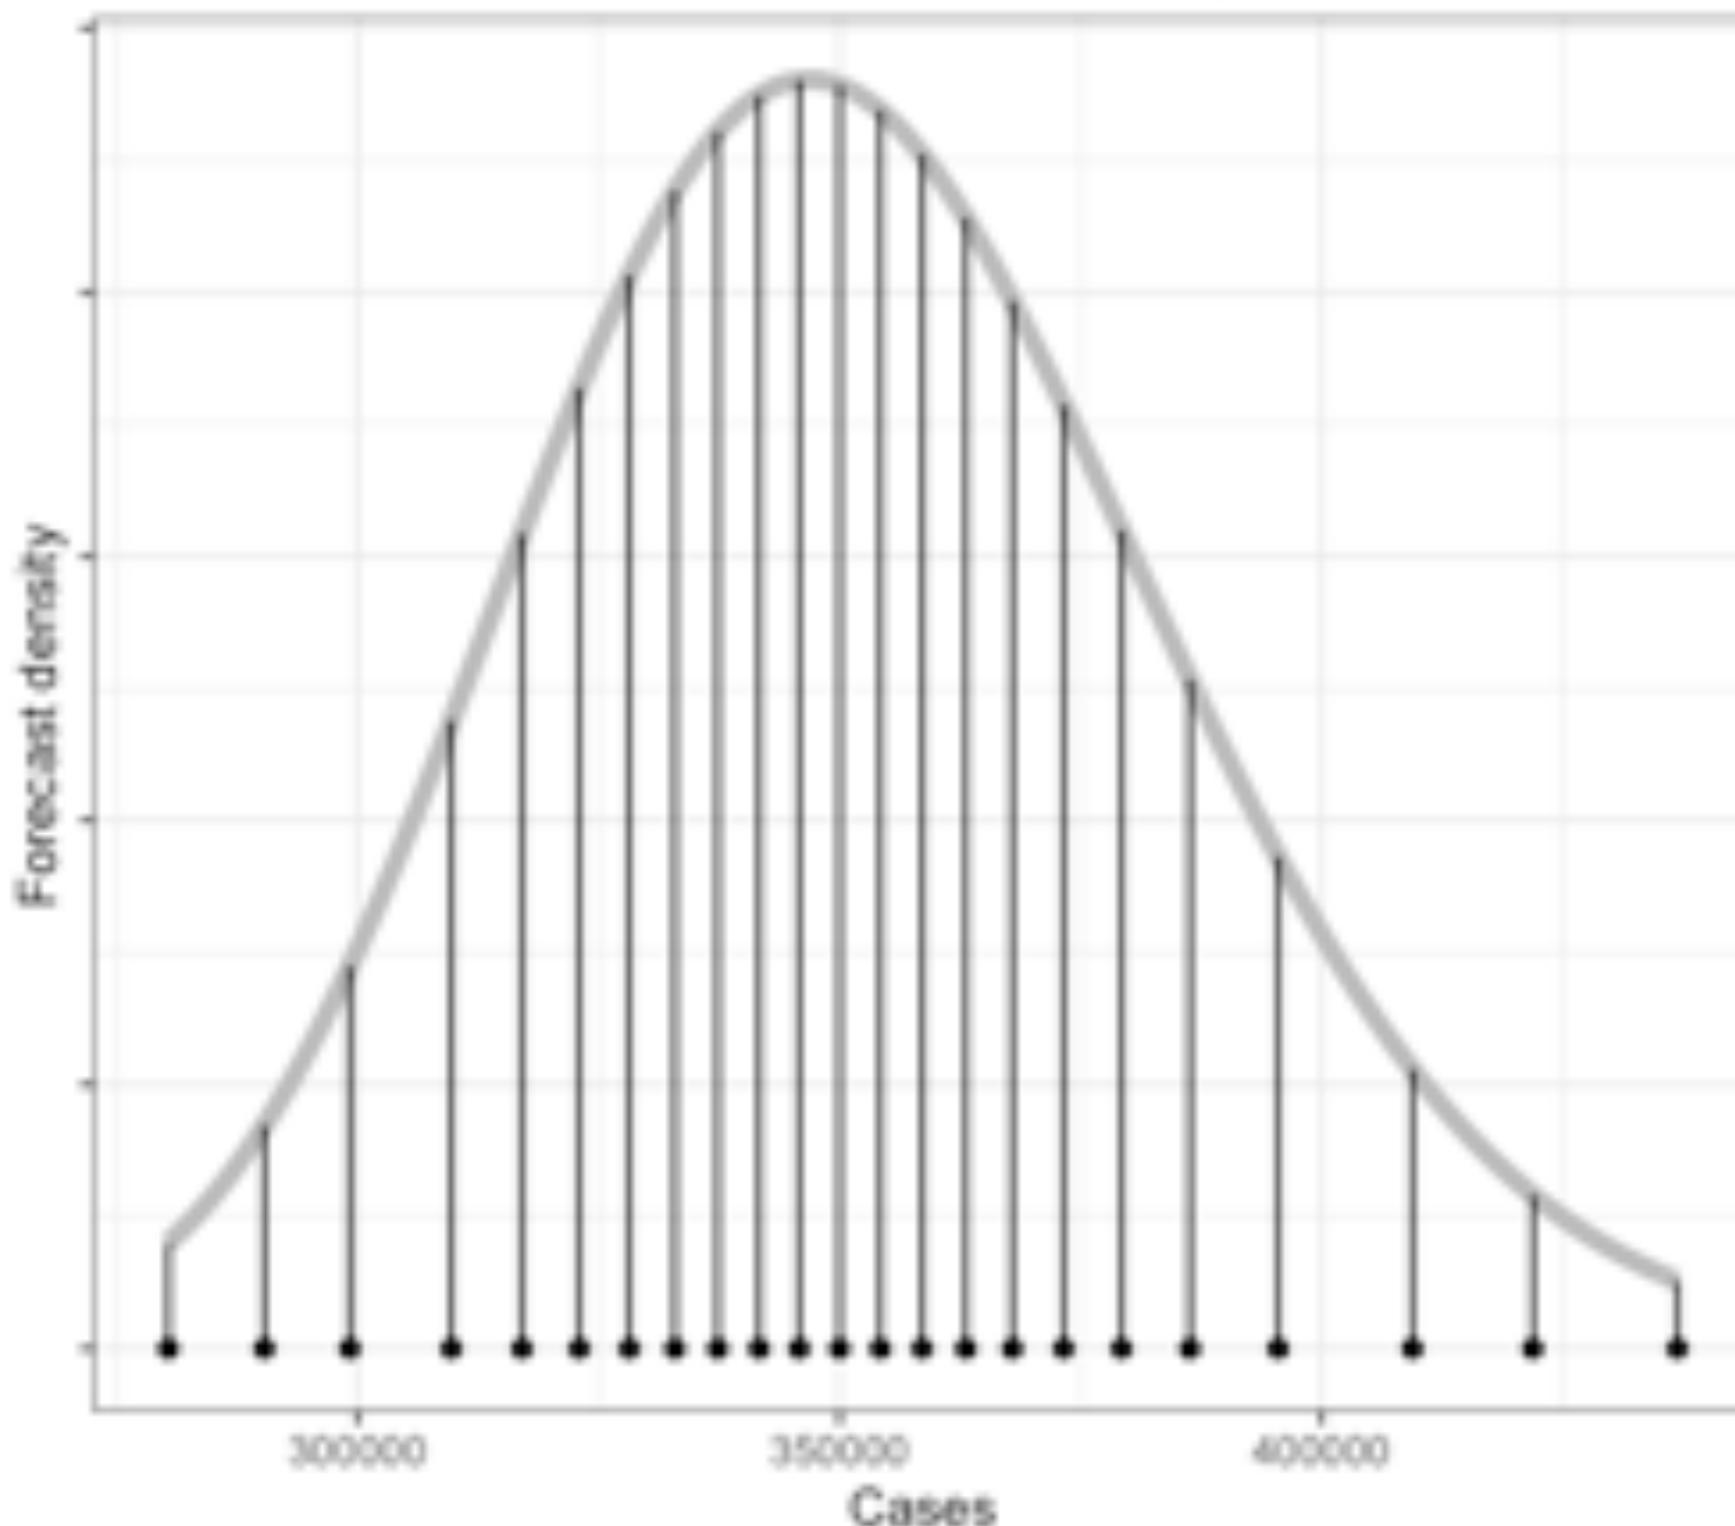
Four Different Forecast Evaluation Challenges:

- ❑ National Point Forecasts
 - One time series (per model)
 - Excellent accuracy over time
- ❑ National Prediction Interval Forecasts
 - Purpose of these is not clear, considering the accuracy of the national point forecasts
- ❑ State Point Forecasts
 - 51+ time series (per model)
 - 3 orders of magnitude difference across the time series
- ❑ State Prediction Interval Forecasts
 - Purpose of these is clear, considering the (in)accuracy of the point forecasts

Probabilistic COVID-19 Forecasts

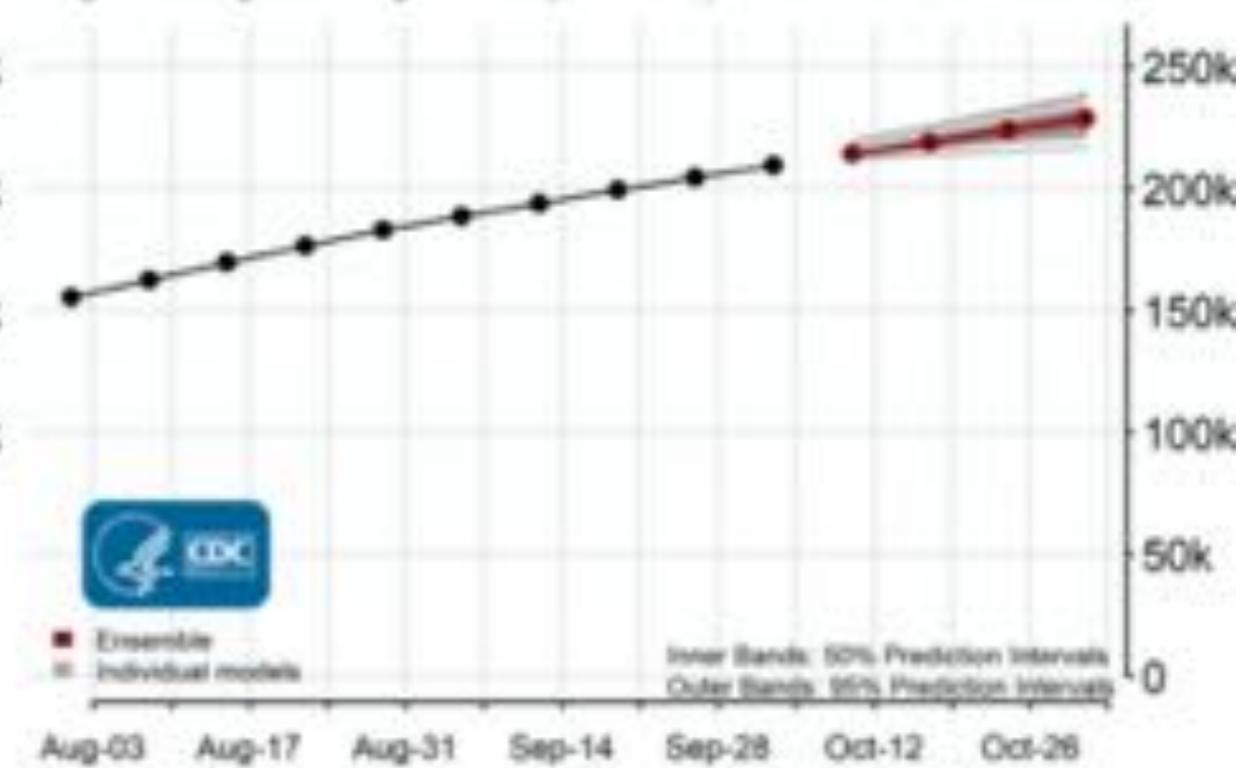
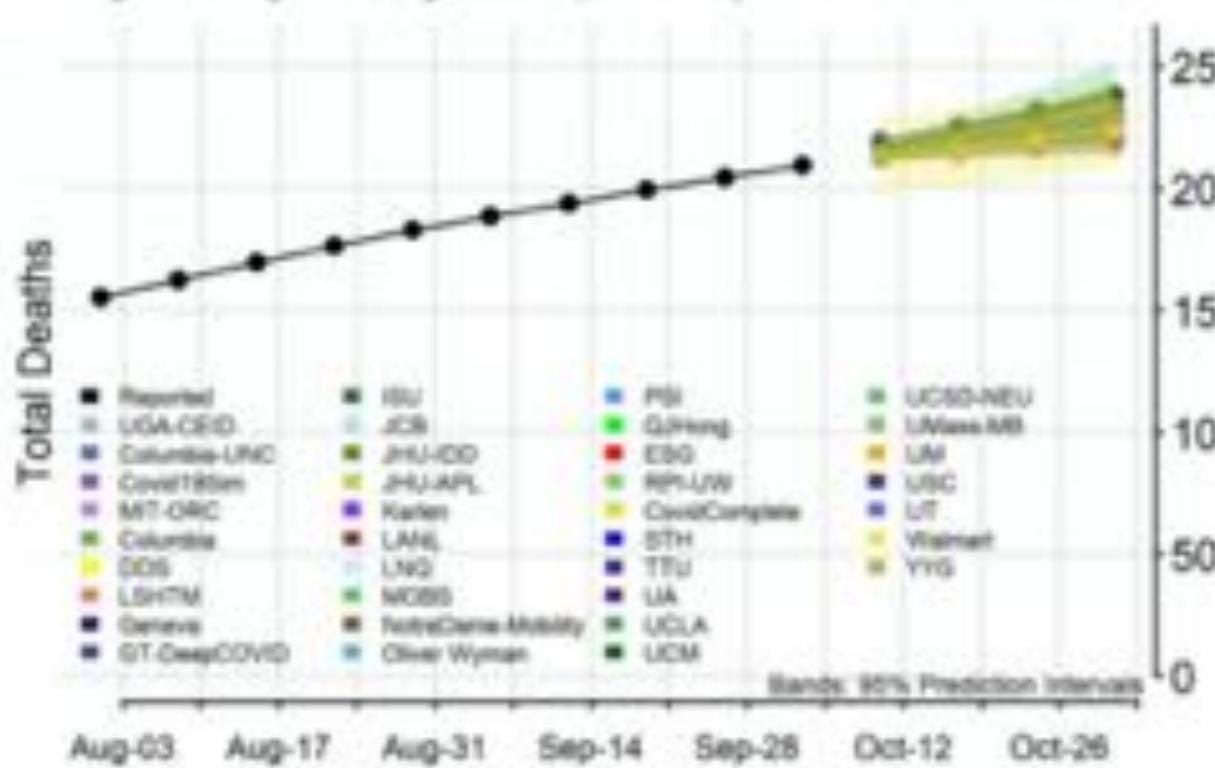
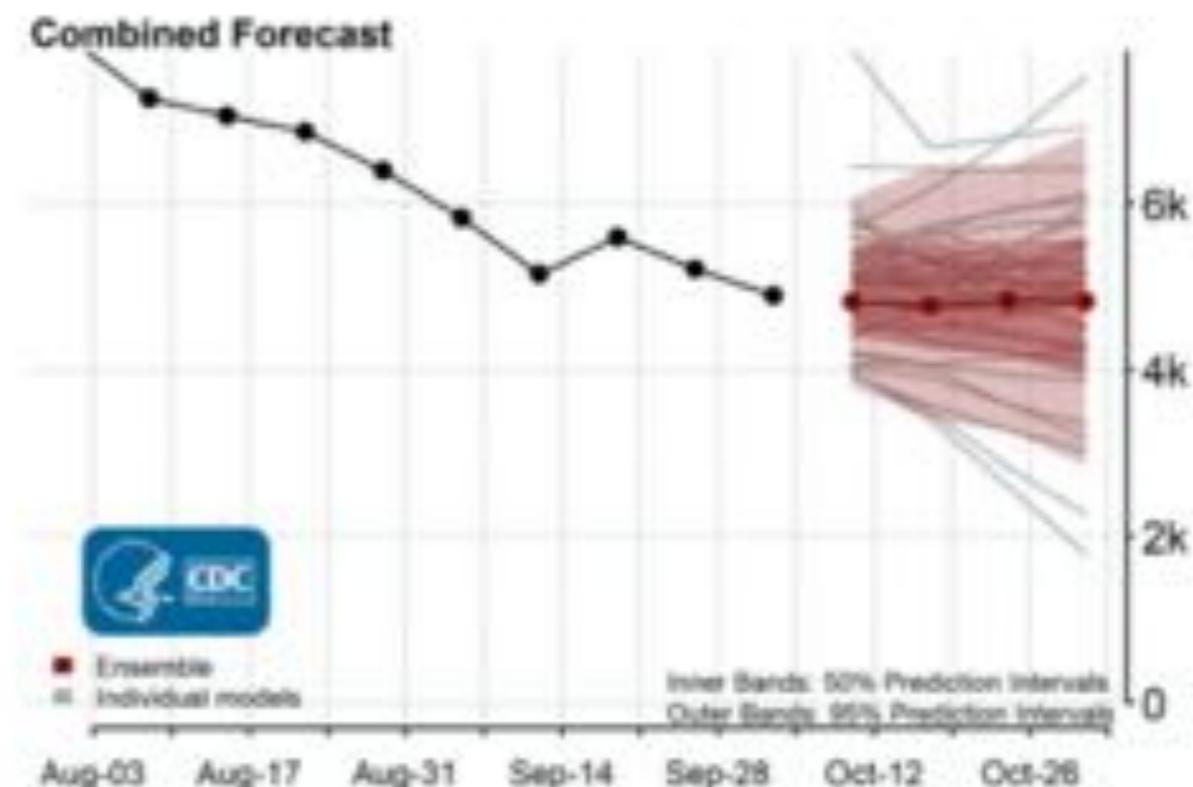
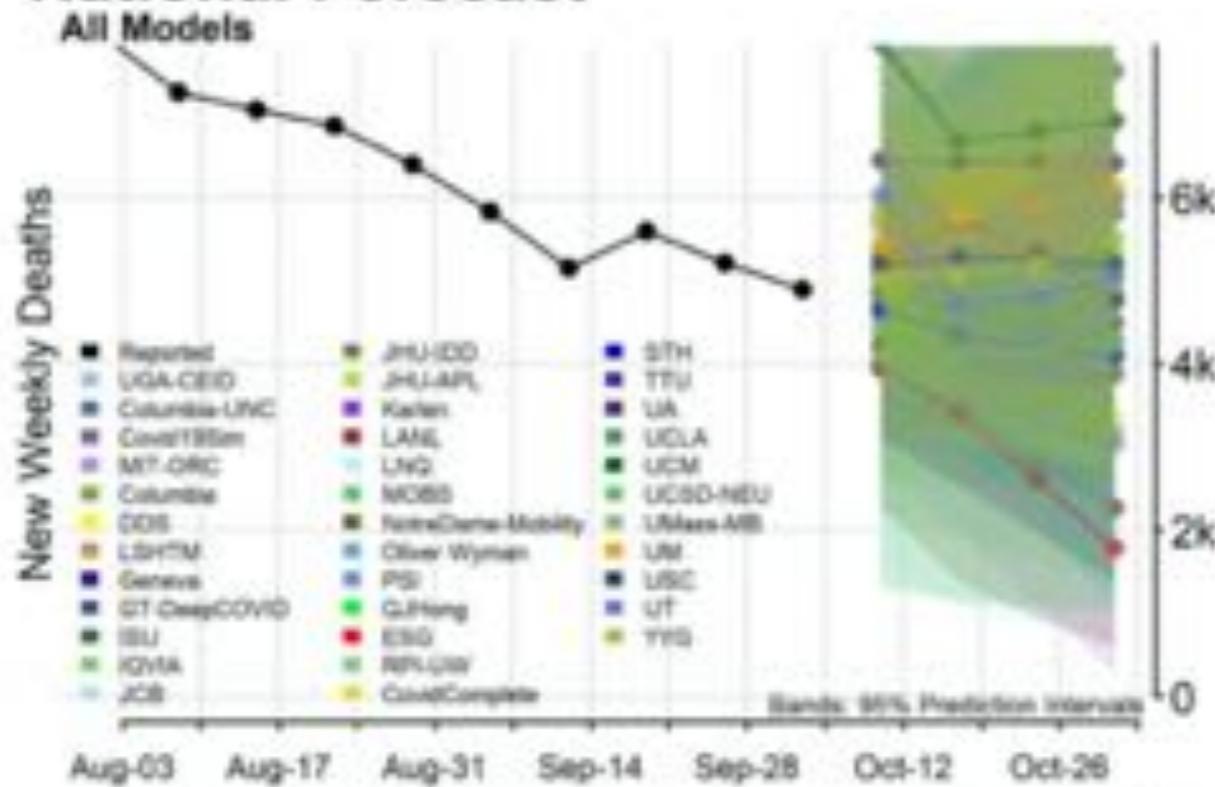
```
quantiles = c(0.01, 0.025, seq(0.05, 0.95, by = 0.05), 0.975, 0.99)
```

Example COVID-19 case forecast density



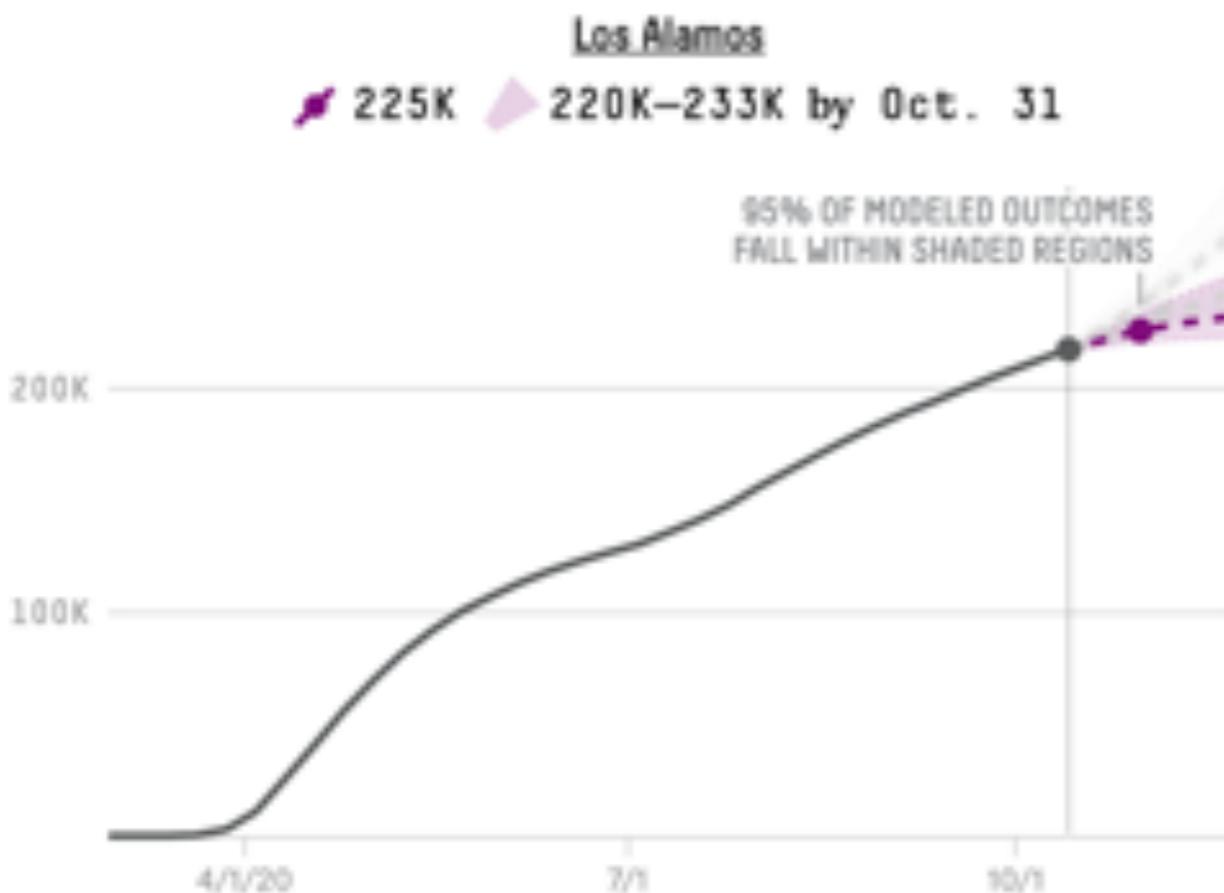
CDC COVID-19 Forecast

National Forecast



538 COVID-19 Forecasts

projects.fivethirtyeight.com/covid-forecasts



This model assumes that there will continue to be interventions, such as stay-at-home orders, but it does not specifically assume what those interventions will be. Instead, it considers various possible interventions to arrive at its forecast, which typically results in wider prediction intervals than a model with stricter assumptions.



This model accounts for state reopenings, and assumes that interventions would be reenacted if cases continue to increase. The model was changed significantly on July 4.

COVID-19 Forecast Questions

Define “forecast date”:

- Date of the last data used
- Date when model computations started
- Date when model computations ended
- Date when forecast is available
 - Relevant for real-time forecasts

Submission date compared to “forecast date”

- Same or next day

COVID-19 Forecast Questions

- Should models be allowed to revise their forecasts?
 - No:
 - Modelers need to stand by their work
 - Modelers could cheat if allowed to revise
 - Downstream forecasts are based on the forecasts received
 - Bugs can be consider a part of the modeling effort
 - Yes:
 - Bugs are inevitable and not representative of the method
 - Erroneous forecasts make the Hub look bad

COVID-19 Forecast Questions

- Should cumulative and incident forecasts be consistent?
 - No:
 - No reason modelers should have to use the same model
 - Due to data delay and computation time, it is impossible
 - Yes:
 - Otherwise ensemble won't be consistent
 - Modelers should fix their bugs

observed_cum	one_week_ahead_cum_forecast	implied_median_inc	actual_median_inc	diff
214370	215097.	727.	5257.	-4531.
214370	218830.	4460.	6566.	-2107
214370	218396.	4026.	5347.	-1321.
214370	219117	4747	5660	-913
214370	219070.	4700.	5514.	-814.
214370	220812.	6442.	7253.	-811.
214370	218476	4106	4738	-632
214370	219244	4874	5245	-371
214370	218598	4228	4500	-272
214370	219213	4843	5068	-225
214370	219594	5224	5033	191
214370	219271	4901	5089	-188
214370	217690.	3320.	3463.	-143.
214370	222008	7638	7731	-93
214370	218461.	4091.	4094.	-3.04
214370	218806	4436	4439	-3
214370	219316	4946	4948	-2

Ensemble construction

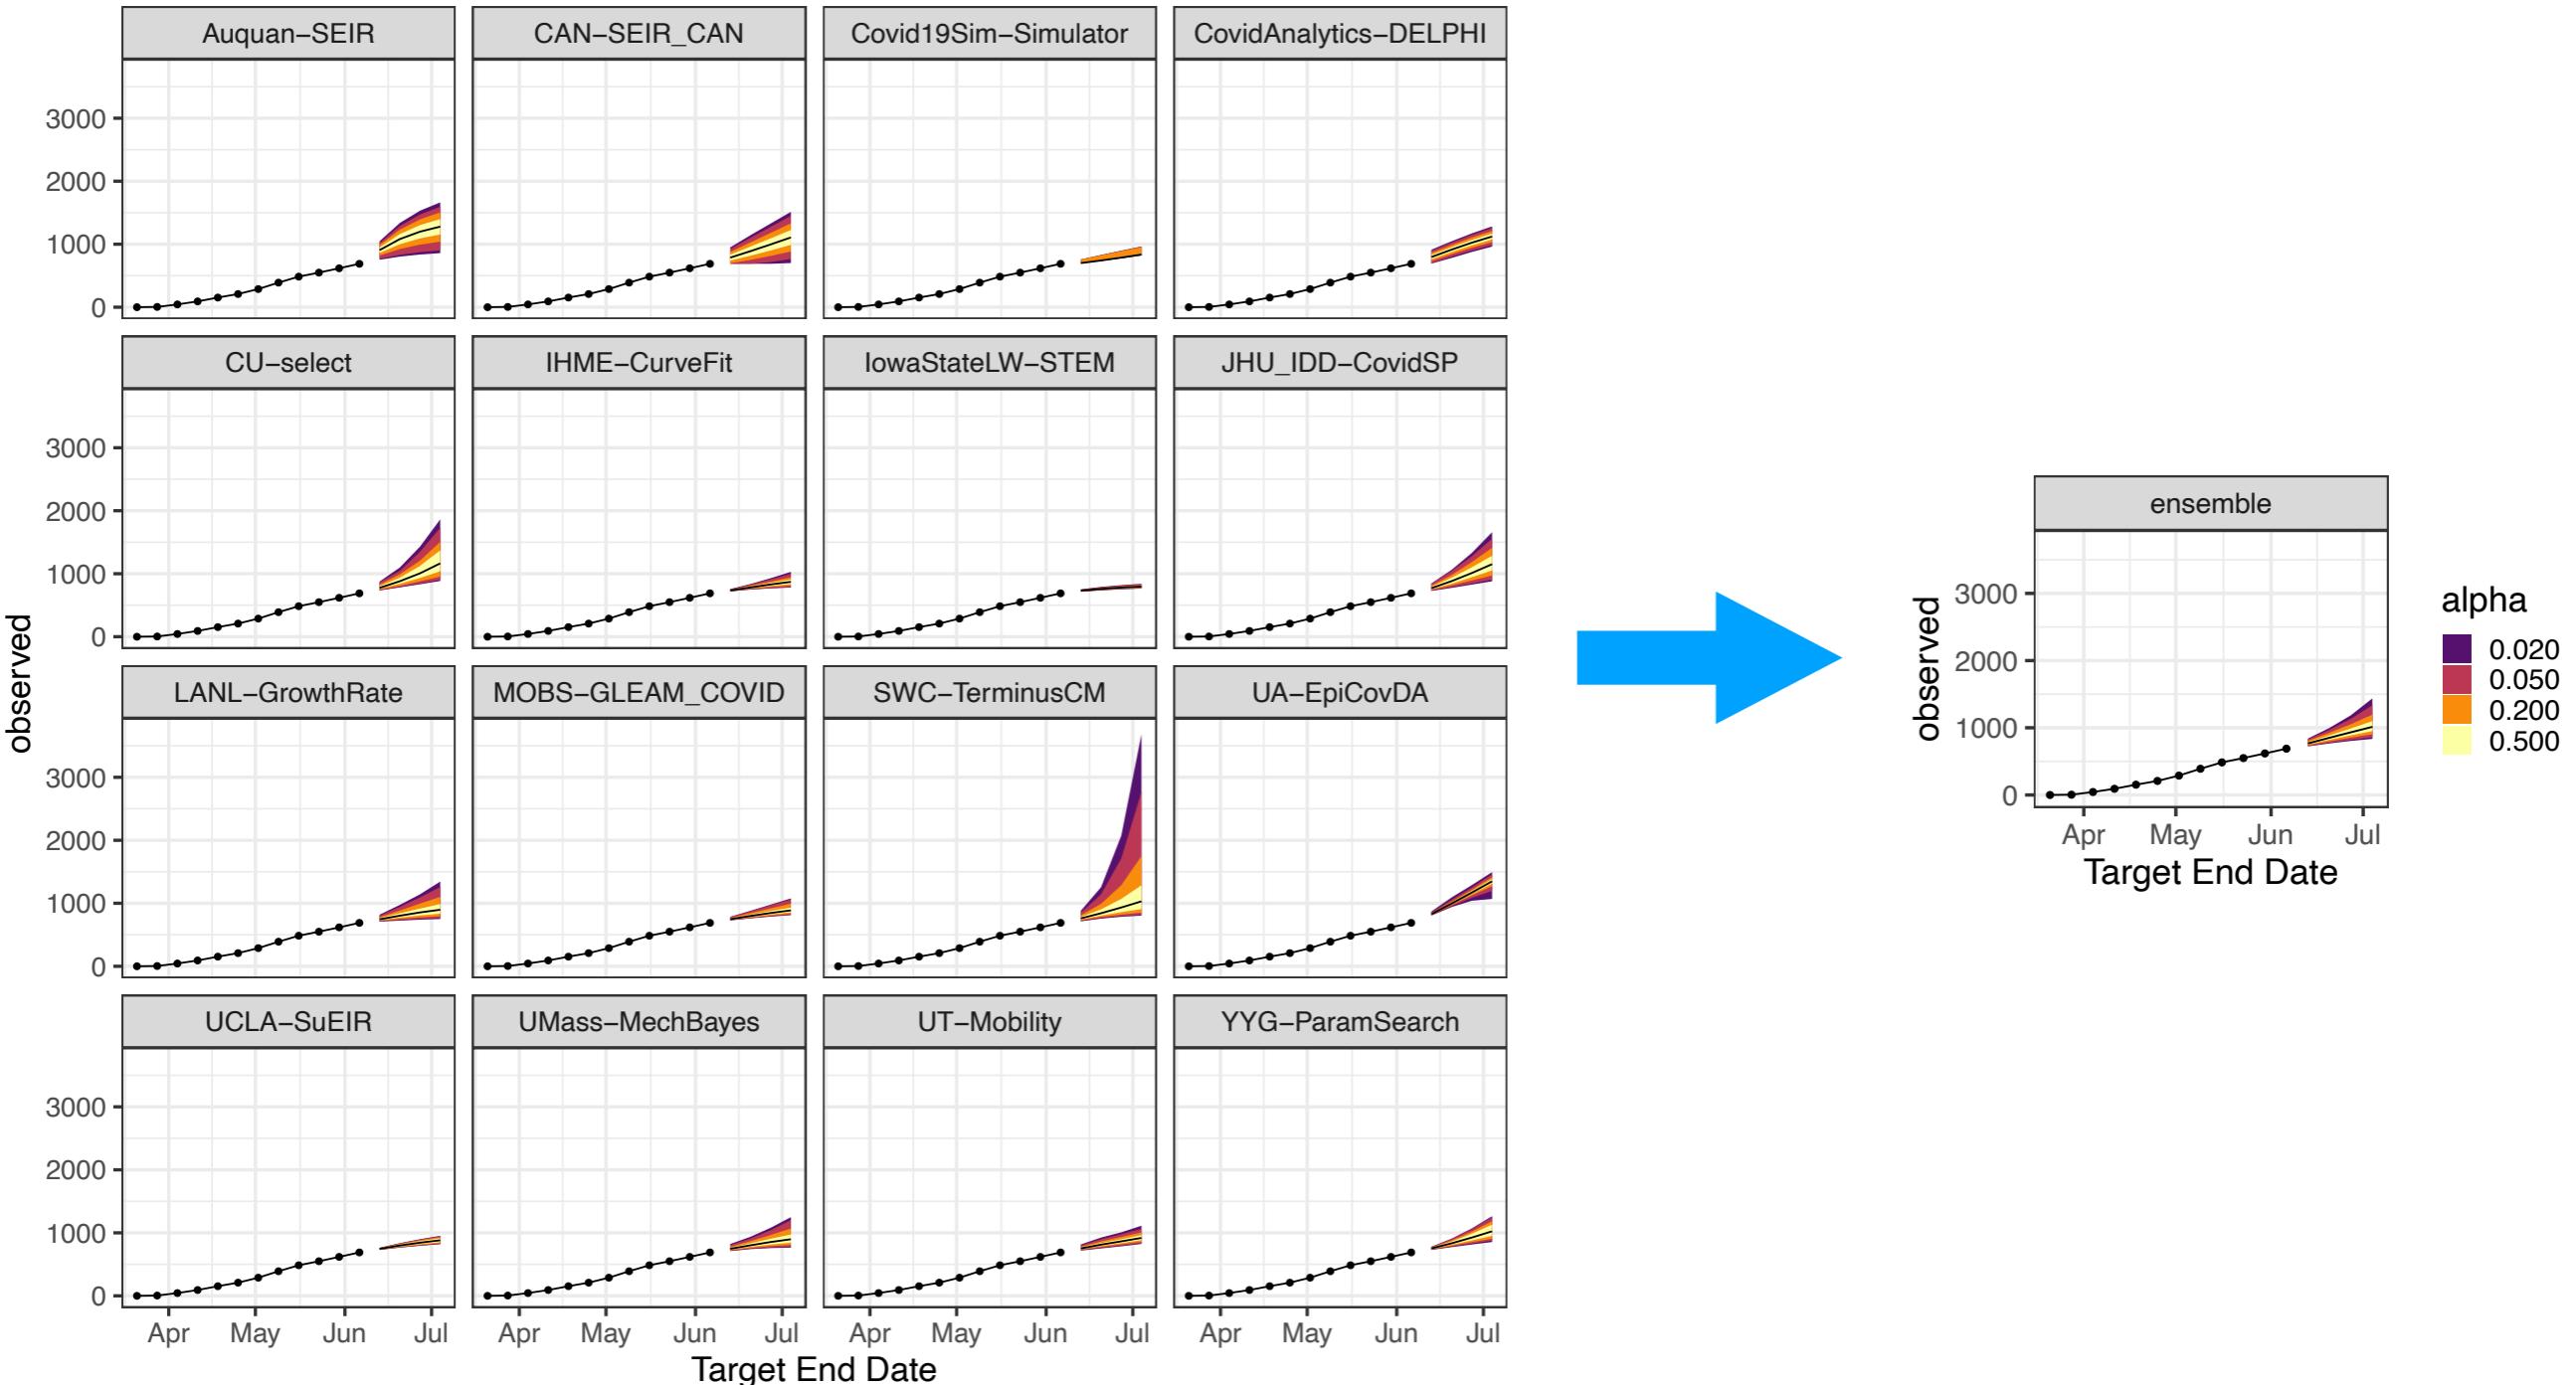


COVID-19
ForecastHub

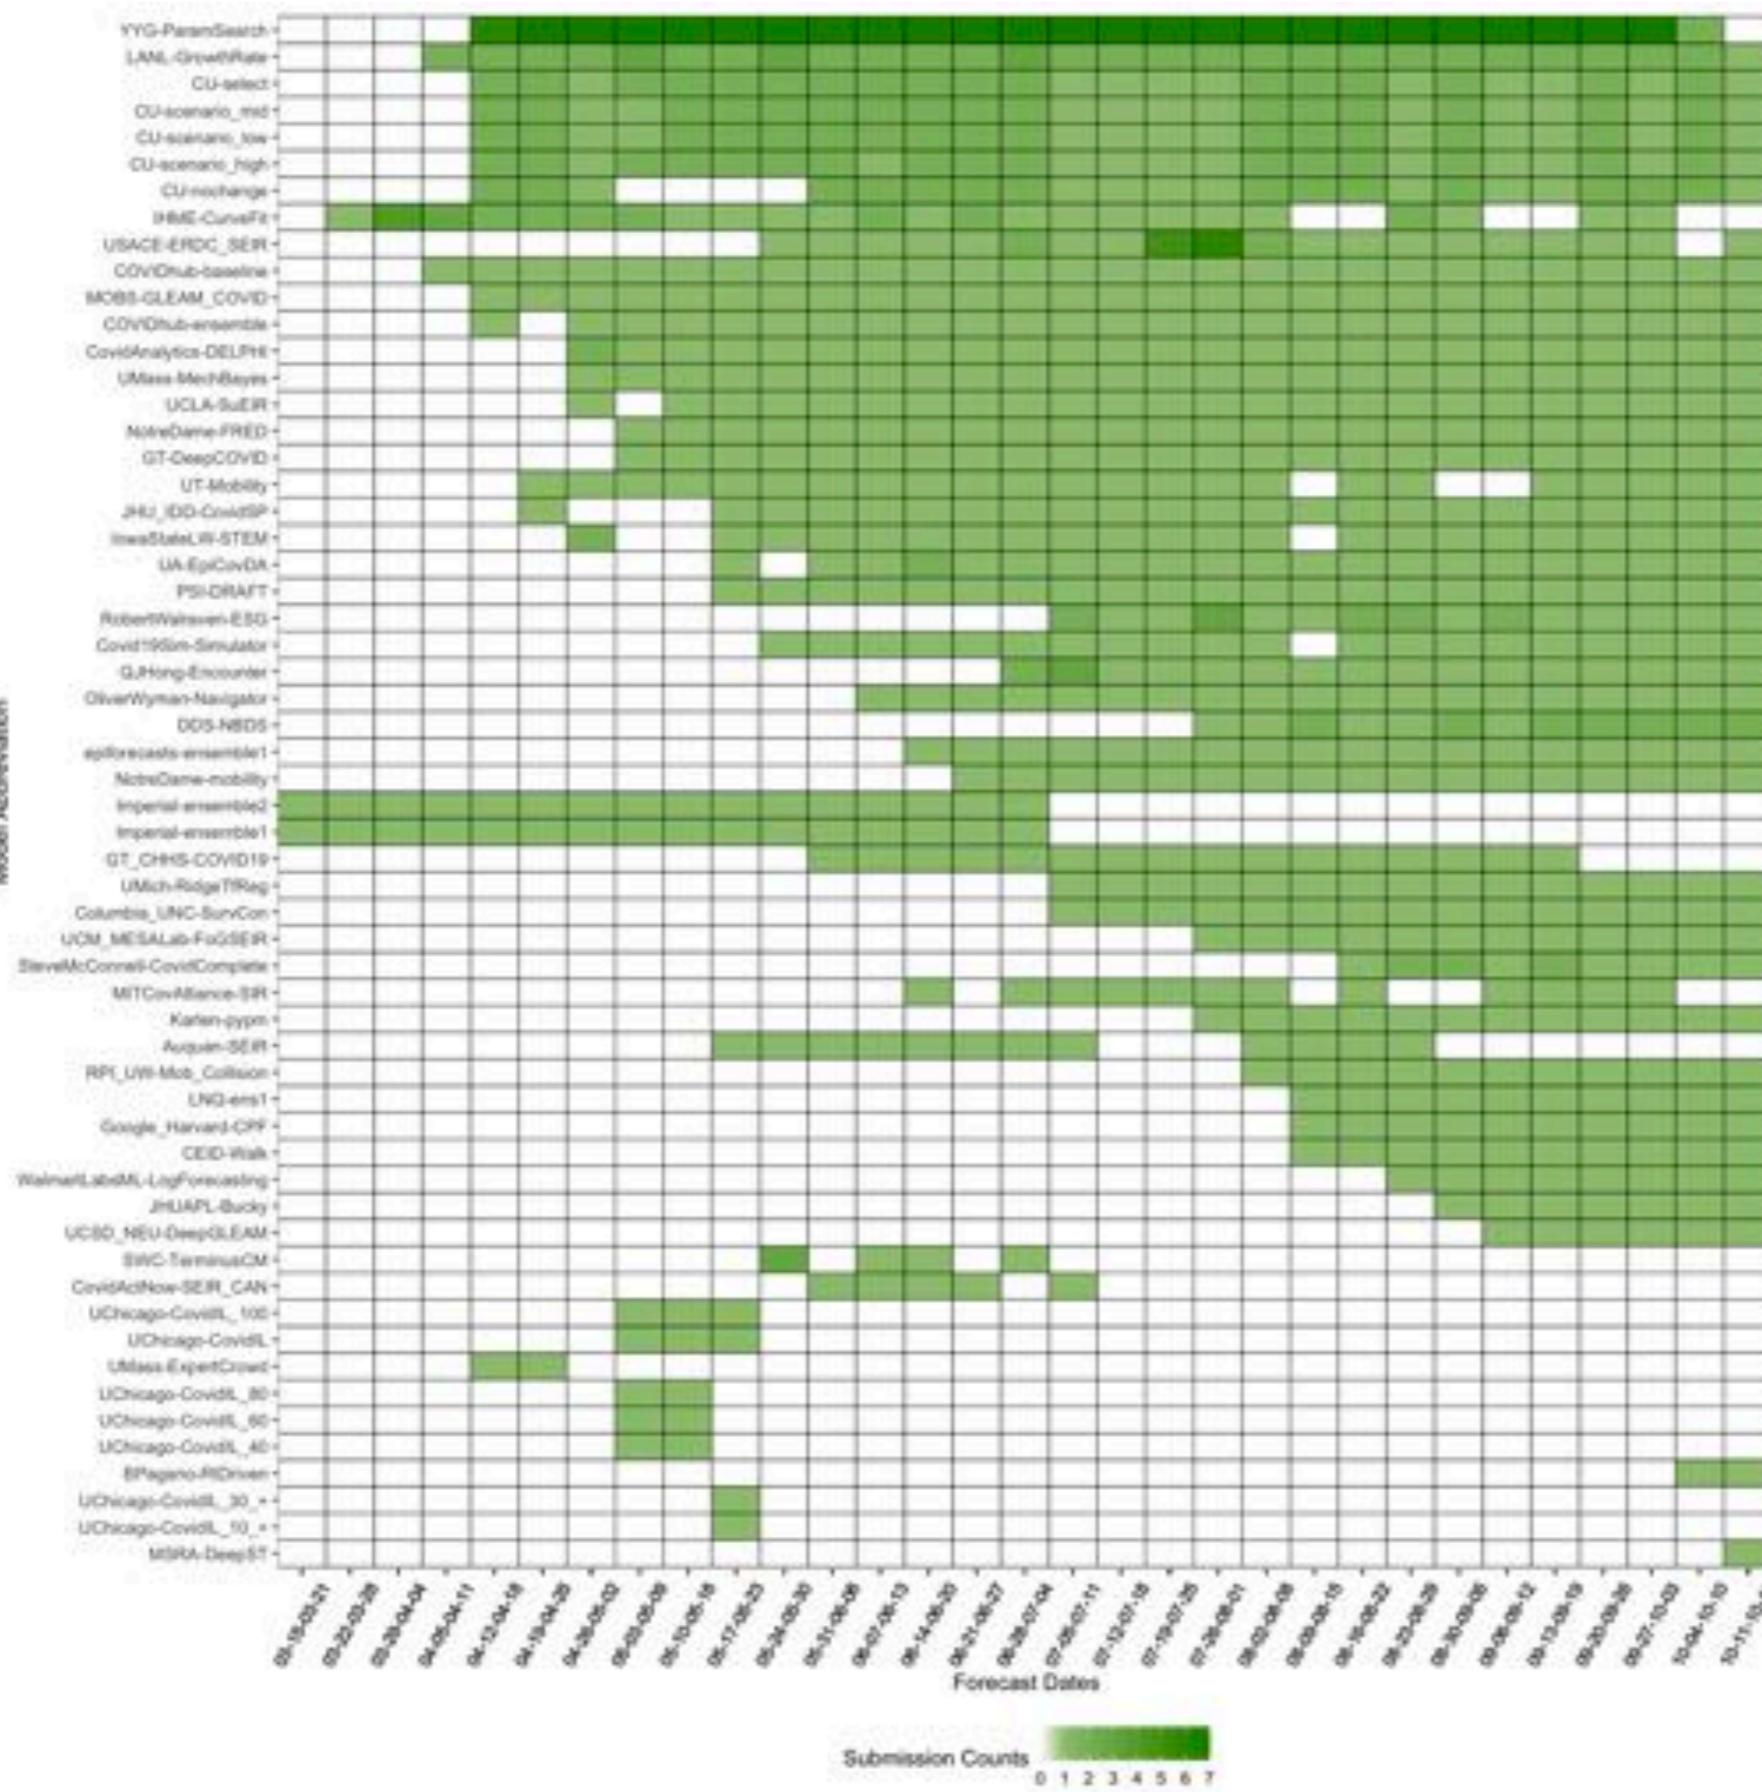
Context and Goals

- Each week we receive forecasts of weekly incident deaths and cumulative deaths due to COVID-19 from approximately 25 teams and 30 models.
- We build an ensemble that combines the predictions from these models.

Alabama

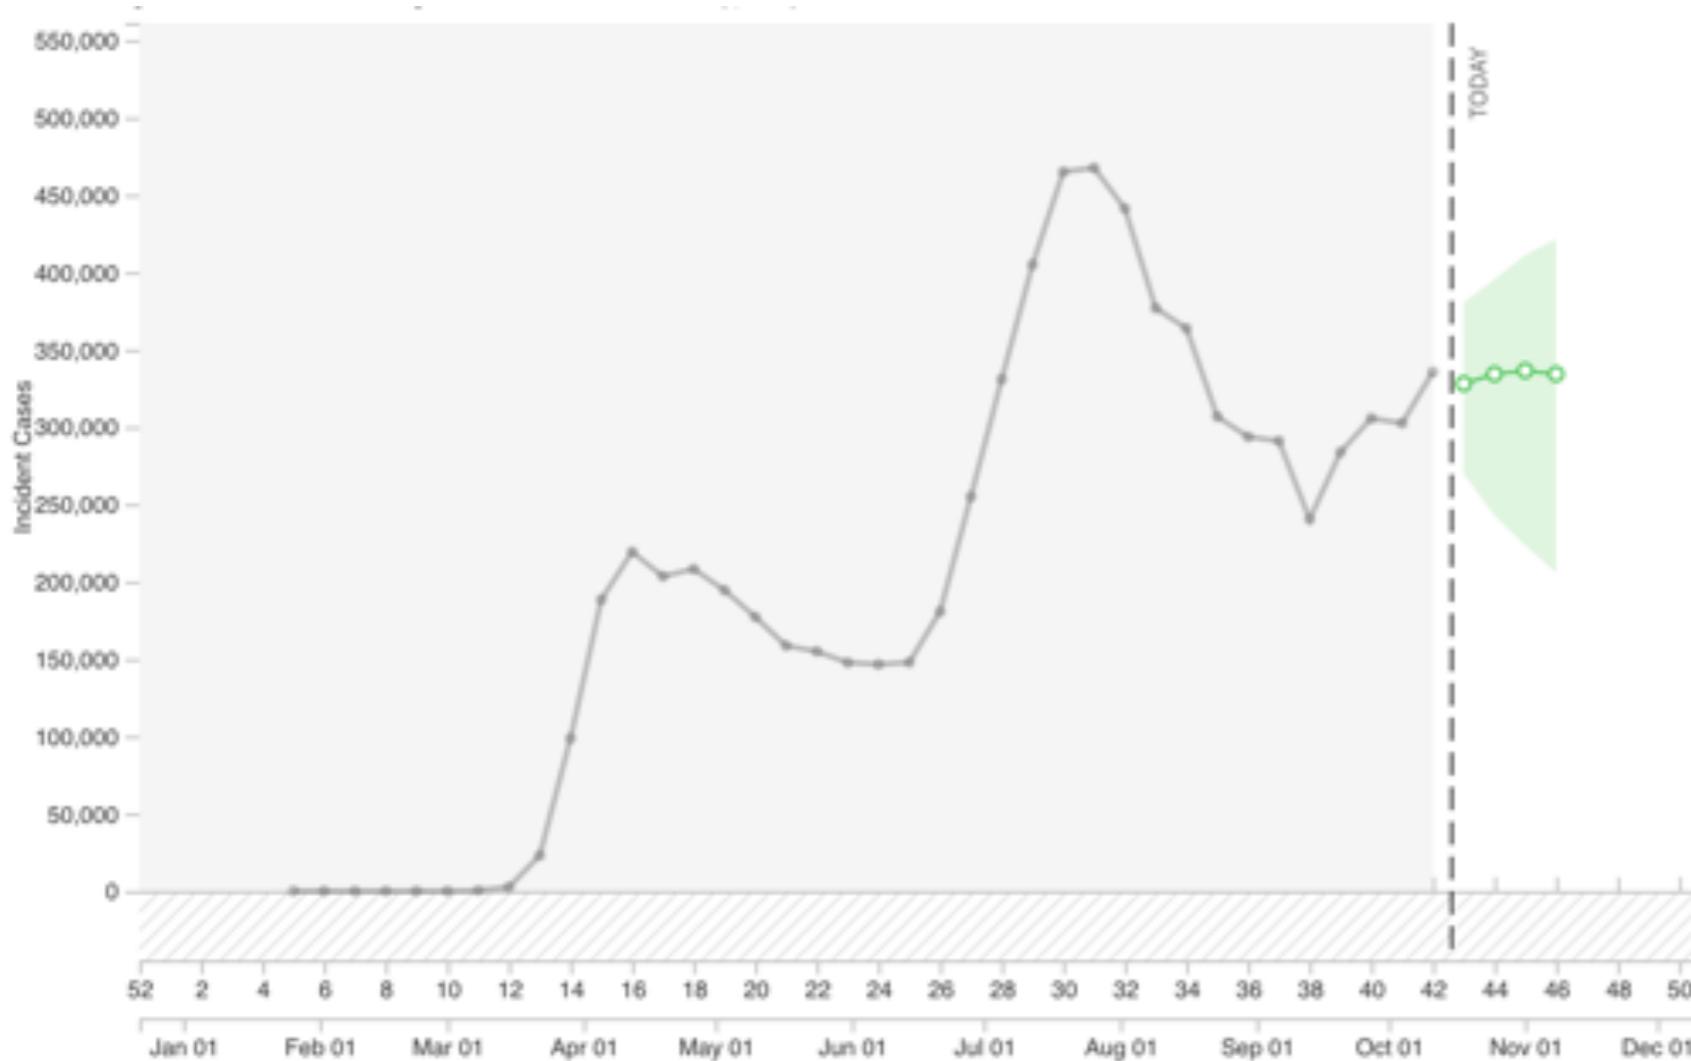


Models and submissions evolve over time



Forecasts collected from models

- Models are required to submit 23 quantiles of a predictive distribution:
$$\widehat{P}(Y \leq q_1) = 0.01, \widehat{P}(Y \leq q_2) = 0.025, \dots, \widehat{P}(Y \leq q_{23}) = 0.99$$
- Each model provides univariate predictive distributions for a set of "step-ahead" forecast targets, e.g. 1-4 week ahead incident reported deaths in a specific location, as reported by the JHU CSSE COVID-19 dashboard.



Weighted Average of Quantiles

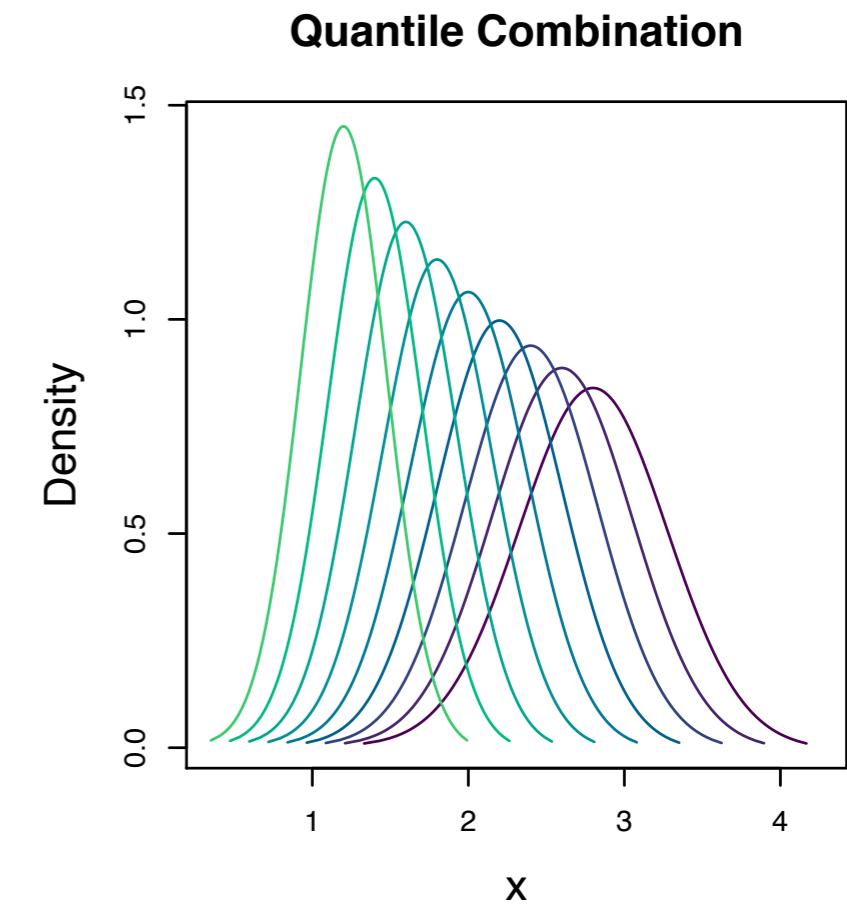
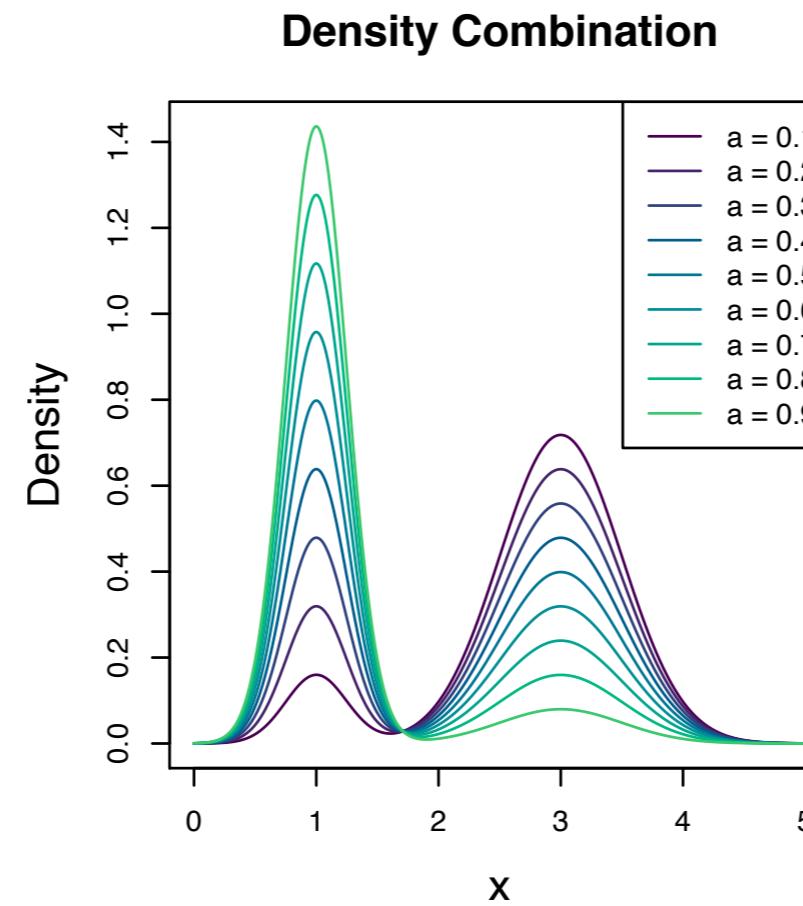
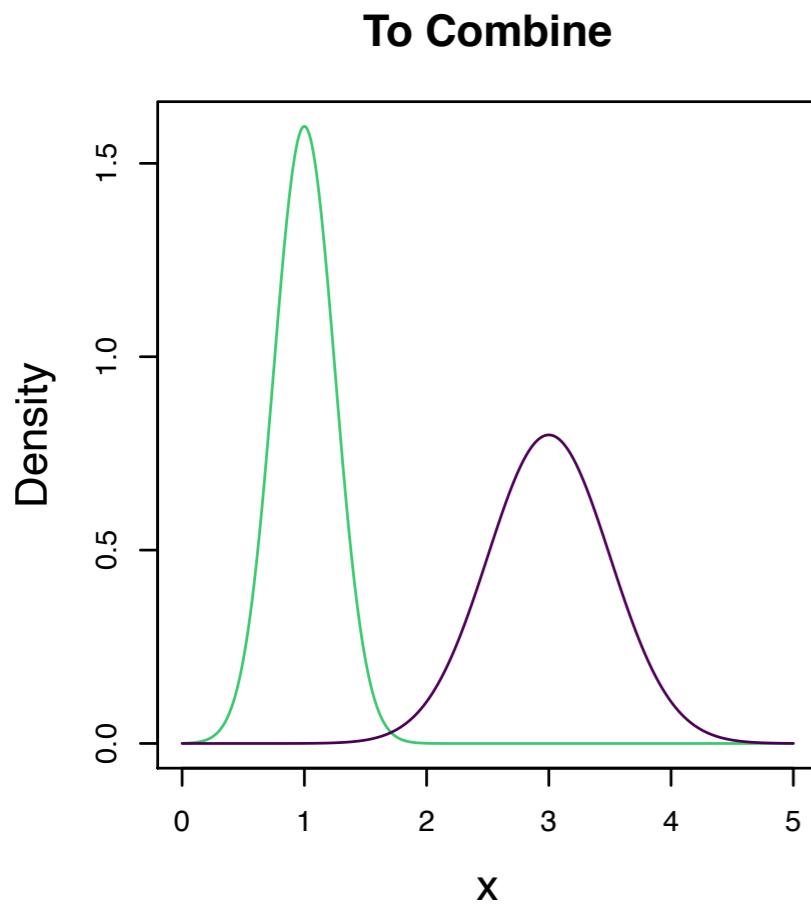
- Models are required to submit 23 quantiles of a predictive distribution:

$$\widehat{P}(Y \leq q_{m,1}) = 0.01, \widehat{P}(Y \leq q_{m,2}) = 0.025, \dots, \widehat{P}(Y \leq q_{m,23}) = 0.99$$

- At each quantile level, we compute the average across all included models

$$q_k = \sum_{m=1}^M w_{m,k} q_{m,k}, \quad k = 1, \dots, 23, \quad \sum_{m=1}^M w_{m,k} = 1, \quad w_{m,k} \geq 0$$

- This is different from computing the average of the predictive density functions (figures adapted from Ryan Tibshirani):



Model evaluation



COVID-19
ForecastHub

Youyang Gu Model Evaluation

covid19-projections.com

		Best-performing models for US state-by-state - 4-week ahead forecasts																					
		We take the weekly projections sent to the CDC for all 50 states (+DC), and compare their incident deaths projections for 4 weeks ahead with each state's true reported incident deaths for those 4 weeks. We take the mean of the absolute errors for each state.																					
		Baseline is a simple model where we take the mean of the previous week's daily deaths to make all future forecasts. For example, for July 20 projections, we use the average daily deaths from July 13 to July 19 for all future projections (this is a constant number).																					
		Projections in red = failed to outperform the baseline model of simply using the previous week's average deaths to make all future projections.																					
		Team-category: Blue = academic institution; purple = private company; orange = government/military; Brown = independent																					
		Note: For fair comparisons, we are only considering models that have 3+ submissions.																					
		Source: https://github.com/yyouyanggu/covid19-forecast-hub-evaluation																					
		Mean absolute error for 4-week ahead forecasts																					
		4/20 - 5/16 4/27 - 5/23 5/4 - 5/10 5/11 - 5/17 5/18 - 5/24 5/25 - 6/1 6/1 - 6/7 6/8 - 6/14 6/15 - 6/21 6/22 - 6/28 6/29 - 7/5 7/6 - 7/12 7/13 - 7/19 7/20 - 7/26 7/27 - 8/2 8/3 - 8/9 8/10 - 8/16 8/17 - 8/23 8/24 - 8/30 8/31 - 8/37 % weeks beat baseline																					
COVIDHub Ensemble (combination of various models)		190 278	100 233	142 132	117 142	122 89	118 99	95 98	87 121	87 138	98 205	175 132	162 140	136 137	113 99	74 80	71 64	53 64	64 67	72 70	330%		
YYG - covid19-projections.com		278	233	132	140	172	142	89	99	98	121	138	205	132	140	137	99	80	64	67	70	330%	
UMass-MB - University of Massachusetts, Amherst		757	272	122	102	62	145	88	97	150	192	176	121	135	194	134	109	76	94	109	79%		
UCLA - University of California, Los Angeles		251	208	107	99	109	120	110	133	193	234	201	200	147	117	105	108	89	89	105	330%		
Oliver Wyman																						95	
MIT_CovidAnalytics - MIT		259	305	298	244	155	127	168	170	152	183	238	184	150	136	145	129	93	122	125	68%		
LANL - Los Alamos National Laboratory		400	318	167	132	129	112	308	100	102	114	160	223	239	254	278	226	161	151	139	132	53%	
Baseline (use previous week's average daily deaths)		351	338	290	309	264	207	171	185	162	165	323	265	233	208	153	120	107	124	110	106	—	
Covid19Sim - COVID-19 Simulator																						77%	
IowaStateUW - Iowa State University		233	321	N/A	153	129	158	130	128	136	231	255	240	267	163	152	N/A	146	172	113	53%		
IHME - Institute for Health Metrics and Evaluation		529	552	238	177	194	N/A	174	134	125	92	137	208	N/A	204	179	165	138	N/A	254	238	47%	
UA - University of Arizona																						47%	
USC - University of Southern California																						80%	
LNC - LockNQuay																						330%	
NEUBS - Northeastern University		620	410	461	463	265	155	133	135	148	183	231	256	249	231	227	420	134	314	125	147	33%	
UT - University of Texas, Austin		564	549	452	N/A	233	129	106	117	127	136	143	229	181	247	454	656	N/A	284	125	N/A	53%	
COVIDHub Baseline		686	332	289	372	274	245	184	175	180	182	343	275	261	230	163	151	134	564	139	129	—	
CMU - Carnegie Mellon University (Delphi)		427	325	221	209	228	208	302	313	271	480	709	293	170	230	225	184	130	171	136	155	25%	
CU-select - Columbia University																		111	111	85	74	81	330%
JOB - John Burau																						330%	
USACE - US Army Engineer Research and Development Center																						44%	
GT-DeepCOVID - Georgia Institute of Technology																						22%	
NotreDame-mobility - University of Notre Dame																						0%	
JHU_JDG - Johns Hopkins University																						6%	
PSI - Predictive Science Inc																						13%	
Avguan - Avguan Data Science																						13%	
Karlen - Karlen Working Group																						14%	
CEID - University of Georgia																						0%	
DDS - Discrete Dynamical Systems																						17%	
Robert Walraven																						0%	
EpiForecasts - LSHTM																						0%	
MITCovAlliance - MIT																						14%	
Median Mean Abs Error		478	328	272	208	233	266	158	152	127	148	209	253	239	207	290	158	129	132	125	213		

Weighted Interval Score

- Consider a single $(1 - \alpha) \times 100\%$ predictive interval $[l, u]$ for the observed response y . The interval score is:

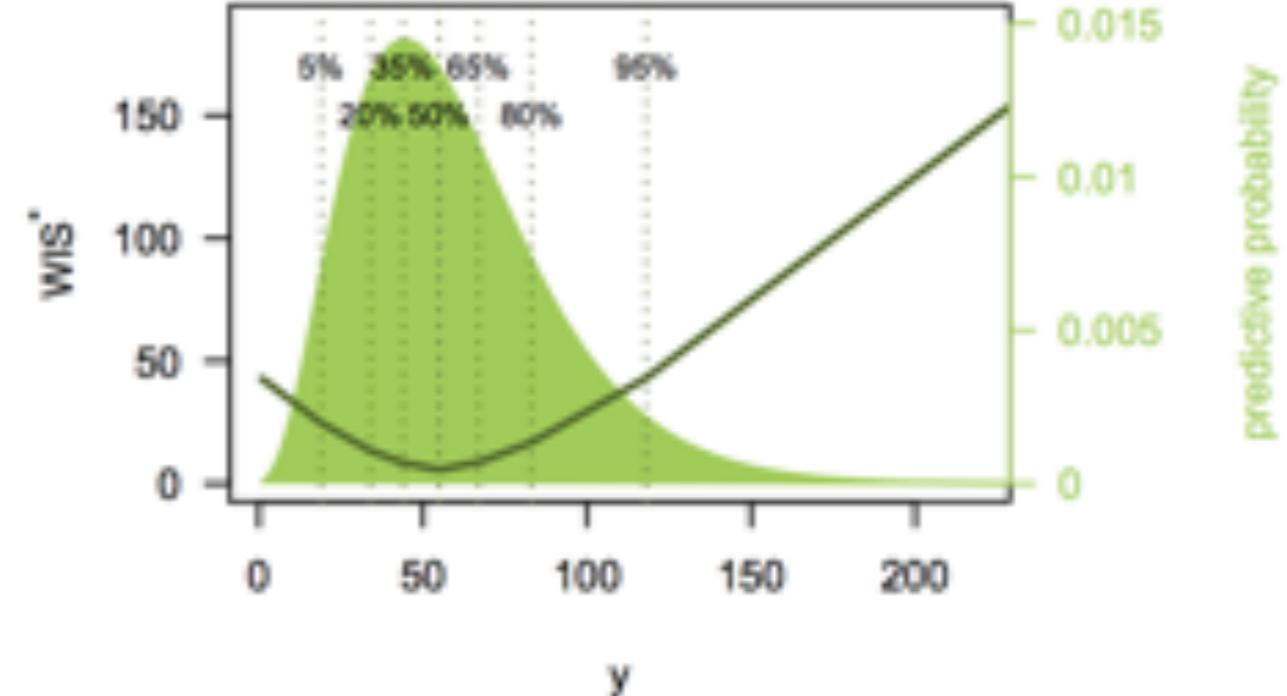
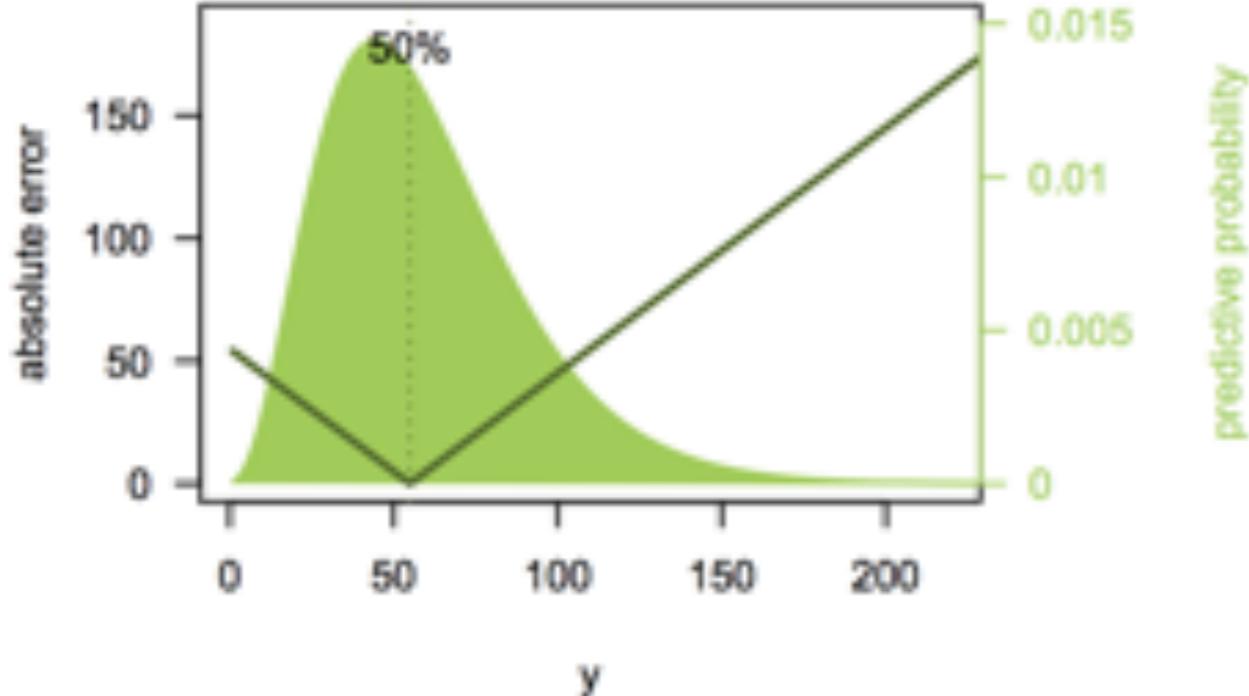
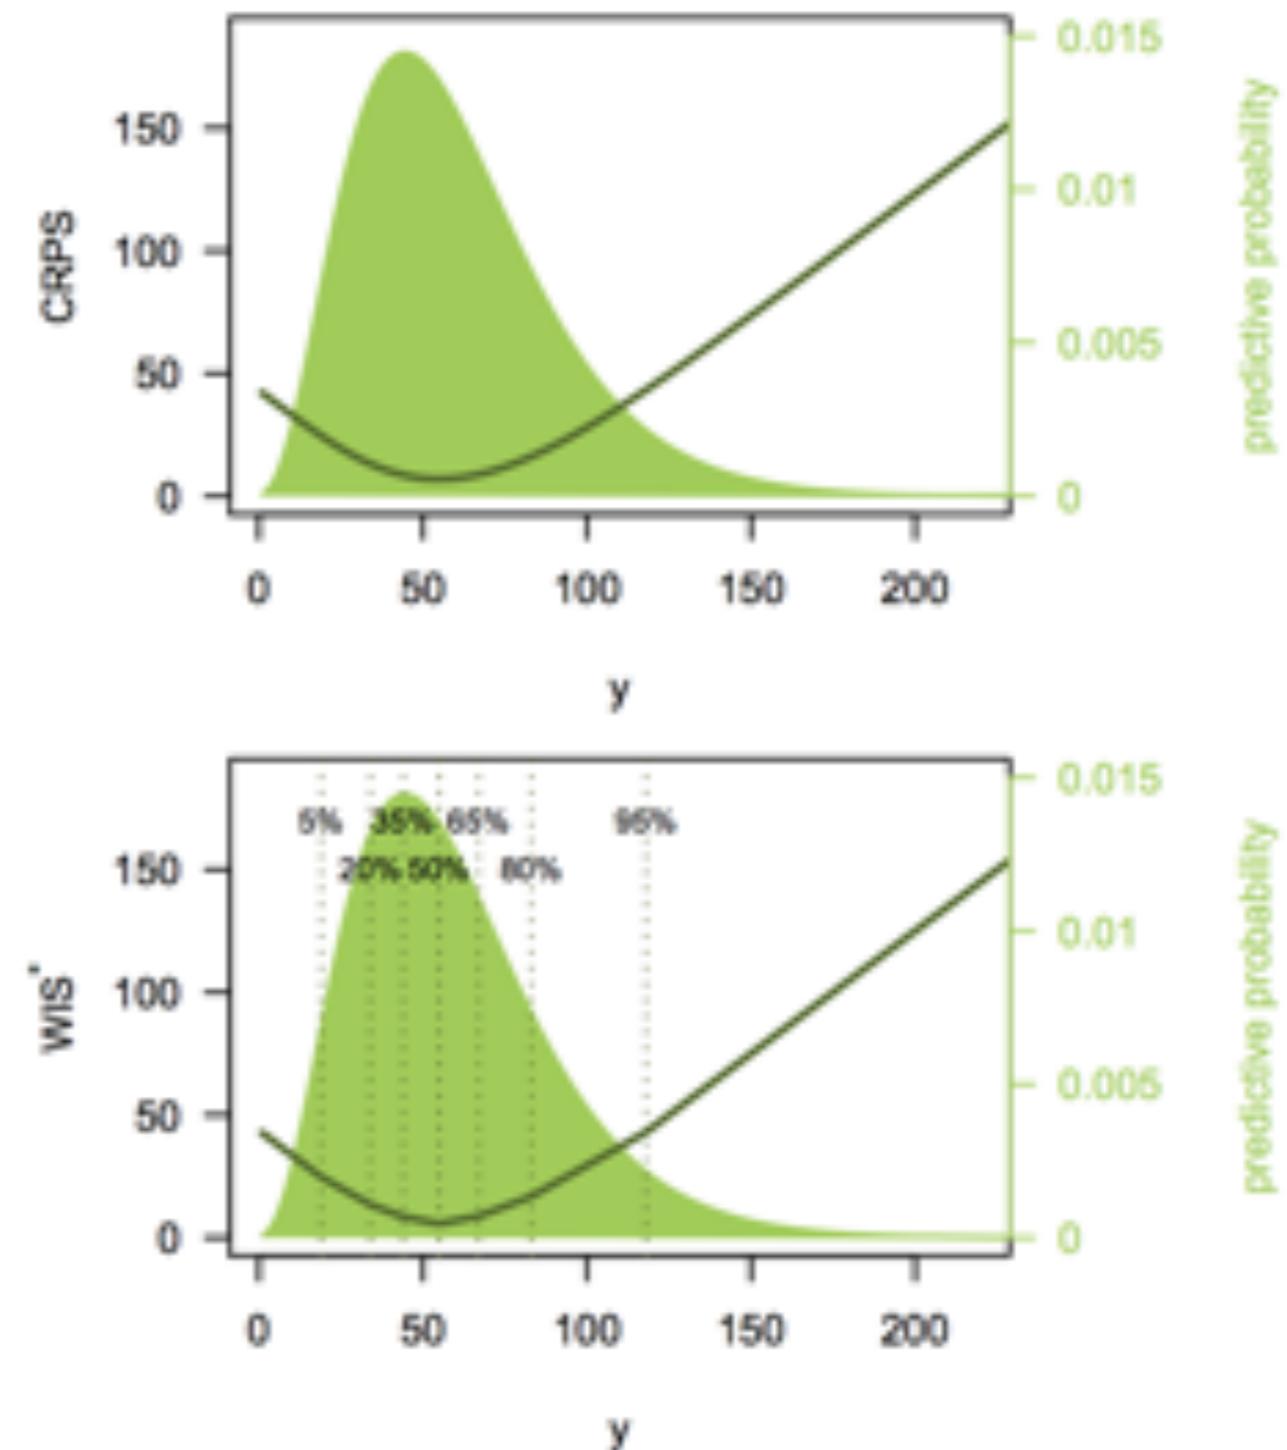
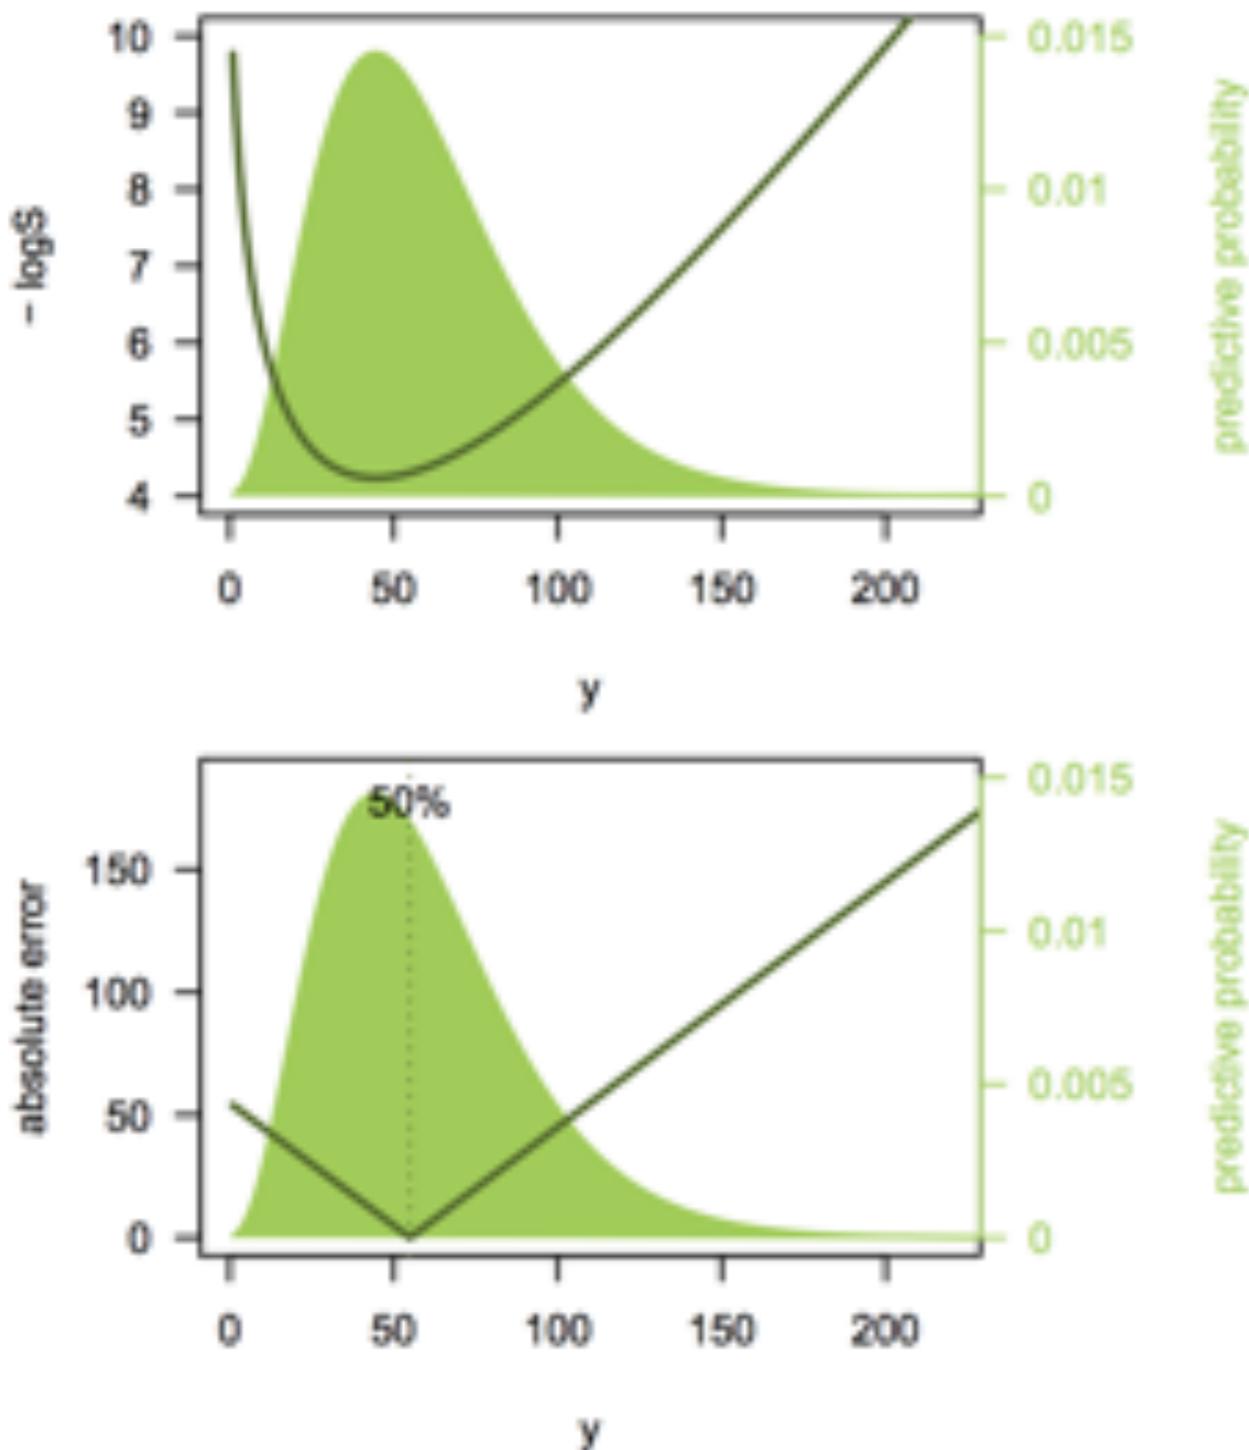
$$\mathbf{IS}_\alpha(F, y) = \underbrace{(u - l)}_{\text{Width of interval}} + \underbrace{\frac{2}{\alpha} \cdot (l - y) \cdot \mathbf{1}(y < l)}_{\text{Penalty if interval is too high}} + \underbrace{\frac{2}{\alpha} \cdot (y - u) \cdot \mathbf{1}(y > u)}_{\text{Penalty if interval is too low}},$$

- Smaller \mathbf{IS}_α is better
- For multiple predictive intervals, we compute a weighted average of \mathbf{IS}_α

$$\mathbf{WIS}_{\alpha_{0:K}}(F, y) = \frac{1}{K+1} \times \left(w_0 \times 2 \times |y - m| + \sum_{k=1}^K (w_k \times \mathbf{IS}_{\alpha_k}(F, y)) \right).$$

- We use weights $w_i = \frac{\alpha_i}{2}$
- The resulting score is **proper** — forecasters don't have any incentive to report anything other than their true beliefs about the future.
- See Bracher et al. (2020) for more: <https://arxiv.org/abs/2005.12881>

WIS compared to other scores



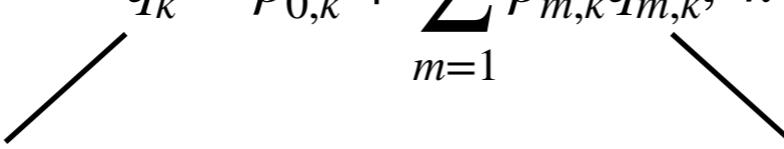
Examples of log score (negatively oriented), continuous ranked probability score (CRPS), absolute error, and the weighted interval score (WIS). Scores are shown as a function of the observed value y . The predictive distribution is Negative Binomial with mean 60 and size 4.

For all scores as shown, smaller values are desirable.

General Quantile Combinations

- General set up:

$$q_k = \beta_{0,k} + \sum_{m=1}^M \beta_{m,k} q_{m,k}, \quad k = 1, \dots, 23$$



k'th predictive quantile for ensemble

k'th predictive quantile for component model number m

- Equally weighted has $\beta_{0,k} = 0$ and $\beta_{m,k} = \frac{1}{M}$
- Estimate parameters by minimizing the Weighted Interval Score achieved by the ensemble in a window of recent weeks subject to some constraints.

Practical challenges

- **Missingness:** not every model makes forecasts at each location or in every week
- **Window size:** How many past weeks of performance are used to estimate model weights? 1, 2, 3, or 4; also, 0 for Equal Weighted
- **Constraints** on parameters

Comparing Ensemble specifications

averaging across 4 targets, all locations

Mean WIS, COVIDhub-ensemble (our actual real-time ensemble forecast): 61.14859

Mean WIS, YYG-ParamSearch (a top performing individual model): 64.94579

Comparing Ensemble specifications

averaging across 4 targets, all locations

Mean WIS, COVIDhub-ensemble (our actual real-time ensemble forecast): 61.14859

Mean WIS, YYG-ParamSearch (a top performing individual model): 64.94579

missingness	constraint	quantile group	window	checks	mean_wis
<chr>	<chr>	<chr>	<int>	<chr>	<dbl>
per location group	equal	per model	3	standard	57.15831
per location group	convex	per quantile	4	standard	57.28399
per location group	convex	3 groups	4	baseline	58.44847
per location group	equal	per model	4	standard	58.80845
per location group	convex	per model	4	standard	59.79603
per location group	convex	3 groups	4	standard	59.92877
impute	convex	per quantile	2	none	60.08238
impute	convex	per quantile	3	standard	60.29250
impute	convex	per quantile	2	standard	60.61362
per location group	equal	per model	0	standard	82.49395

Top approaches have comparable error to our equal-weighted approach.

Comparing Ensemble specifications

averaging across 4 targets, all locations

Mean WIS, COVIDhub-ensemble (our actual real-time ensemble forecast): 61.14859

Mean WIS, YYG-ParamSearch (a top performing individual model): 64.94579

missingness	constraint	quantile group	window	checks	mean_wis
<chr>	<chr>	<chr>	<int>	<chr>	<dbl>
per location group	equal	per model		3 standard	57.15831
per location group	convex	per quantile		4 standard	57.28399
per location group	convex	3 groups		4 baseline	58.44847
per location group	equal	per model		4 standard	58.80845
per location group	convex	per model		4 standard	59.79603
per location group	convex	3 groups		4 standard	59.92877
impute	convex	per quantile	2	none	60.08238
impute	convex	per quantile	3	standard	60.29250
impute	convex	per quantile	2	standard	60.61362
per location group	equal	per model	0	standard	82.49395

Top-ranked approaches
require submission history
and drop models that fail
automated checks

Comparing Ensemble specifications

averaging across 4 targets, all locations

Mean WIS, COVIDhub-ensemble (our actual real-time ensemble forecast): 61.14859

Mean WIS, YYG-ParamSearch (a top performing individual model): 64.94579

missingness <chr>	constraint <chr>	quantile group <chr>	window <int>	checks <chr>	mean_wis <dbl>
per location group	equal	per model	3	standard	57.15831
per location group	convex	per quantile	4	standard	57.28399
per location group	convex	3 groups	4	baseline	58.44847
per location group	equal	per model	4	standard	58.80845
per location group	convex	per model	4	standard	59.79603
per location group	convex	3 groups	4	standard	59.92877
impute	convex	per quantile	2	none	60.08238
impute	convex	per quantile	3	standard	60.29250
impute	convex	per quantile	2	standard	60.61362
per location group	equal	per model	0	standard	82.49395

**COVIDhub-ensemble if we didn't manually
eliminate models based on visual mis-alignment
--> human-in-the-loop has helped.**

Comparing Ensemble specifications

averaging across 4 targets, all locations

Mean WIS, COVIDhub-ensemble (our actual real-time ensemble forecast): 61.14859

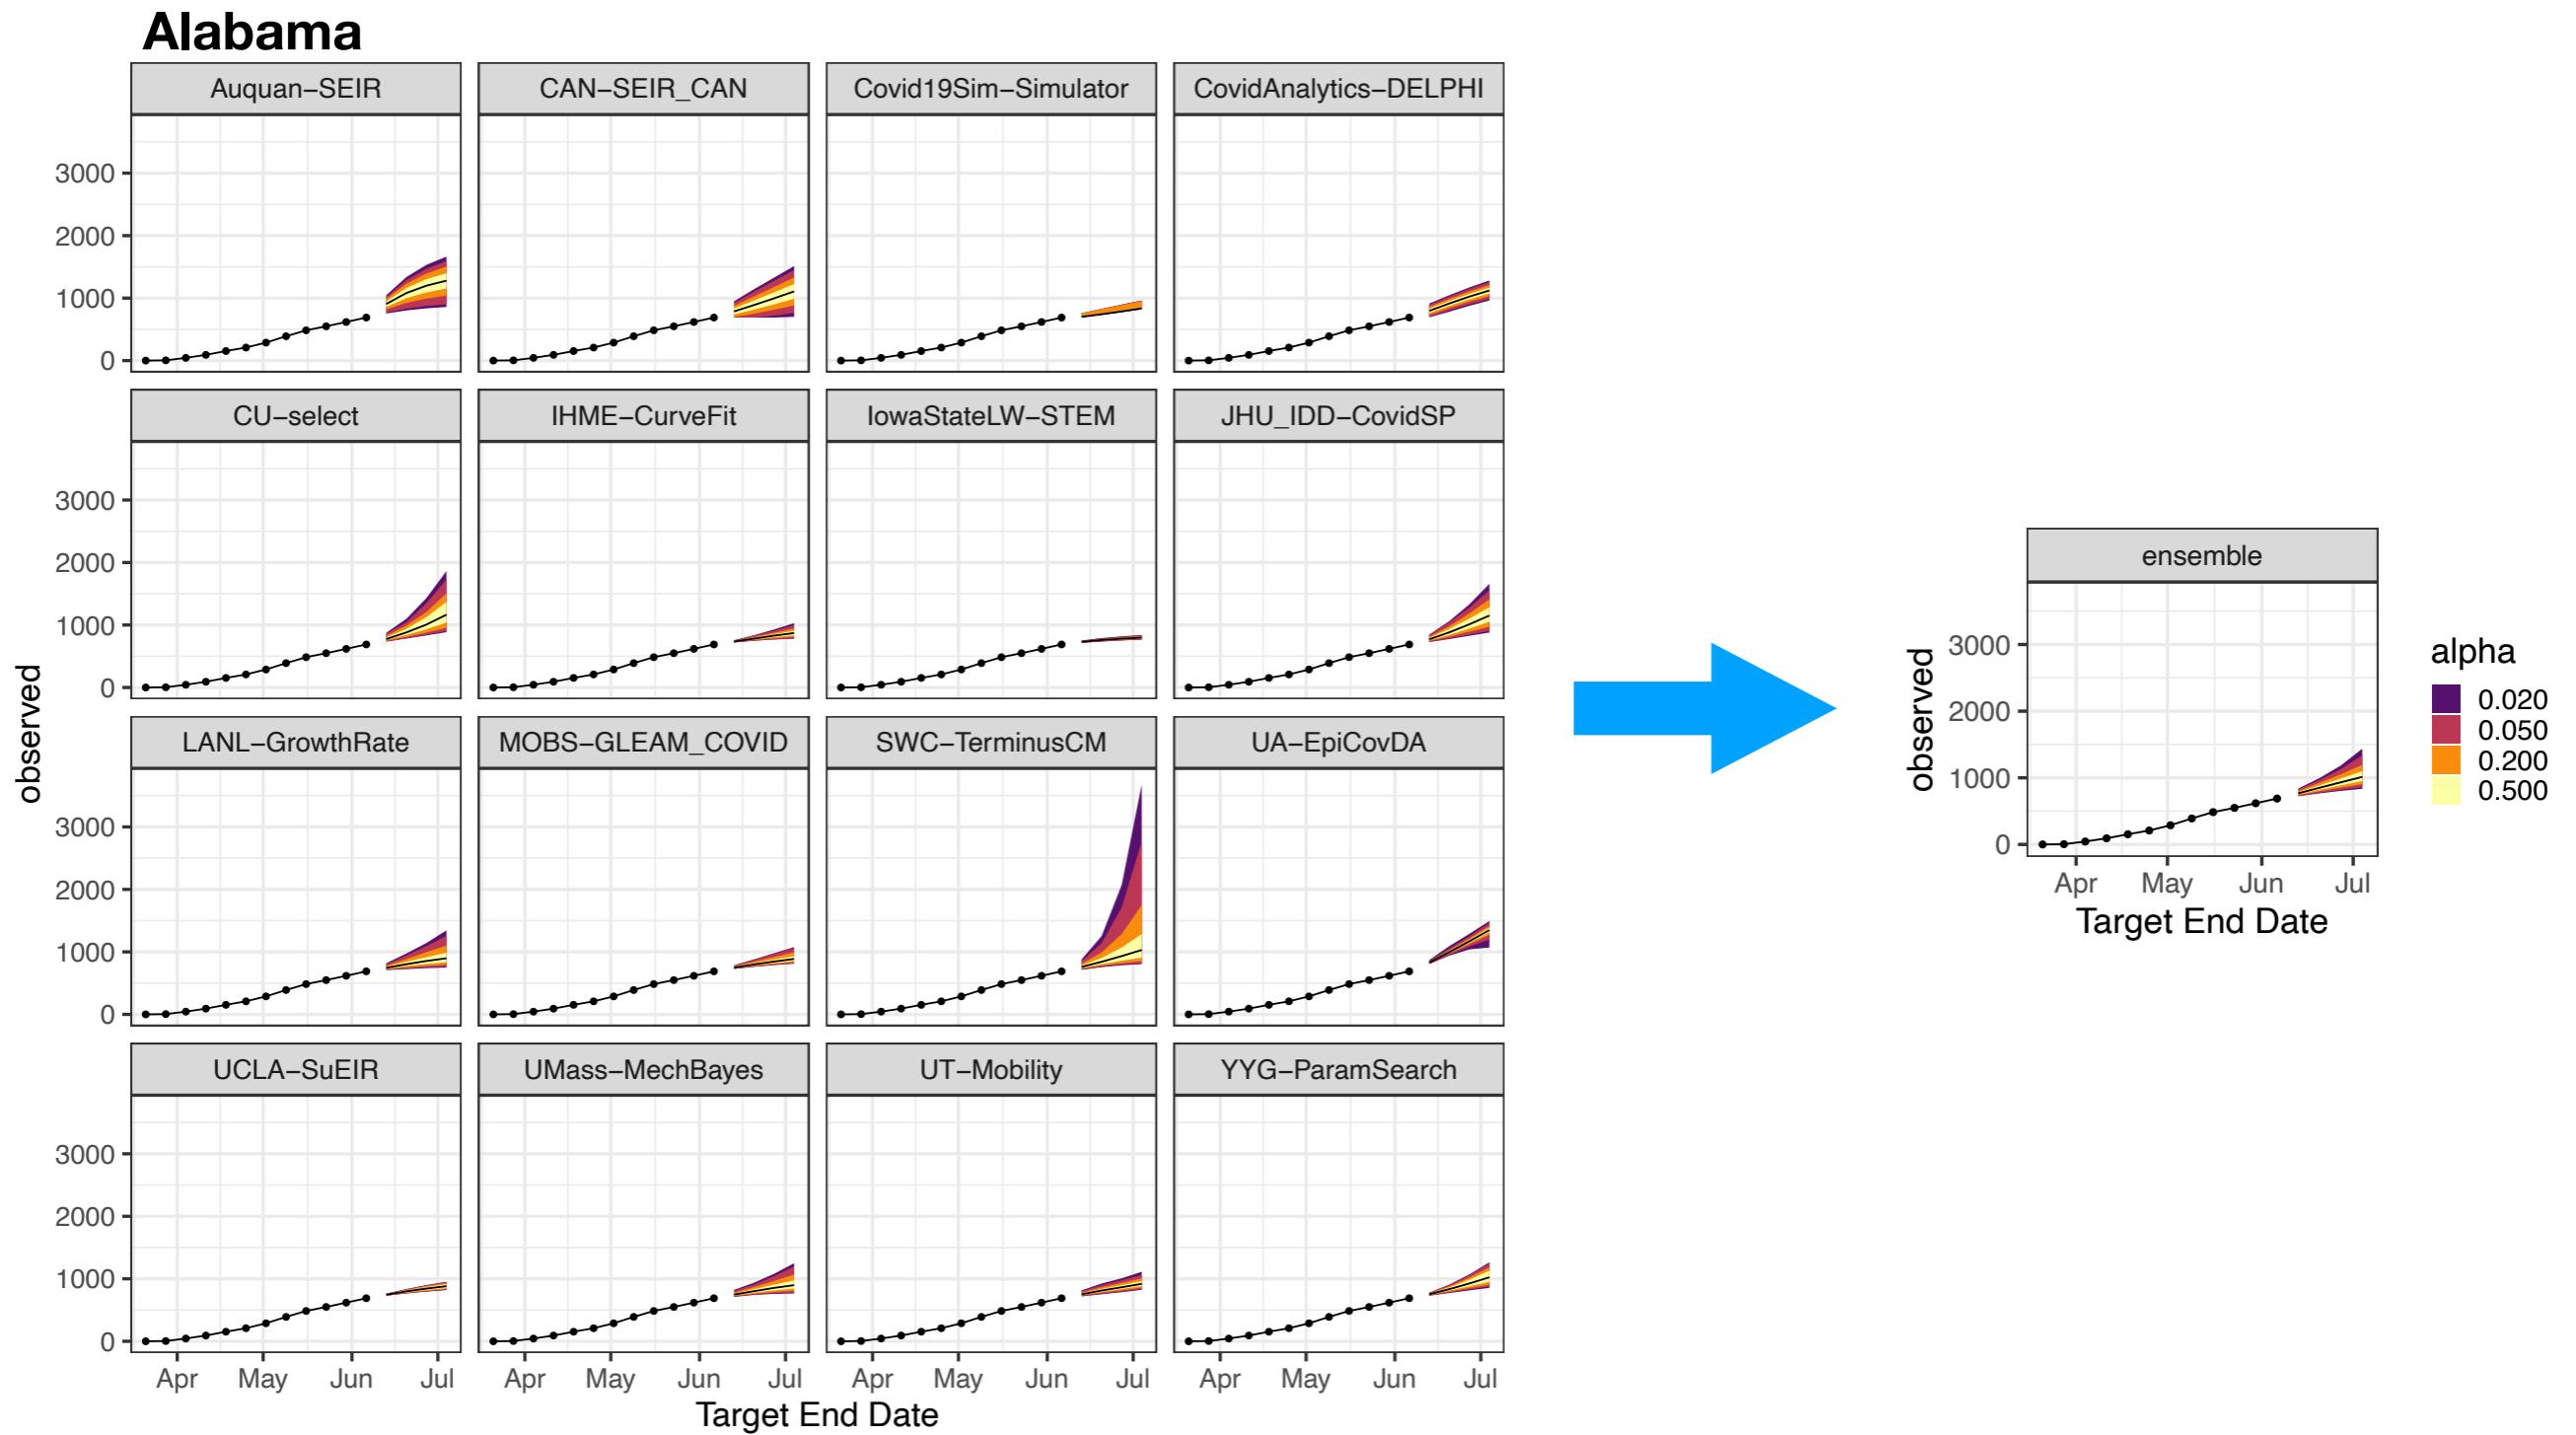
Mean WIS, YYG-ParamSearch (a top performing individual model): 64.94579

missingness	constraint	quantile group	window	checks	mean_wis
<chr>	<chr>	<chr>	<int>	<chr>	<dbl>
per location group	equal	per model	3	standard	57.15831
per location group	convex	per quantile	4	standard	57.28399
per location group	convex	3 groups	4	baseline	58.44847
per location group	equal	per model	4	standard	58.80845
per location group	convex	per model	4	standard	59.79603
per location group	convex	3 groups	4	standard	59.92877
impute	convex	per quantile	2	none	60.08238
impute	convex	per quantile	3	standard	60.29250
impute	convex	per quantile	2	standard	60.61362
per location group	equal	per model	0	standard	82.49395

In recent weeks, approaches that rely only on estimation
to remove models have started to show competitive
performance. (This was not true early on.)

Note: still not much better than our current approach

Result in Pictures



VERY preliminary accuracy results

- Nearly all submitted models **do not beat** baseline model.
- COVID-19 Forecast Hub ensemble
 - is one of a few models that consistently outperforms the baseline
 - has approximately nominal coverage for death forecasts
 - has lower than nominal coverage for case forecasts

Next steps

- Continued crank-turning on collection and aggregation of short-term forecasts for national and state levels.
- Collecting county-level forecasts of cases, in hopes of informing COVID-19 vaccine site selection efforts.
- With funding and collaborative support from CDC, we will continue to use statistical approaches to make accurate short-term forecasts of where this pandemic is taking us.



COVID-19 ForecastHub

www.medrxiv.org/content/10.1101/2020.08.19.20177493v1

Ensemble Forecasts of Coronavirus Disease 2019 (COVID-19) in the U.S.

✉ Evan L Ray, Nutchawattanachit, ✉ Jarad Niemi, Abdul Hannan Kanji, Katie House, ✉ Estee Y Cramer,
✉ Johannes Bracher, Andrew Zheng, ✉ Teresa K Yamana, Xinyue Xiong, Spencer Woody, Yuanjia Wang,
✉ Lily Wang, Robert L Walraven, Vishal Tomar, Katherine Sherratt, Daniel Sheldon, Robert C Reiner,
B. Aditya Prakash, ✉ Dave Osthus, Michael Lingzhi Li, Elizabeth C Lee, Ugur Koayluoglu, Pinar Keskinocak,
Youyang Gu, Quanquan Gu, Glover E George, Guido España, Sabrina Corsetti, ✉ Jagpreet Chhatwal,
Sean Cavany, Hannah Biegel, Michal Ben-Nun, Jo Walker, Rachel Slayton, Velma Lopez, Matthew Biggerstaff,
✉ Michael A Johansson, ✉ Nicholas G Reich, COVID-19 Forecast Hub Consortium

doi: <https://doi.org/10.1101/2020.08.19.20177493>

Questions?